


WHAT HAS MY COMPILER DONE FOR ME LATELY?

UNBOLTING THE COMPILER'S LID

*Matt Godbolt, DRW Trading
@mattgodbolt matt@godbolt.org*

ABOUT ME



MY GOALS

- Un-scary-fy assembler
- Appreciate your compiler!

OUTLINE

- Compiler Explorer story
- Assembly 101
- What has my compiler done for me lately?
- Behind the scenes of Compiler Explorer

BACKSTORY

```
int sum(const vector<int> &v) {  
    int result = 0;  
    for (size_t i = 0; i < v.size(); ++i)  
        result += v[i];  
    return result;  
}
```

```
int sum(const vector<int> &v) {  
    int result = 0;  
    for (int x : v) result += x;  
    return result;  
}
```

Is one better than the other?

WARNING


- Reading assembly alone can be misleading
- *Always* measure too
- Google Benchmark
- quick-bench.com

X86 ASSEMBLY 101

REGISTERS

- rax, rbx, rcx, rdx, rsp, rbp, rsi, rdi, r8-r15
- xmm0-xmm15
- rdi, rsi, rdx... arguments
- rax is return value

REGISTERS



INSTRUCTIONS

```
op  
op dest  
op dest, src  
op dest, src1, src2
```

- op is e.g. call, ret, add, sub, cmp...
- dest, src is register or memory reference:
[base + reg1_{opt} + reg2*(1, 2, 4 or 8)_{opt}]

(Intel asm syntax)

INSTRUCTIONS

```
mov  eax, DWORD PTR [r14]
add  rax, rdi
add  eax, DWORD PTR [r14+4]
sub  eax, DWORD PTR [r14+4*rbx]
lea  rax, [r14+4*rbx]
xor  edx, edx
```

```
int  eax = *r14;    // int *r14;
rax += rdi;
eax += r14[1];
eax -= r14[rbx];
int  *rax = &r14[rbx];
edx = 0;
```

SUMMARY

- Registers: rax, rbx, rcx ...
- Size: rax, eax, ax ...
- Params: rdi, rsi, rdx, rcx ...
- Result: rax
- op dest, src
- dest, src are registers or memory

WHERE WERE WE?

```
int sum(const vector<int> &v) {  
    int result = 0;  
    for (size_t i = 0; i < v.size(); ++i)  
        result += v[i];  
    return result;  
}
```

```
int sum(const vector<int> &v) {  
    int result = 0;  
    for (int x : v) result += x;  
    return result;  
}
```

Which is better?

COMPILER EXPLORER V0.1

```
$ g++ /tmp/test.cc -O2 -c -S -o - -masm=intel \  
  | c++filt \  
  | grep -vE '\s+\.'
```


```
sum(std::vector<int, std::allocator<int> > const&):  
.LFB786:  
    mov rcx, QWORD PTR [rdi]  
    mov rax, QWORD PTR 8[rdi]  
    sub rax, rcx  
    shr rax, 2  
    mov rsi, rax  
    ...
```

COMPILER EXPLORER V0.1

Not very pretty
TO THE WEB!

DEMO

A ▾



```
1 // setup...
5 int sum(const vector<int> &v) {
6     int result = 0;
7     for (size_t i = 0; i < v.size(
8         result += v[i];
9     return result;
10 }
11
```

x86-64 gcc 7.1 ▾

-O2 -std=c++1z -march=haswell

11010

.LX0:

.text


//

\s+

Intel

Demangle

A ▾



1

[Edit on C++ Compiler Explorer](#)
(/)

WALKTHROUGH

```
; rdi = const vector<int> *  
mov rdx, QWORD PTR [rdi] ; rdx = *rdi      ≡ begin()  
mov rcx, QWORD PTR [rdi+8] ; rcx = *(rdi+8) ≡ end()
```

```
template<typename T> struct _Vector_impl {  
    T *_M_start;  
    T *_M_finish;  
    T *_M_end_of_storage;  
};
```

TRADITIONAL

```
sub rcx, rdx ; rcx = end-begin
mov rax, rcx
shr rax, 2   ; (end-begin)/4
je .L4
add rcx, rdx
xor eax, eax
```

```
size_t size() const noexcept {
    return _M_finish - _M_start;
}
```

RANGE

```
xor eax, eax
cmp rdx, rcx ; begin==end?
je .L4
```

```
auto __begin = begin(v);
auto __end = end(v);
for (auto __it = __begin;
     __it != __end;
     ++it)
```

WALKTHROUGH

```
; rcx ≡ end, rdx = begin, eax = 0
.L3:
    add eax, DWORD PTR [rdx]    ; eax += *rdx
    add rdx, 4                  ; rdx += sizeof(int)
    cmp rdx, rcx                ; is rdx == end?
    jne .L3                     ; if not, loop
    ret                         ; we're done
```

BACKSTORY

SO, WHICH APPROACH IS BEST?

ALSO

- Optimizer settings are important
- `std::accumulate`

**WHAT HAS MY COMPILER DONE
FOR ME LATELY?**

MULTIPLICATION

```
int mulByY(int x, int y) {  
    return x * y;  
}
```

[View](#)

```
mulByY(int, int):  
    mov eax, edi  
    imul eax, esi  
    ret
```

MULTIPLICATION

$$\begin{array}{r} 1101 \quad (13) \\ \times 0101 \quad (5) \\ \hline 1101 \\ 0000 \\ 1101 \\ + 0000 \\ \hline 01000001 \quad (65) \end{array}$$

That's a lot of additions!
Haswell 32-bit multiply - 4 cycles

MULTIPLICATION

```
int mulByConstant(int x) { return x * 2; }
```

[View](#)

MULTIPLICATION

```
int mulBy65599(int a) {  
    return (a << 16) + (a << 6) - a;  
    //           ^           ^  
    //      a * 65536      |  
    //                   a * 64  
    // 65536a + 64a - 1a = 65599a  
}
```

[View](#)

DIVISION

```
int divByY(int x, int y) {  
    return x / y;  
}  
int modByY(int x, int y) {  
    return x % y;  
}
```

[View](#)

Haswell 32-bit divide - 22-
29 cycles!

```
divByY(int, int):  
    mov eax, edi  
    cdq  
    idiv esi  
    ret  
modByY(int, int):  
    mov eax, edi  
    cdq  
    idiv esi  
    mov eax, edx  
    ret
```

DIVISION

```
unsigned divByConstant(unsigned x) { return x / 2; }
```

[View](#)

DIVISION

```
mov  eax, edi          ; eax = x
mov  edx, 0xaaaaaaaaab
mul  edx               ; eax:edx = x * 0xaaaaaaaaab
mov  eax, edx          ; (x * 0xaaaaaaaaab) >> 32
                        ; ≡ (x * 0xaaaaaaaaab) / 0x100000000
                        ; ≡ x * 0.66666666667
shr  eax               ; x * 0.3333333333
ret
```

MODULUS

```
int modBy3(unsigned x) {  
    return x % 3;  
}
```

[View](#)

```
mov  eax, edi  
mov  edx, 0xaaaaaaaaab  
mul  edx  
mov  eax, edx  
shr  eax  
lea  eax, [rdx+rdx*2]  
sub  edi, eax  
mov  eax, edi  
ret
```

WHY MODULUS?

- Bucket selection in hash maps
- `libc++` special-cases power-of-two
- `boost multi_index`

COUNTING BITS

```
int countSetBits(int a) {  
    int count = 0;  
    while (a) {  
        count++;  
        a &= (a-1);  
    }  
    return count;  
}
```

[View](#)

SUMMATION

```
constexpr int sumTo(int x) {  
    int sum = 0;  
    for (int i = 0; i <= x; ++i)  
        sum += i;  
    return sum;  
}  
int main(int argc, const char *argv[]) {  
    return sumTo(20);  
}
```

[View](#)

SUM(X)

$$\begin{aligned}\sum x &\equiv \frac{x(x+1)}{2} \\ &\equiv x + \frac{x(x-1)}{2}\end{aligned}$$

WHAT HAS MY COMPILER DONE FOR ME LATELY?

A lot!

HOW IT WORKS

HOW IT WORKS - BACKEND

- Written in `node.js`
- Runs on Amazon

NODE.JS

```
function compile(req, res, next) {  
  // exec compiler, feed it req.body, parse output  
}  
var webServer = express();  
var apiHandler = express.Router();  
apiHandler.param('compiler',  
  function (req, res, next, compiler) {  
    req.compiler = compiler;  
    next();  
  });  
apiHandler.post('/compiler/:compiler/compile', compile);  
webServer.use('/api', apiHandler);  
webServer.listen(10240);
```

AMAZON EC2

- Edge cache
- Load balancer
- Virtual machines
- Docker images
- Shared compiler storage

THE COMPILERS

- Built through docker images
- Compilers stored on S3
- OSS ones publically available
- MS compilers via WINE

HOW IT WORKS - SECURITY

- Compilers: huge attack vector
- Principle of "what's the worst could happen"
- Docker
- LD_PRELOAD

HOW IT WORKS - FRONTEND

- Microsoft's Monaco
- GoldenLayout

THE CODE

- github.com/mattgodbolt/compiler-explorer
- github.com/mattgodbolt/compiler-explorer-image
- Running locally is easy!

```
$ make
```

- More in next C++ Weekly

OTHER USES

- Code pastebin
- Compiler development
- C++ REPL
- Training resource

COMING SOON...

- CFG viewer
- Unified languages
- Execution support

THANKS

- Thanks to contributors:
 - Rubén Rincón
 - Gabriel Devillers
 - Simon Brand, Johan Engelen, Jared Wyles, Chedy Najjar
 - ...and the rest!
- Thanks to Patreon folks
- Thanks to awesome C++ community
- Thanks to you!

GO READ SOME ASSEMBLY!

{ gcc , cppx , d , swift , haskell , go , ispc }.godbolt.org

(AND THANK A COMPILER DEVELOPER)