

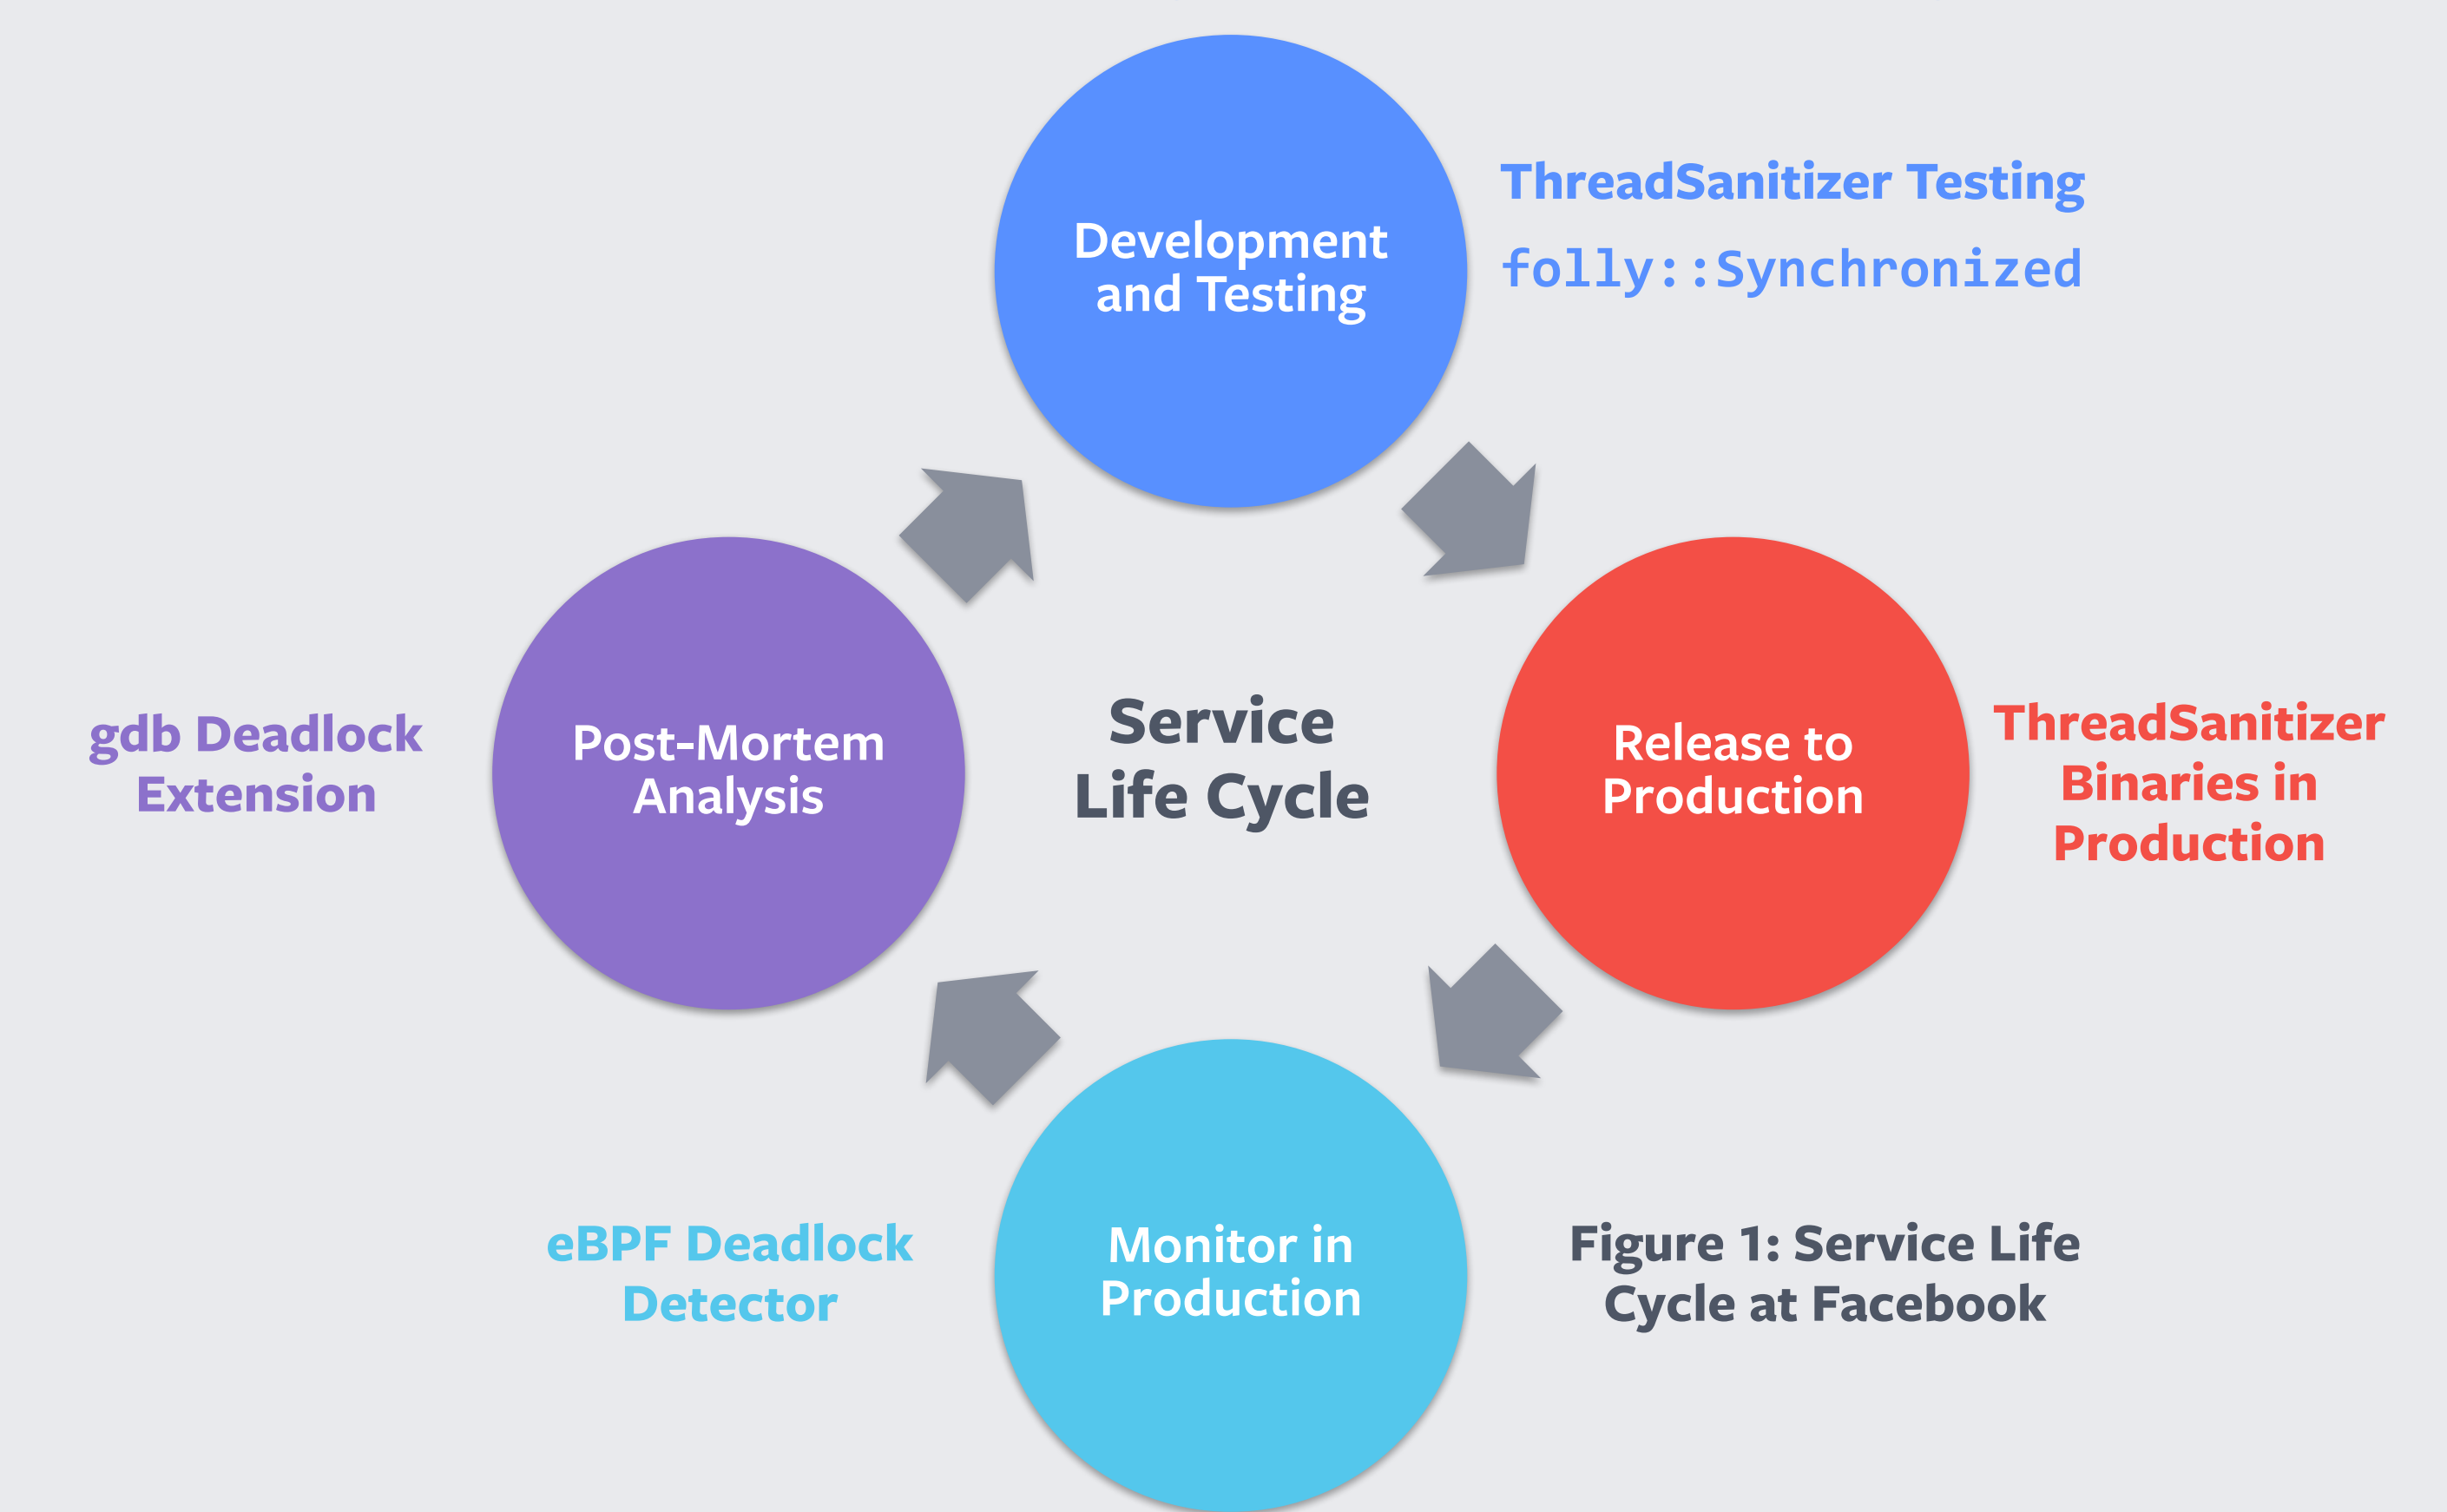
Motivation: C++ and Deadlocks at Facebook

Facebook has an enormous and complex C++ codebase. Facebook uses C++ primarily for developing backend services. Examples of backend services include newsfeed ranking, ads personalization, search, and messaging.

With such a large codebase and huge scale of deployments, it can be difficult to detect and prevent bugs in production services at Facebook. Deadlocks in particular are worrisome—they can be difficult to detect, as the service may appear alive but is failing requests. Further, deadlocks can be easily introduced unintentionally and can be difficult to identify in a large codebase. What tools can we use to make detecting, debugging, and preventing deadlocks easier?

Approach: “End-to-end” Debugging Tools

Service developers at Facebook are responsible for all phases of the service life cycle. We have developed end-to-end tooling to help service developers detect deadlocks throughout all of the phases of the service life cycle (Figure 1). At Facebook, these tools have helped catch existing deadlocks in production for multiple services, and have helped prevent new regressions from ever reaching production.



Tool	Need to recompile?	Can use on dead process?	Overhead?	Need to rewrite code?	Works for...	Bugs it can catch
folly::Synchronized	Yes	No	None	Yes	Lockable, SharedLockable traits	Data race
ThreadSanitizer	Yes	No	Huge. Up to 10x memory, 15x slowdown	No	std::mutex, and other mutexes with TSAN annotations	Many types of synchronization bugs, including lock inversion
eBPF deadlock detector	No	No	None when tool is not running. Huge when on.	No	pthread_mutex_t, std::mutex	Lock inversion
gdb deadlock extension	No	Yes	None	No	pthread_mutex_t, std::mutex, other mutexes with metadata about ownership	Deadlock

Figure 2: Pros and cons of the deadlock debugging tools and when to use each tool

How can we prevent deadlocks during development and testing?

Development and Testing

Approach: Use folly::Synchronized.

- C++ library makes it impossible to access data without acquiring corresponding mutex.
- API encourages smaller critical sections, reducing the probability of deadlocks.

Figure 3: folly::Synchronized example to implement a counter with the history of changes to the count.

```
class Counter {
void add(int n) {
state->wlock([&n](auto& state) {
state.counter += n;
state.deltas.push_back(n);
});
}
int get() const {
return state->rlock([](const auto& state) {
return state.counter;
});
}
private:
struct State {
int counter;
std::vector<int> deltas;
};
folly::Synchronized<
State, std::shared_mutex> state_;
};
```

The lines in red are the critical sections.

Approach: Run all tests with ThreadSanitizer (TSAN).

- ThreadSanitizer instruments all memory access and mutex acquisitions and can report data races, lock inversions, and other types of synchronization bugs.
- Recompile and run all existing tests with TSAN. Developers must wait for all tests to pass during code review before landing the change.

Success (6 Success) Hide Builds

Figure 4: Screenshot of Facebook's continuous integration builds during code review

How can we detect deadlocks before releasing to all of production?

Release to Production

Approach: Push ThreadSanitizer builds of release candidates to a subset of production (canary).

- Send a subset of production traffic or shadow traffic to ThreadSanitizer builds.
- This allows us to exercise production code paths with ThreadSanitizer and find bugs that cannot be caught with normal testing.

1. Subset of production traffic, or

2. Shadow Traffic: record and replay old requests, discard responses

Figure 5: How we canary TSAN builds to a subset of production

How can we monitor for deadlocks in a running production service?

Monitor in Production

Constraints: We cannot restart or recompile the service.

Approach: Use Linux eBPF to trace mutex acquisitions.

- Linux eBPF allows us to run custom hooks whenever a process hits the specified userspace or kernel symbol.
- Maintain a map of thread to currently held mutexes for that thread. Anytime a mutex B is acquired while mutex A is already held, add edge (A, B). There is a lock inversion (potential deadlock) if there is a cycle in the graph.
- To monitor for deadlocks in production, we periodically run the detector on a subset of production machines and log the results if a lock inversion is detected.

Figure 6: Overview of the eBPF deadlock detector tracing a running process

How can we detect deadlocks after the service has deadlocked or crashed?

Post-Mortem Analysis

Approach: Create a gdb extension to analyze mutex internals.

- Our gdb extension introduces a new deadlock command that can be used with a hung process or core dump.
- The NPTL implementation of pthread_mutex_t keeps track of the owner of the mutex. The extension uses this information to build the waits-for graph.
- Edge (A, B) exists if thread A is waiting on a mutex held by thread B. There is a deadlock if there is a cycle in this graph.

```
$ gdb -p 3406029
(gdb) deadlock
Found deadlock!
Thread 4 (LWP 3406032) is waiting on mutex (0x00007ffffd0bed7c8) held by Thread 5 (LWP 3406033)
Thread 5 (LWP 3406033) is waiting on mutex (0x00007ffffd0bed818) held by Thread 3 (LWP 3406031)
Thread 3 (LWP 3406031) is waiting on mutex (0x00007ffffd0bed7f0) held by Thread 4 (LWP 3406032)
```

```
(gdb) thread 3
(gdb) bt
#0 __lll_lock_wait ()
#1 0x00007fb051d4cf44 in pthread_mutex_lock (mutex=0x7ffffd0bed7f0)
...
(gdb) frame 1
(gdb) p *mutex
$1 = {__data = {__lock = 2, __count = 0, __owner = 3406032, ... } ...},
(gdb) thread find 3406032
Thread 4 has target id 'Thread 0x7fb04f14a700 (LWP 3406032)'
```

Figure 7: Example output from the gdb deadlock extension. The extension inspects the mutex internals to find the current owner of a mutex.