

API & ABI **versioning**

How to handle compatibility with
your C++ libraries





**When I change my code,
what are the impacts?**



About this talk

- Semver!
- Binary compatibility!





About this **talk** (for real)

- Changes and impacts on API & ABI
- Categorizing changes
- Avoiding impacts
- Advertising change through versioning



Hello!

*I am **Mathieu Ropert***

I'm a C++ developer at Murex where I work on internal frameworks and Open Source initiatives.

You can reach me at:

✉ mro@puchiko.net

🐦 [@MatRopert](https://twitter.com/MatRopert)

🐙 [@mropert](https://github.com/mropert)

1

Library lifecycle

Asking yourself the right questions



So you want to publish a **library**

- Will all users' code belong to the same repo as your library?
- If yes, versioning is not mandatory
- But even then, it will not hurt to think about the impacts



So you want to publish a **library**

- Will you ever break backward compatibility?
- Remember that removing old / deprecated features is still breaking compatibility
- If you do it, even rarely, you need a way to distinguish changes



So you want to publish a **library**

- Do you want your users to be able to hotswap your library in production?
- Not an option for header-only libraries
- If the answer is “yes”, you will have to monitor ABI changes



Things to keep in **mind**

- It's important to be careful when changing API
 - Even if you can patch all your clients at once
- If binary compatibility is a concern, you also need to keep an eye on ABI impacts
- You'll need to inform your users about changes and their impacts



Versioning

Communication between maintainers and users
about the changes in a software



Reasonable use

- Some users will expect unreasonable guarantees from your code
 - Line numbers
 - Symbol addresses (and being able to get them)
 - Real type of `auto` types
 - Layout of private members
- This talk is not about how to handle that

2

Changes in API

Contracts and how not to breach them



What's an **API**?

- An API is a contract between the maintainer and the user
- It's divided in two parts
 - Pre-conditions: what the caller must provide
 - Post-conditions: what the callee will ensure if the pre-conditions are met

std::swap

Defined in header <algorithm>
Defined in header <utility>

(until C++11)
(since C++11)

```
template< class T >
void swap( T& a, T& b );
template< class T >
void swap( T& a, T& b ) noexcept( /* see below */ );
template< class T2, std::size_t N >
void swap( T2 (&a)[N], T2 (&b)[N] ) noexcept( /* see below */ );
```

(until C++11)

(1)

(since C++11)

(2)

(since C++11)

Exchanges the given values.

1) Swaps the values a and b. This overload does not participate in overload resolution unless `std::is_move_constructible_v<T>` && `std::is_move_assignable_v<T>` is `true`. (since C++17)

2) Swaps the arrays a and b. In effect calls `std::swap_ranges(a, a+N, b)`. This overload does not participate in overload resolution unless `std::is_swappable_v<T2>` is `true`. (since C++17)

Parameters


a, b - the values to be swapped

Type requirements

- T must meet the requirements of `MoveAssignable` and `MoveConstructible`.
- T2 must meet the requirements of `Swappable`.

Return value

(none)





API in C++ terms

Internal

- Names
- Signatures
- Declarations locations

External

- Pre-conditions
- Post-conditions
- Misc guarantees



API in C++ terms

- Not all parts of an API are part of the language or seen by the compiler
- You must rely on some form of documentation to express the rest
- Caution is advised when changing parts not covered by the language itself



API changes by **impacts**

- API breaking change
 - Clients must adapt their code
- API non-breaking change
 - Guaranteed to be backward compatible, but not always forward compatible
- No change to API
 - Guaranteed to be both backward and forward compatible



Changes with no **impact**

- Any change that does not add, remove or change a contract
- Changes to implementation
 - Bugfixes
 - Performance tuning
 - Refactoring



Changes with no **impact**

- No name or signature has changed or moved
- Defined behaviour is still the same...
- ...including specific guarantees
 - Complexity
 - Iterator validity



API non-breaking changes

- Adding a new contract
 - New function
 - New overload(*)
 - New type
 - New namespace



API non-breaking changes

- Relaxing an existing contract
 - New default argument to a function(*) or template
 - New struct member
 - Relaxing pre-conditions
 - Narrowing post-conditions
 - Narrowing guarantees
 - Defining undefined behaviour



API breaking changes

- Changing a signature
 - Argument types or order
 - Return type
- Renaming
- Moving declaration to another header file



API breaking changes

- Narrowing a contract
 - Narrowing pre-conditions
 - Relaxing post-conditions
 - Relaxing existing guarantees



API breaking changes

- Narrowing a contract
 - Narrowing pre-conditions
 - Relaxing post-conditions
 - Relaxing existing guarantees
- Evil!



API breaking changes

- Narrowing a contract
 - Narrowing pre-conditions
 - Relaxing post-conditions
 - Relaxing existing guarantees
- Evil!
- Seriously, don't do that



Invisible breaking change

Before

```
// Sorts a vector of integers
// Complexity:  $O(n * \log n)$ 
void foo(std::vector<int>& v) {
    std::sort(begin(v), end(v));
}
```

After

```
// Sorts a vector of integers
// Complexity:  $O(n!)$ 
void foo(std::vector<int>& v) {
    while (!std::is_sorted(begin(v), end(v)))
        std::random_shuffle(begin(v), end(v));
}
```

3

Changes in ABI

Compatibility between binaries



What is ABI?

- Application Binary Interface
- Defines how binary components talk to each others
- Not covered by the C++ Standard(*)



ABI in C++ terms

Infrastructure

- Calling convention
- Exception handling
- Mangling
- C++ runtime

Code

- Symbol names
- Binary representation of API types
- vtable layout



Symbol names

- Each exported symbol has an id:

Name + Signature => id

```
void foo(int)      => _Z3fooi
```

```
void foo(double) => _Z3food
```



Symbol **names**

- Changing the id of any public symbol will break ABI
- Public symbols are all API symbols *and* all symbols used by inline functions in public headers



Implementation changes

Before

```
namespace details {  
    MY_EXPORT void bar();  
};  
  
inline void foo() {  
    details::bar();  
}
```

After

```
namespace details {  
    MY_EXPORT void bar(int);  
};  
  
inline void foo() {  
    details::bar(0);  
}
```





vtable **Layout**

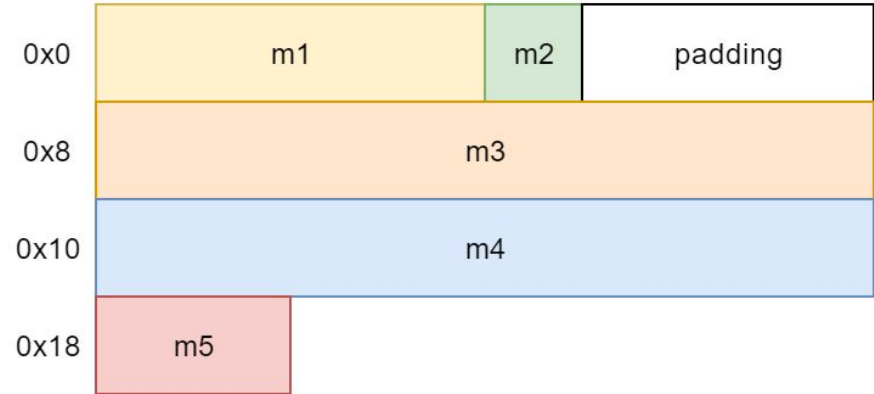
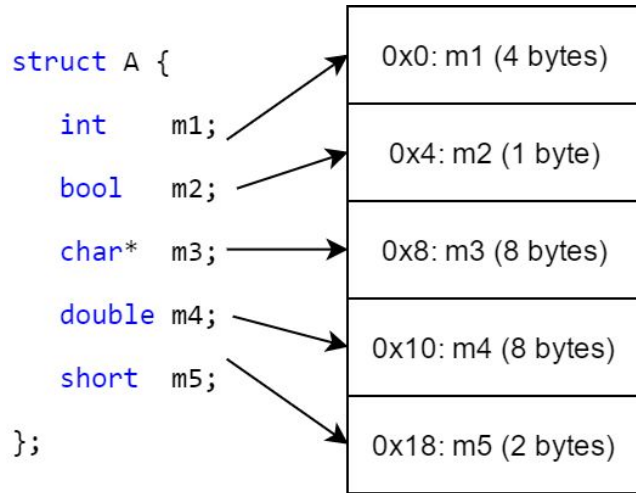

- How pointers to virtual methods are stored
- Depends on the compiler
 - Usually one standard per OS
- Breaks when you reorder virtual methods
- Or when you add a new one



Binary representation

- Each public structure has a particular layout in the ABI
 - Structure size
 - Size of each member
 - Starting offset of each member
- Actual layout depends on various platform rules







Binary representation

- Changing the type or the order of members in a struct will break ABI
- Adding a new member will break it too
- Changing a member visibility may also break ABI

4

C++ Versioning

Semver reloaded



Semantic Versioning

- A formal convention to express compatibility between versions
- Created in 2011 by Tom Preston-Werner
- 3 numbers sequence: X.Y.Z
- X is major release
- Y is minor release
- Z is patch release



Major release

- 1st component of *semver* convention
- Indicates an important and *non-backward-compatible* change
- Users will have to change their code to upgrade or downgrade
- Some features may not be available anymore



Minor release

- 2nd component of *semver* convention
- Indicates the addition of new features that do not impact the existing ones
- Existing users can safely upgrade without changing their code
- Downgrading is also possible if the new features are not used



Patch release

- 3rd component of *semver* convention
- Indicates the release contains only bugfixes
- Existing users can safely upgrade without changing their code
- Downgrading is also possible (although usually not recommended)



How to include **API** in versioning?

- Follow semver convention
- Maintain a changelog
- Document your contracts
- Avoid invisible breaking changes



How to include **ABI** in versioning?

- Don't!
 - If your clients always recompile, don't bother
 - If your library is header only
 - But make it clear in your documentation



How to include **ABI** in versioning?

- Don't!
 - If your clients always recompile, don't bother
 - If your library is header only
 - But make it clear in your documentation
- Or adapt semver convention to include ABI



Semver **reloaded**

- API or ABI breaking change: major revision
- API or ABI non-breaking change: minor revision
- No change: patch revision



What about **dependencies**?

- Changing the major version of a public dependency will break API
 - ... and possibly ABI too
- Changing the major version of a private dependency will break ABI




Can I do **more**?

- Advertising change is important, but you can go the extra mile by providing migration scripts
- For example, Clang based refactoring tools
- This will encourage clients to upgrade quickly



What the future may **hold**

- Contracts TS should help you detect changes to your API more easily
- Modules TS may change the way you distribute binaries (possibly without headers)



Quizz

Did you follow everything?



Quizz #1

Before

```
void foo(int);
```

After

```
void foo(int, bool);
```

**Breaking API change
& breaking ABI change**





Quizz #2

Before

```
int foo(int);
```

After

```
int foo(long);
```

**API change
& breaking ABI change**





Quizz #3

Before

```
struct A {  
    int i;  
    char *s;  
};
```

After

```
struct A {  
    char *s;  
    int i;  
};
```

Breaking API change
& breaking ABI change





Quizz #4

Before

```
struct A {  
    void foo();  
    void bar();  
};
```



After

```
struct A {  
    void bar();  
    void foo();  
};
```

No change




Quizz #5

Before

```
int foo(int a, int b) {  
    return a + b;  
}
```

After

```
int foo(int a, int b) {  
    return a > b ? a : b;  
}
```



Invisible breaking API change





Quizz #6

Before

```
struct A {  
    virtual void foo();  
    virtual void bar();  
};
```

After

```
struct A {  
    virtual void bar();  
    virtual void foo();  
};
```

Breaking ABI change





Quizz #7

Before

```
struct A {  
    int i;  
    bool b;  
    char *s;  
};
```

After

```
struct A {  
    int i;  
    bool b;  
    char t[2];  
    char *s;  
};
```

ABI change(*)





Quizz #8

Before

```
void foo(int);
```

After

```
void foo(int, bool = false);
```

**API change
& breaking ABI change**





Quizz #9

Before

```
void foo(int);
```

After

```
void bar(int);
```

**Breaking API
& breaking ABI change**





Quizz #10

Before

```
struct A {  
    int i;  
    char *s;  
};
```

After

```
struct A {  
    int i;  
    char *str;  
};
```

Breaking API change





Quizz #10 and half

Before

```
namespace details {  
    int bar(int);  
}  
  
inline int foo(int x) {  
    return details::bar(x);  
}
```

After

```
namespace details {  
    int bazz(int);  
}  
  
inline int foo(int x) {  
    return details::bazz(x);  
}
```

Breaking ABI change



*No system became successful by
breaking backward
compatibility...*

*... especially without warning its
users beforehand*

“



Versioning

Communication between maintainers and users
about the changes in a software



Thanks!

Any **questions** ?

You can reach me at

✉ mro@puchiko.net

🐦 [@MatRopert](https://twitter.com/MatRopert)

🐙 [@mropert](https://github.com/mropert)