



C++ AND PERSISTENT MEMORY TECHNOLOGIES, LIKE INTEL'S 3D-XPOINT

Tomasz Kapela, @tomaszkapela

CppCon, Bellevue 2017



C++ AND PERSISTENT MEMORY ~~TECHNOLOGIES, LIKE INTEL'S 3D XPOINT~~

Tomasz Kapela, @tomaszkapela

CppCon, Bellevue 2017

**THIS IS NOT ABOUT 3D-XPOINT, IT'S ABOUT
C++**

THE MODEL

Programming Model:

At least four meanings...

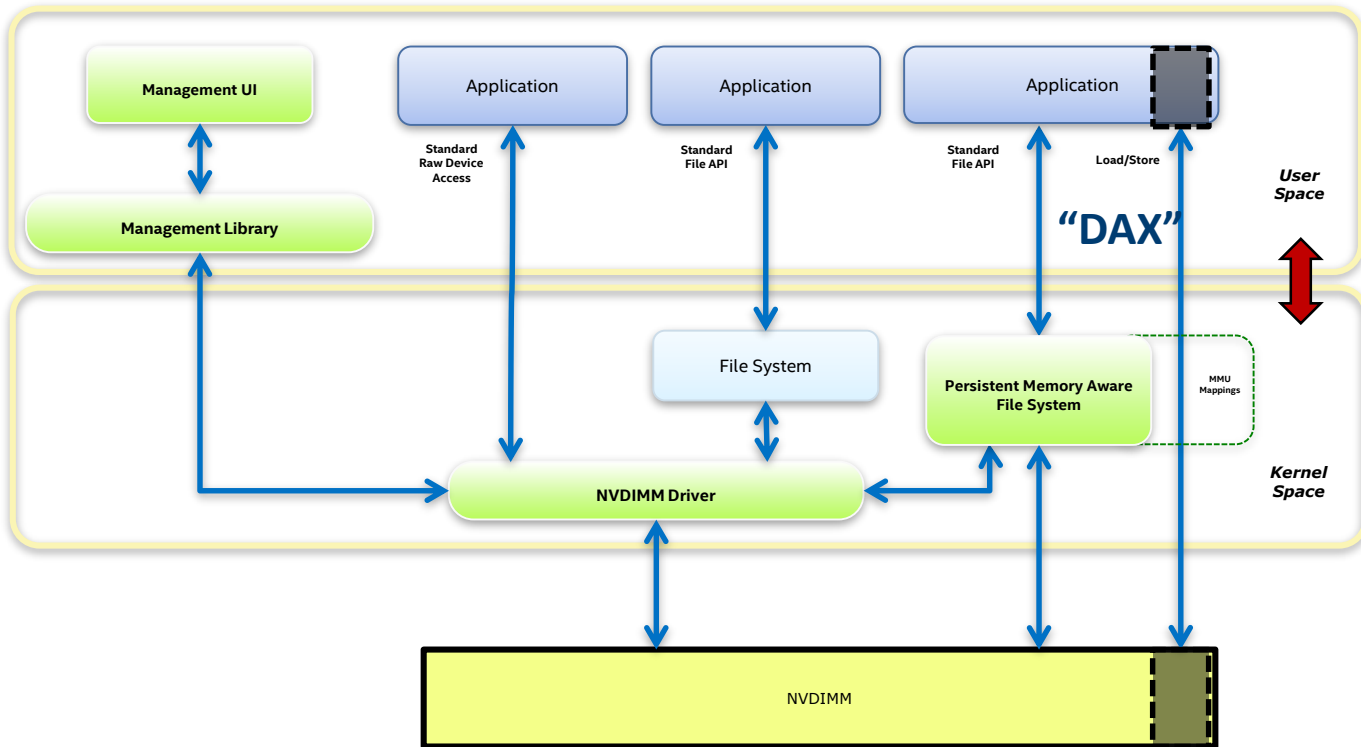
1. Interface between HW and SW
2. Instruction Set Architecture (ISA)
3. Exposed to Applications (by the OS)
4. The Programmer Experience

Programming Model:

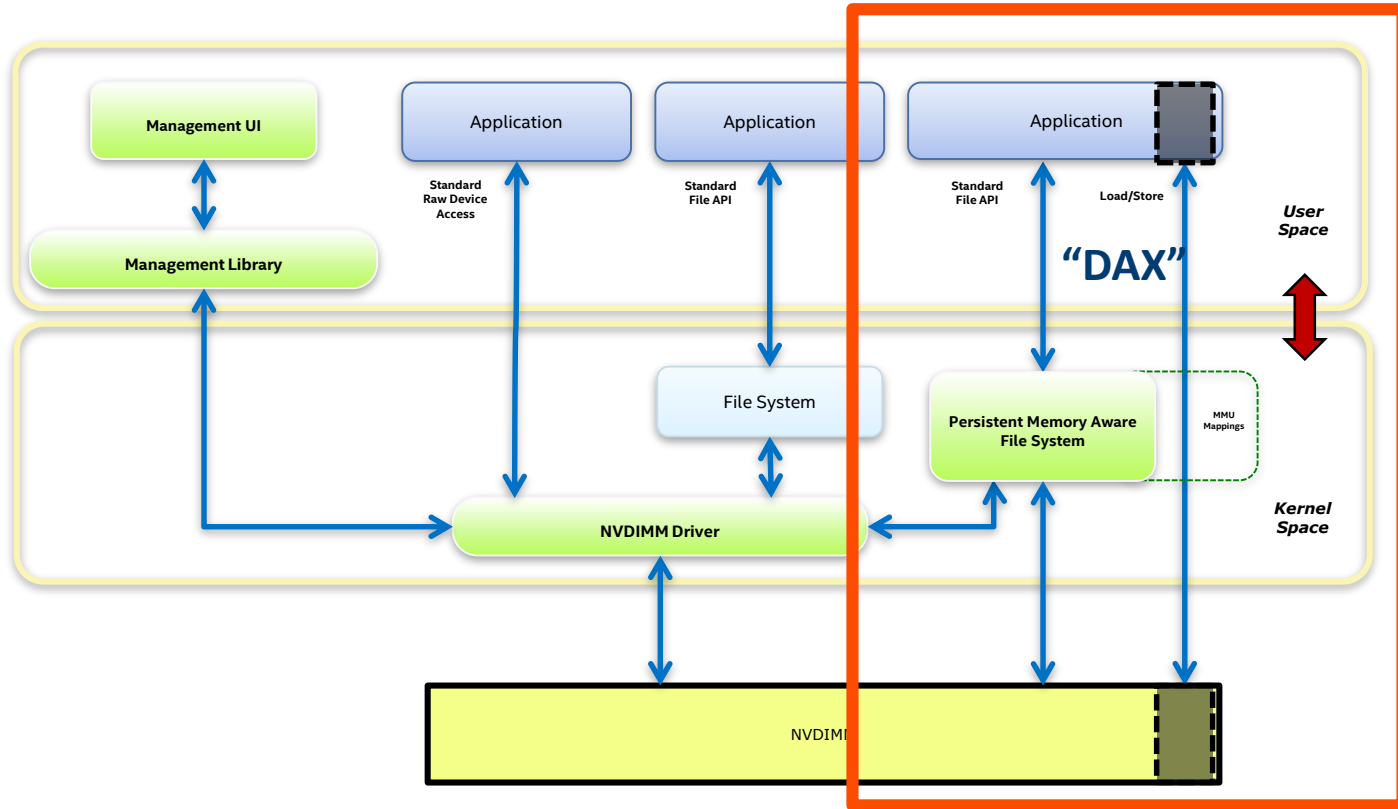
At least four meanings...

1. Interface between HW and SW
2. Instruction Set Architecture (ISA)
3. Exposed to Applications (by the OS)
4. The Programmer Experience

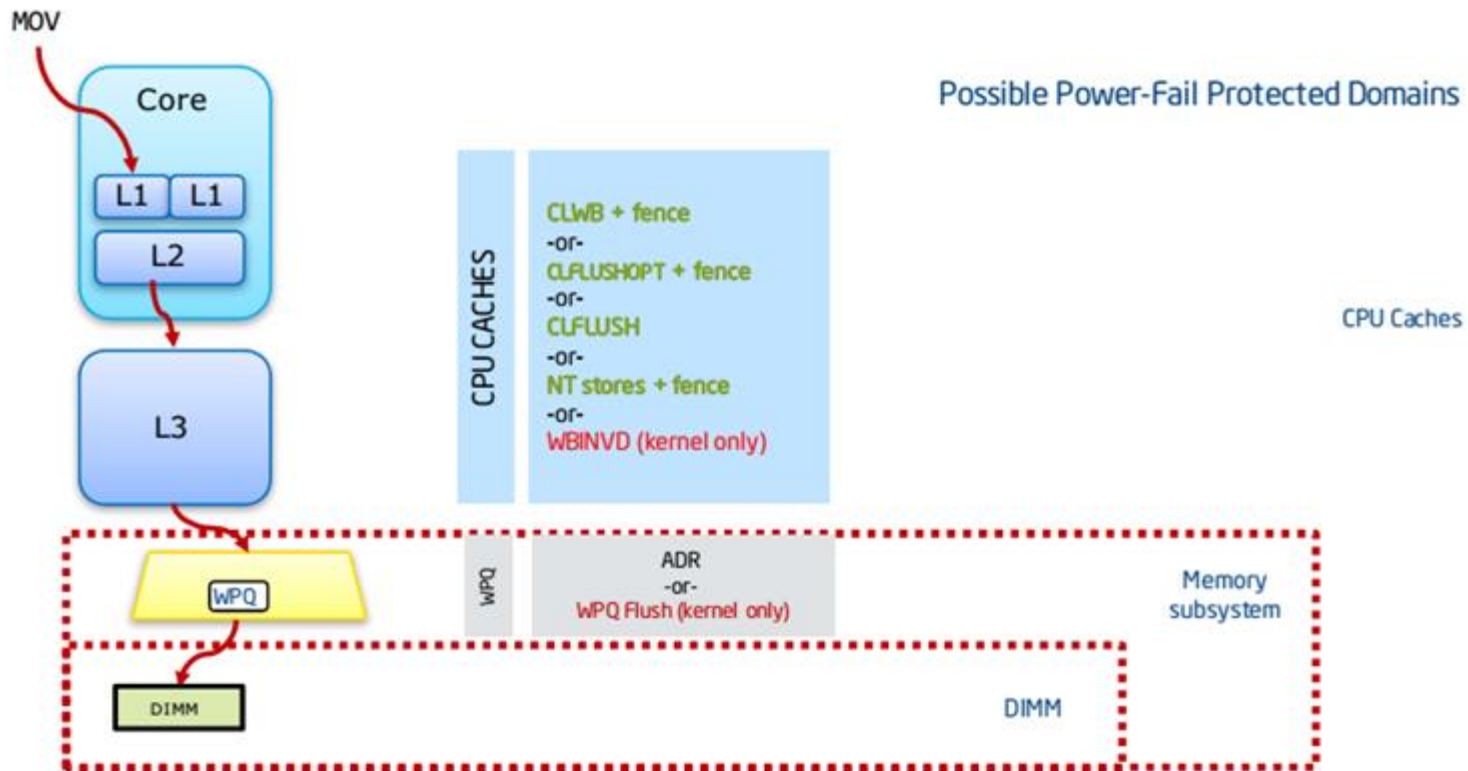
Programming Model (meaning 3): Exposing to Applications



Programming Model (meaning 3): Exposing to Applications




Data persistency



Atomicity

Flushing to Persistence

```
open(...);  
mmap(...);  
strcpy(pmem, „Hello, World!");  
pmem_persist(pmem, 14);
```




crash

Atomicity

Flushing to Persistence

```
open(...);  
mmap(...);  
strcpy(pmem, „Hello, World!");  
pmem_persist(pmem, 14);
```



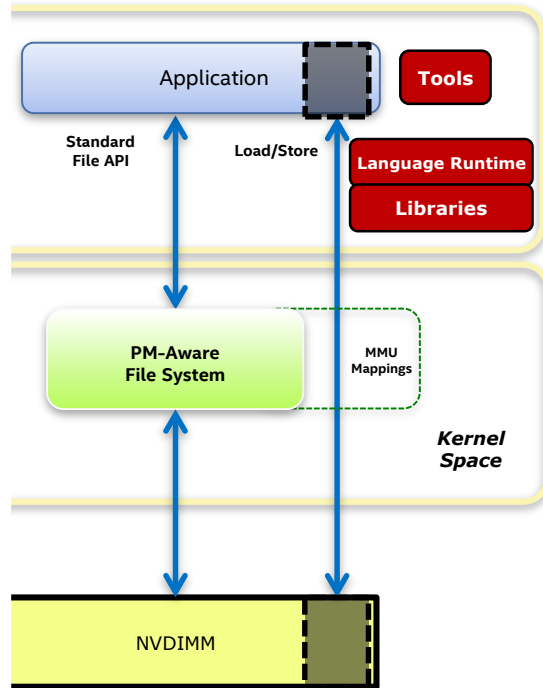
crash

Crossing the 8-Byte Store

Result?

1. „\0\0\0\0\0\0\0\0\0\0...”
2. „Hello, W\0\0\0\0\0\0...”
3. „\0\0\0\0\0\0\0\0\0orld!\0”
4. „Hello, \0\0\0\0\0\0\0\0”
5. „Hello, World!\0”

Programming Model (meaning 4): The Programmer Experience



Result:
Safer, less error-prone,
idiomatic in common languages

C++ std::vector

```
using pvector = vector<p<int>, pallocator>;  
pvector persistent_vector = allocate_from_persistent_memory();  
  
...  
persistent_vector.push_back(42);
```

C++ std::vector

```
std::vector<p<int>> persistent_vector = allocate_from_persistent_memory();
```

```
...
```

```
persistent_vector.push_back(42);
```

No flush calls.
Transactional.
C++ library handles it all.

C++ std::vector

```
std::vector<p<int>> persistent_vector = allocate_from_persistent_memory();  
...  
persistent_vector.push_back(42);
```

No flush calls.
Transactional.
C++ library handles it all.

See “pilot” project at: <https://github.com/pmem/libcxx>

GROUND RULES

Glossary

- Persistent Memory/Storage Class Memory/Non-Volatile Memory – fast, byte addressable memory which retains contents between power cycles
- Pool – a contiguous blob of persistent memory in a process' virtual address space

LET'S START WITH A QUEUE

Queue Entry

```
/* entry in the list */  
struct entry {  
    shared_ptr<entry> next;  
    int value;  
};  
  
// somewhere in the queue class  
  
shared_ptr<entry> head;  
shared_ptr<entry> tail;
```

Push

```
void push(int value)
{
    auto n = make_shared<entry>(value, nullptr);

    if (head == nullptr) {
        head = tail = n;
    } else {
        tail->next = n;
        tail = n;
    }
}
```

Pop

```
int pop()
{
    if (head == nullptr)
        throw runtime_error("Nothing to pop");

    auto ret = head->value;
    head = head->next;

    if (head == nullptr)
        tail = nullptr;

    return ret;
}
```

TRANSACTIONS

Push

```
void push(int value)
{
    transaction::exec_tx(pool, [this, &value] {
        auto n = make_shared<entry>(value, nullptr);

        if (head == nullptr) {
            head = tail = n;
        } else {
            tail->next = n;
            tail = n;
        }
    });
}
```

Scope of the
transaction

Pop

```
int pop()
{
    transaction::exec_tx(pool, [this] {
        if (head == nullptr)
            throw runtime_error("Nothing to pop");

        auto ret = head->value;
        head = head->next;
        if (head == nullptr)
            tail = nullptr;

        return ret;
    });
}
```

Scope of the
transaction

Transactions

- Undo log based transactions
- ACID like properties
- Can be nested
- In case of interruption it is rolled-back or completed upon next pool open
- Locks are held until the end of a transaction

Manual Transaction Example

```
auto pop = pool<root>::open("/path/to/poolfile", "layout string");  
{  
    transaction::manual(pop, persistent_mtx, persistent_shmtx);  
  
    // do some work...  
  
    transaction::commit();  
}
```

```
auto aborted = transaction::get_last_tx_error();
```

Manual Transaction Example

Pool handle

```
auto pop = pool<root>::open("/path/to/poolfile", "layout string");  
{  
    transaction::manual(pop, persistent_mtx, persistent_shmtx);  
  
    // do some work...  
  
    transaction::commit();  
}
```

```
auto aborted = transaction::get_last_tx_error();
```

Manual Transaction Example

Root object type

```
auto pop = pool<root>::open("/path/to/poolfile", "layout string");  
{  
    transaction::manual(pop, persistent_mtx, persistent_shmtx);  
  
    // do some work...  
  
    transaction::commit();  
}
```

```
auto aborted = transaction::get_last_tx_error();
```

Manual Transaction Example

Custom layout
string

```
auto pop = pool<root>::open("/path/to/poolfile", "layout string");  
{  
    transaction::manual(pop, persistent_mtx, persistent_shmtx);  
  
    // do some work...  
  
    transaction::commit();  
}
```

```
auto aborted = transaction::get_last_tx_error();
```

Manual Transaction Example

```
auto pop = pool<root>::open("/path/to/poolfile", "layout string");  
{  
    transaction::manual(pop, persistent_mtx, persistent_shmtx);  
  
    // do some work...  
  
    transaction::commit();  
}
```

Synchronization
point

```
auto aborted = transaction::get_last_tx_error();
```

Manual Transaction Example

```
auto pop = pool<root>::open("/path/to/poolfile", "layout string");  
{  
    transaction::manual(pop, persistent_mtx, persistent_shmtx);  
  
    // do some work...  
  
    transaction::commit();  
}
```

Has to be
explicit

```
auto aborted = transaction::get_last_tx_error();
```

Manual Transaction Example

```
auto pop = pool<root>::open("/path/to/poolfile", "layout string");  
{  
    transaction::manual(pop, persistent_mtx, persistent_shmtx);  
  
    // do some work...  
  
    transaction::commit();  
}
```



Locks drop here

```
auto aborted = transaction::get_last_tx_error();
```


Manual Transaction Example

```
auto pop = pool<root>::open("/path/to/poolfile", "layout string");  
{  
    transaction::manual(pop, persistent_mtx, persistent_shmtx);  
  
    // do some work...  
  
    transaction::commit();  
}
```

Well... This is embarrassing

```
auto aborted = transaction::get_last_tx_error();
```

Manual Transactions

- Based on the familiar RAII concept
- Fairly easy to use
 - Explicit transaction commit because of `std::uncaught_exception`
 - Does not throw an exception on transaction abort
- By default aborts to account for third-party exceptions or amnesia
- Accepts an arbitrary number of (persistent memory resident) locks

Automatic Transaction Example

```
auto pop = pool<root>::open("/path/to/poolfile", "layout string");  
try {  
    transaction::automatic(pop, persistent_mtx, persistent_shmtx);  
    // do some work...  
} catch (...) {  
    // do something meaningful  
}  
  
auto aborted = transaction::get_last_tx_error();
```

Automatic Transaction Example

```
auto pop = pool<root>::open("/path/to/poolfile", "layout string");  
try {  
    transaction tx(pop, persistent_mtx, persistent_shmtx);  
    // do some work...  
} catch (...) {  
    // do something meaningful  
}  
  
auto aborted = transaction::get_last_tx_error();
```

No more
commit

Automatic Transaction Example

```
auto pop = pool<root>::open("/path/to/poolfile", "layout string");  
try {  
    transaction::automatic(pop, persistent_mtx, persistent_shmtx);  
    // Locks drop here ...  
} catch (...) {  
    // do something meaningful  
}  
  
auto aborted = transaction::get_last_tx_error();
```

Automatic Transaction Example

```
auto pop = pool<root>::open("/path/to/poolfile", "layout string");  
try {  
    transaction::automatic(pop, persistent_mtx, persistent_shmtx);  
    // Locks drop here ...  
} catch (...) {  
    // do something meaningful and mind your locks!  
}  
  
auto aborted = transaction::get_last_tx_error();
```

Automatic Transaction Example

```
auto pop = pool<root>::open("/path/to/poolfile", "layout string");  
try {  
    transaction::automatic(pop, persistent_mtx, persistent_shmtx);  
    // do some work...  
} catch (...) {  
    // do something meaningful and mind your locks!  
}  
  
auto aborted = transaction::get_last_tx_error();
```

Still
embarrassing?

Automatic Transactions

- Functionally and semantically identical to the manual transaction
- No explicit transaction commit
- Need C++17
 - Relies on `std::uncaught_exception`s

Closure Transaction Example

```
auto pop = pool<root>::open("/path/to/poolfile", "layout string");
```

```
transaction::exec_tx(pop, [] {  
    // do some work...  
}, persistent_mtx, persistent_shmtx);
```

```
auto aborted = transaction::get_last_tx_error();
```

Closure Transaction Example

```
auto pop = pool<root>::open("/path/to/poolfile", "layout string");
```

```
transaction::exec_tx(pop, [] {  
    // do some work...  
}, persistent_mtx, persistent_shmtx);
```

Awkward

```
auto aborted = transaction::get_last_tx_error();
```

Closure Transaction Example

```
auto pop = pool<root>::open("/path/to/poolfile", "layout string");
```

```
transaction::exec_tx(pop, [] {  
    // do some work...  
}, persistent_mtx, persistent_shmtx);
```

Not embarrassing
anymore

```
auto aborted = transaction::get_last_tx_error();
```

Closure Transactions

- Take an `std::function` object as transaction body
- No explicit transaction commit
- Available with every C++11 compliant compiler
- Throw an exception when the transaction is aborted
- Take an arbitrary number of locks
 - Unfortunately at the very end



Resides in persistent memory

Queue Entry Recap

```
/* entry in the list */  
struct entry {  
    shared_ptr<entry> next;  
    int value;  
};
```

Queue Entry Recap

```
/* entry in the list */  
struct entry {  
    shared_ptr<entry> next;  
    int value;  
};
```



Allocated, modified and
deleted within a transaction

Queue Entry Recap

```
/* entry in the list */  
struct entry {  
    shared_ptr<entry> next;  
    int value;  
};
```

Needs to be
snapshot

Allocated, modified and
deleted within a transaction

Partially Persistent Queue Entry

```
/* entry in the list */  
struct entry {  
    shared_ptr<entry> next;  
    p<int> value;  
};
```

Does the actual
snapshot

Allocated, modified and
deleted within a transaction

The p<> Property

- AKA the workhorse
- Overloads `operator=` for snapshotting in a transaction
- Overloads a bunch of other operators for seamless integration
 - Arithmetic
 - Logical
- Should be used for fundamental types
 - No convenient way to access members of aggregate types
 - No `operator.` to overload

PERSISTENT POINTER

Partially Persistent Queue Entry Recap

```
/* entry in the list */  
struct entry {  
    shared_ptr<entry> next;  
    p<int> value;  
};
```



Allocated, modified and
deleted within a transaction

Partially Persistent Queue Entry Recap

/* entry in the list

struct entry {

shared_ptr<entry> next;

p<int> value;

};

Needs to be
snapshot

Allocated, modified and
deleted within a transaction

Fully Persistent Queue Entry

```
/* entry in the list */  
struct entry {  
    persistent_ptr<entry> next;  
    p<int> value;  
};
```

} Allocated, modified and
deleted within a transaction

Fully Persistent Queue Entry

```
/* entry */  
struct  
    persistent_ptr<entry> next;  
    p<int> value;  
};
```

Now will be
snapshot

Allocated, modified and
deleted within a transaction

Persistent Smart Pointer

- Points to objects within a persistent memory pool
 - Manages translating persistent addresses to runtime addresses
- Is a random access iterator
- Has primitives for flushing contents to persistence
- `std::allocator` friendly

Persistent Smart Pointer

- Does not manage object lifetime
- Does not automatically add contents to the transaction
 - But it does add itself to the transaction
- Does not point to polymorphic objects
 - No good way to rebuild runtime state after pool reboot

Persistent Queue Entry

```
/* entry in the list */  
struct entry {  
    persistent_ptr<entry> next;  
    p<int> value;  
};
```

```
// somewhere in the queue class
```

```
persistent_ptr<entry> head;  
persistent_ptr<entry> tail;
```

Push Recap

```
void push(int value)
{
    transaction::exec_tx(pool, [this] {
        auto n = make_shared<entry>(value, nullptr);

        if (head == nullptr) {
            head = tail = n;
        } else {
            tail->next = n;
            tail = n;
        }
    });
}
```

Persistent Push

```
void push(int value)
{
    transaction::exec_tx(pool, [this] {
        auto n = make_persistent<entry>(value, nullptr);

        if (head == nullptr) {
            head = tail = n;
        } else {
            tail->next = n;
            tail = n;
        }
    });
}
```

Persistent Push

```
void push(int value)
{
    transaction::exec_tx(
        auto n = make_persistent<entry>(value, nullptr);

        if (head == nullptr) {
            head = tail = n;
        } else {
            tail->next = n;
            tail = n;
        }
    });
}
```

Allocation now part
of the transaction

Persistent Push

The pool handle for completeness

```
void push(pool_base &pool, int value)
{
    transaction::exec_tx(pool, [this] {
        auto n = make_persistent<entry>(value, nullptr);

        if (head == nullptr) {
            head = tail = n;
        } else {
            tail->next = n;
            tail = n;
        }
    });
}
```

Pop Recap

```
int pop() {  
    transaction::exec_tx(pool, [this] {  
        if (head == nullptr)  
            throw runtime_error("Nothing to pop");  
  
        auto ret = head->value;  
        head = head->next;  
        if (head == nullptr)  
            tail = nullptr;  
  
        return ret;  
    });  
}
```

Persistent Pop

```
int pop() {  
    transaction::exec_tx(pool, [this] {  
        auto ret = head->value;  
        auto tmp_entry = head->next;  
  
        delete_persistent<entry>(head);  
  
        head = tmp_entry;  
  
        if (head == nullptr)  
            tail = nullptr;  
        return ret;  
    });  
}
```

Code removed to
improve readability

Persistent Pop

```
int pop() {  
    transactional_op [this] {  
        auto ret = head->value;  
        auto tmp_entry = head->next;  
  
        delete_persistent<entry>(head);  
  
        head = tmp_entry;  
  
        if (head == nullptr)  
            tail = nullptr;  
        return ret;  
    });  
}
```

Temporary
entry pointer

Persistent Pop

```
int pop() {  
    transaction::exec_tx(pool, [this] {  
        auto  
        auto  
        delete_persistent<entry>(head);  
  
        head = tmp_entry;  
  
        if (head == nullptr)  
            tail = nullptr;  
        return ret;  
    });  
}
```

No object lifetime
management

Persistent Pop

```
int pop() {  
    transaction::exec_tx(pool, [this] {  
        auto ret = head->value;  
        auto tmp_entry = head->next;  
  
        Update head tent<entry>(head);  
        head = tmp_entry;  
  
        if (head == nullptr)  
            tail = nullptr;  
        return ret;  
    });  
}
```

Persistent Pop

Complete example

```
int pop(pool_base &pool) {  
    transaction::exec_tx(pool, [this] {  
        if (head == nullptr)  
            throw runtime_error("Nothing to pop");  
  
        auto ret = head->value;  
        auto tmp_entry = head->next;  
        delete_persistent<entry>(head);  
        head = tmp_entry;  
        if (head == nullptr)  
            tail = nullptr;  
        return ret;  
    });  
}
```

Transactional Allocation Functions

- Can be used only within transactions
 - Use transaction logic to enable allocation/delete rollback of persistent state
- `make_persistent` calls appropriate constructor
 - Syntax similar to `std::make_shared`
- `delete_persistent` calls the destructor
 - Not similar to anything found in `std`

THREAD SYNCHRONIZATION

Persistent memory resident

Persistent Memory Synchronization

- Types:
 - mutex
 - shared_mutex
 - timed_mutex
 - condition_variable
- All with an interface similar to their `std` counterparts
- Auto reinitializing
- Can be used with transactions

ALLOCATOR

Details

- Standard compliant implemetation
 - allocate/deallocate
 - construct/destroy
 - max_size
 - rebind
- Uses persistent_ptr as its pointer type (really important)
- Transactional only
- Not much left to say...

CONTAINERS ENABLING

Why?

- The pointer type is implemented
- The allocator is implemented
- The std containers are already there
 - Almost immediately usable
- Widely used with a familiar interface
- Maintained outside of NVML

std::vector

```
typedef typename __alloc_traits::pointer      pointer;  
  
pointer      __begin_;  
pointer      __end_;  
pointer      __end_cap_;
```

std::map (AKA RB Tree)

```
template <class _VoidPtr>
class __tree_node_base
{
    // many other things
    pointer      __right_;
    __parent_pointer __parent_;
    bool __is_black_;
```

std::map (AKA RB Tree)

```
template <class _VoidPtr>
class __tree_node_base
{
    // many other things
    pointer      __right_;
    __parent_pointer __parent_;
    bool __is_black_;
```

Vital node
metadata

std::map (AKA RB Tree)

```
template <class _VoidPtr>
class __tree_node_base
{
    // many other things
    pointer      __right_;
    __parent_pointer __parent_;
    bool __is_black_;
```

Will not get
snapshot

std::map (AKA RB Tree)

```
template <class _VoidPtr>  
class __tree_node_base  
{
```

Where's the
allocator?

```
    // many other things  
    pointer      __right_;  
    __parent_pointer __parent_;  
    bool __is_black_;
```


<memory>

```
template <class _Ptr>
struct pointer_traits
{
    typedef _Ptr                pointer;

    // many other things

    typedef typename __pointer_traits_persistency_type<element_type,
        pointer>::type          persistency_type;
```

Injects p<>

std::map revisited

```
template <class _VoidPtr>
class __tree_node_base
{

typedef typename __rebind_persistency_type<pointer, bool>::type bool_type;
// many other things
bool_type __is_black_;
```

std::map revisited

```
template <class _VoidPtr>
class __tree_node_base
{
```

```
typedef typename __rebind_persistency_type<pointer, bool>::type bool_type;
// many other things
bool_type __is_black_;
```




p<bool>

KNOWN ISSUES

Issues

- Static data members
 - The DATA/BSS section is not in persistent memory
- Standard library layout versioning
- Reliable vtable data rebuild
- No `std::string` equivalent - yet
- No persistent references – core language change



Handy links

Persistent memory programming homepage:

<http://pmem.io>

Intel developer zone:

<https://software.intel.com/en-us/persistent-memory>

Libpmemobj++ documentation:

http://pmem.io/nvml/cpp_obj/

Pmem Google Groups page:

<https://groups.google.com/forum/#!forum/pmem>