

From security to performance to GPU programming: Exploring modern allocators

Sergey Zubkov

Morgan Stanley

What do they do?

```
template<class T, class Allocator> class vector;
```

1. Allocation policy
2. Deallocation policy
3. Construction policy
4. Destruction policy
5. Addressing* policy (“fancy” or “synthetic” pointers)

*must have mapping to heap memory to use standard containers

How are they used?

1. STL's static allocators (allocator is part of the type)

```
std::vector<T, Allocator> v;
```

2. C++17's dynamic allocators (supplied at run-time)

```
std::pmr::vector<T> v(&pool);
```

Enough theory, let's see what this can do

Just a few of the many off-the-shelf allocators available now:

- (various libraries) `secure_allocator`
- `std::pmr::polymorphic_allocator`
- `tbb::scalable_allocator`

- `boost::interprocess::allocator` (fancy pointers inside!)
 - (with a shout-out to `std::scoped_allocator_adaptor`)
- `boost::compute::pinned_allocator` (even fancier pointers inside!)

secure_allocator

...as reinvented in Bitcoin, Botan, MongoDB, and even JsonCpp

- On allocation:
 - locks memory so that it cannot be swapped out (mlock+madvise/VirtualLock)
- On deallocation:
 - wipes memory before freeing (OPENSSL_cleanse, RtlSecureZeroMemory, etc)

secure_allocator

```
{  
    std::string pwd = "correct horse battery staple";  
    ...  
}  
p = ...; // hackers got this address  
for (int n = 0; n < 28; ++n)  
    if (isprint(p[n])) std::cout << p[n]; else std::cout << "❖";
```

heap contents at 0x2534c20:

❖❖❖❖❖❖❖horse battery staple

secure_allocator

```
#include "support/allocators/secure.h"
{
    SecureString pwd = "correct horse battery staple";
    ...
}
p = ...; // hackers got this address
for (int n = 0; n < 28; ++n)
    if (isprint(p[n])) std::cout << p[n]; else std::cout << "❖";
```

heap contents at 0x7f7ac0236fe0:

secure_allocator

```
#include "support/allocators/secure.h"
{
    SecureString pwd = "correct horse battery staple";
    ...
}
p = .
for (
    if
        template <typename T>
        struct secure_allocator /* we're not going there! */;

        typedef std::basic_string<char,
                                std::char_traits<char>,
                                secure_allocator<char> > SecureString;
```

heap contents at 0x7f7ac0236fe0:

?? ??

Based on a true story...

```
struct Event { std::vector<int> data = std::vector<int>(512); };
std::list<std::shared_ptr<Event>> q;
void producer() {
    for (int n = 0; n != /* LOTS */; ++n) {
        std::lock_guard<std::mutex> lk(m);
        q.push_back(std::make_shared<Event>());
        cv.notify_one();
    }
}
void consumer() { /* might resize the event... */ }
```

What happens after running this for one year?

Based on a true story...

After this demo completes:

- 50k events x 16 produces x 4 consumers
 - Total non-mmapped bytes (arena): 1,625,690,112
 - Total allocated space (uordblks): 119,632
 - Total free space (fordblks): 1,625,570,480
- 100k events x 16 producers x 4 consumers
 - Total non-mmapped bytes (arena): 3,310,215,168
 - Total allocated space (uordblks): 121,552
 - Total free space (fordblks): 3,310,093,616

std::pmr::polymorphic_allocator

C++17 to the rescue!

```
std::pmr::synchronized_pool_resource pool;
struct Event {
    std::pmr::vector<int> data = std::pmr::vector<int>(512, &pool);
};
std::pmr::list<std::shared_ptr<Event>> q(&pool);
void producer() {
    for (int n = 0; n != /* LOTS */; ++n) {
        std::lock_guard<std::mutex> lk(m);
        q.push_back(std::allocate_shared<Event,
                                std::pmr::polymorphic_allocator<Event>>(&pool));
        cv.notify_one();
    }
}
```

std::pmr::polymorphic_allocator

After this demo completes:

- 50k events x 16 produces x 4 consumers
 - Total non-mmapped bytes (arena): 2,908,160
 - Total allocated space (uordblks): 118,832
 - Total free space (fordblks): 2,789,328
- 100k events x 16 producers x 4 consumers
 - Total non-mmapped bytes (arena): 2,908,160 (1000x!)
 - Total allocated space (uordblks): 118,832
 - Total free space (fordblks): 2,789,328

tbb::scalable_allocator


Size-segregated thread-private heaps, and many optimizations[*]

- Alexey Kukanov and Michael J. Voss “The Foundations for Scalable Multi-core Software in Intel® Threading Building Blocks”, Intel Technology Journal, Volume 11, Issue 4, 2007

```
struct Event { std::vector<int, tbb::scalable_allocator<T>> data{512}; };
std::list<std::shared_ptr<Event>, tbb::scalable_allocator<T>> q;
void producer() {
    for (int n = 0; n != /* LOTS */; ++n) {
        std::lock_guard<std::mutex> lk(m);
        q.push_back(std::allocate_shared<Event,
                                tbb::scalable_allocator<Event>>({}));
        cv.notify_one();
    }
}
```

tbb::scalable_allocator

Time to transmit 160k events in the previous demo vs number of threads



boost::interprocess::allocator

How do you share a C++ container between processes?

```
namespace ipc = boost::interprocess;
ipc::managed_shared_memory segment{ /* ... */ };

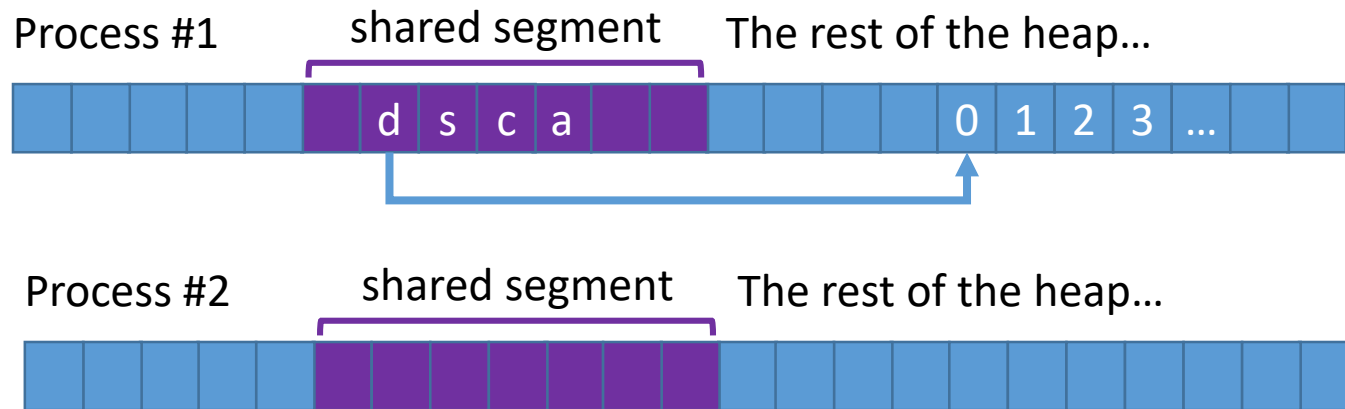
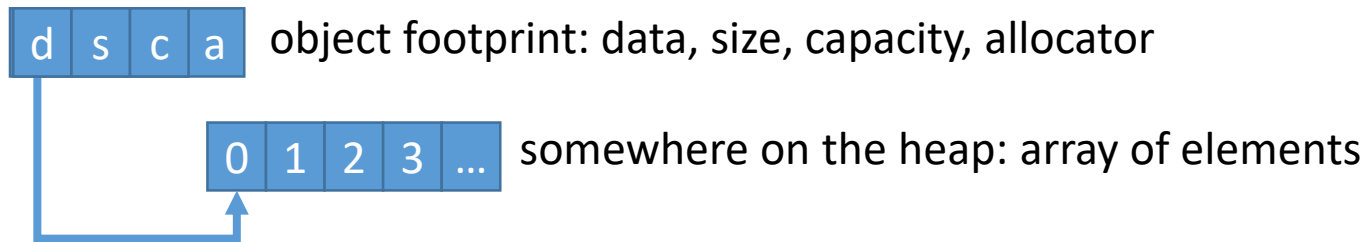
using vec_t = std::vector<int>;
vec_t* v = segment.construct<vec_t>("vec")(in.begin(), in.end());
```

```
... in another process ...
vec_t* m = segment.find<vec_t>("vec").first;
for (int& n : *m) std::cout << n << ' '; // SEGFAULT!
```

boost::interprocess::allocator

How do you share a C++ container between processes?


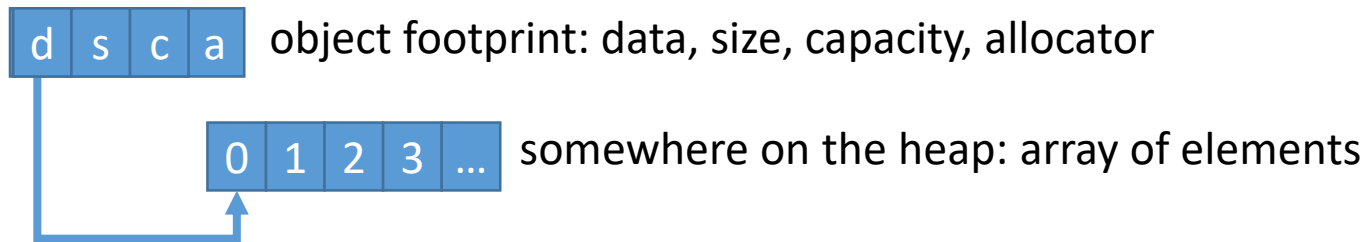
What's in a `vector<int>`?



boost::interprocess::allocator

How do you share a C++ container between processes?



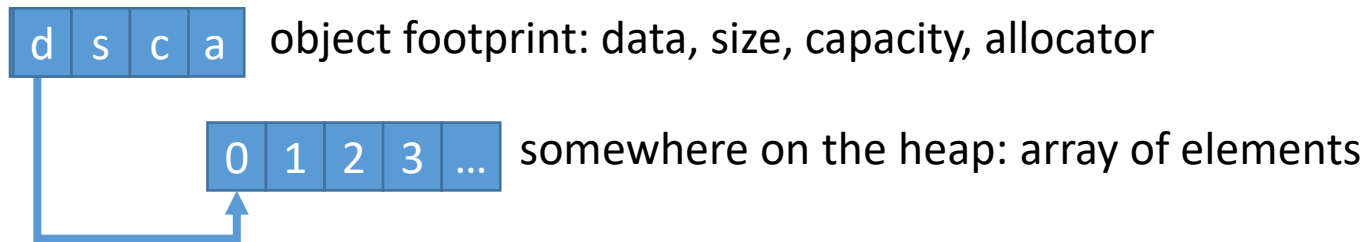
What's in a `vector<int>`?



boost::interprocess::allocator

How do you share a C++ container between processes?

What's in a `vector<int>`?




boost::interprocess::allocator

How do you share a C++ container between processes?

```
using alloc_t = ipc::allocator<int, ipc::managed_shared_memory::segment_manager>;  
alloc_t a{ segment.get_segment_manager() };
```

```
using vec_t = std::vector<int, alloc_t>;  
vec_t* v = segment.construct<vec_t>("vec")(in.begin(), in.end(), a);
```




```
... in another process ...  
vec_t* m = segment.find<vec_t>("vec").first;  
for (int& n : *m) std::cout << n << ' '; // WORKS
```

boost::interprocess::allocator

What if there is a container inside that container?

```
using mat_t = std::vector<std::vector<int>, alloc_t<std::vector<int>>>>;  
mat_t& m = *segment.construct<mat_t>("matrix")(in.begin(), in.end(), a);
```



```
... in another process ...  
using mat_t = std::vector<std::vector<int>, alloc_t<std::vector<int>>>>;  
mat_t* m = segment.find<mat_t>("matrix").first;  
for (auto& r : *m)  
    for (int n : r)  
        std::cout << n << ' '; // SEGFAULT!
```

boost::interprocess::allocator

What if there is a container inside that container?

```
using vec_t = std::vector<int, alloc_t<int>>>;
using mat_t = std::vector<vec_t, std::scoped_allocator_adaptor<alloc_t<vec_t>>>>;
mat_t& m = *segment.construct<mat_t>("matrix")(3, a);

m[0].resize(3); // allocates the row inside the same segment
m[2].assign({ 6, 7, 8 });
```

```
... in another process ...

using vec_t = std::vector<int, alloc_t<int>>>;
using mat_t = std::vector<vec_t, std::scoped_allocator_adaptor<alloc_t<vec_t>>>>;
mat_t* m = segment.find<mat_t>("matrix").first;
for (auto& r : *m)
    for (int n : r)
        std::cout << n << ' '; // WORKS
```

Fancy pointer



Our vector's data pointer is not T^* !

It is `boost::interprocess::offset_ptr<T>`

https://github.com/boostorg/interprocess/blob/develop/include/boost/interprocess/offset_ptr.hpp#L207

```
//!A smart pointer that stores the offset between between the pointer and the  
//!the object it points. This allows offset allows special properties, since  
//!the pointer is independent from the address of the pointee, if the  
//!pointer and the pointee are still separated by the same offset. This feature  
//!converts offset_ptr in a smart pointer that can be placed in shared memory and  
//!memory mapped files mapped in different addresses in every process.
```

See C++ Now 2017: Bob Steagall “Testing the Limits of Allocator Awareness”

<https://www.youtube.com/watch?v=fmJfKm9ano8>

https://github.com/boostcon/cppnow_presentations_2017/tree/master/05-18-2017_thursday

boost::compute::pinned_allocator

- On allocation:
 - clCreateBuffer(CL_MEM_ALLOC_HOST_PTR), clRetainMemObject
- On deallocation:
 - clReleaseMemObject
- Fancy pointer
 - device_ptr<T> (holds a reference to a device buffer and an index)

```
namespace bc = boost::compute;  
bc::device d = bc::system::default_device();  
bc::command_queue queue(bc::context{ d }, d);
```

```
bc::vector<float, bc::pinned_allocator<float>> v({ 1.f, 2.f, 3.f, 4.f }, queue);  
bc::transform(v.begin(), v.end(), v.begin(), bc::sqrt<float>(), queue);  
for (float f : v) std::cout << f << ' ';
```

and many more...

- Allocators are reusable, packaged, policies that you can pick and choose.
- Hundreds of them are in use, in production.

cache_aligned_allocator
node_allocator
private_node_allocator
cached_node_allocator
adaptive_pool
private_adaptive_pool
cached_adaptive_pool
pool_allocator
fast_pool_allocator
...

any_interprocess_allocator
stack_allocator
__mt_alloc
__pool_alloc
_ExtPtr_allocator
array_allocator
bitmap_allocator
debug_allocator
throw_allocator
...

mmap_allocator
HugeAllocator
LowAllocator
SpecialAllocator
StackAllocator
aligned_allocator_cpp11
gc_allocator
cacheline_allocator
pyr_pool_compile_allocator
...

tlsf_allocator
aligned_allocator
tracker_alloc
traceable_allocator
libc_allocator_with_realloc
throw_allocator_random
casacore_allocator
mallocator
recycling_allocator
...

Allocator track on CppCon 2017

- 9/25 3:15 pm Alisdair Meredith “An allocator model for std2”
- 9/25 4:45 pm <you are here>
- 9/26 2:00 pm John Lakos “Local (‘Arena’) Memory Allocators”
- 9/28 2:00 pm Bob Steagall “How to Write a Custom Allocator”
- 9/29 9:00 am Pablo Halpern “Modern Allocators: The Good Parts”
- 9/29 1:30 pm Marshall Clow “Customizing the Standard Containers”