

# Compile Once Debug Twice

Samy Al Bahra / @0xF390

# Samy Al Bahra

**@0xF390** or **sbahra@backtrace.io**

**Co-founder of Backtrace.** Building a modern debugging platform for natively compiled software.

**AppNexus, Message Systems, GWU HPCL.**

Scalability, reliability, productivity, multicore, heterogeneous architectures and PGAS.

**Concurrency Kit.** Advanced synchronization primitives for the research, design and implementation of high-performance concurrent systems.

# Introduction

```
(lldb) frame variable  
(void *) unused = <variable not available>  
  
(int) n_state = <variable not available>  
  
(uint64_t) a = 0  
(uint64_t) j = <variable not available>
```

# Introduction

Most people select their compiler based on:

- Availability
- Correctness
- Compilation Speed
- Code Generation Quality

# Introduction

Most people select their compiler based on:

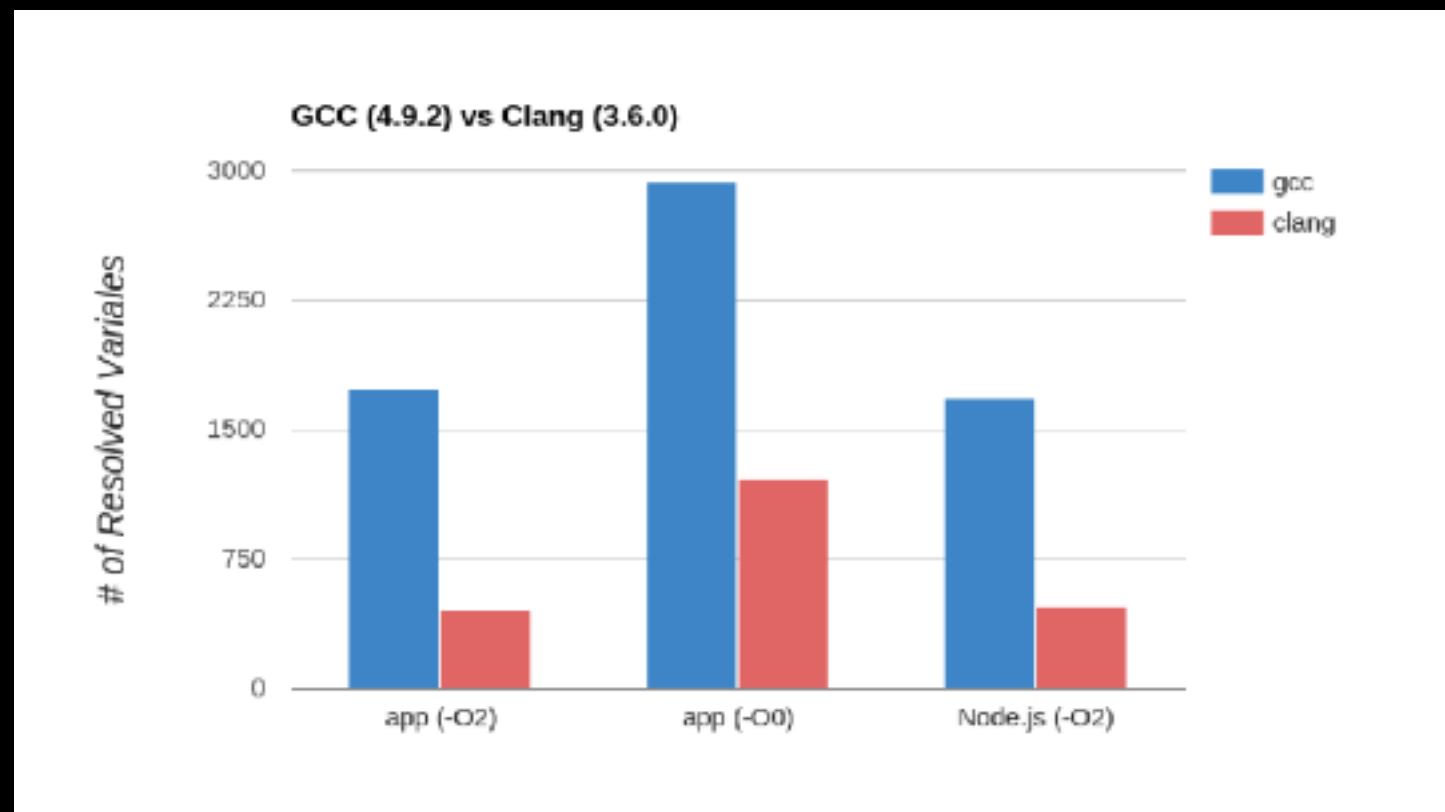
- Availability
- Correctness
- Compilation Speed
- Code Generation Quality

A facet that is commonly ignored is:

**Debug Information Quality.**

# Introduction


Debug information quality **is not** proportionally related to the quality of the generated executable code and wildly varies from compiler to compiler.



# Introduction

Being able to symbolically debug your software effectively directly impacts your ability to resolve many classes of bugs quickly.

Resolution time can be the difference between success and failure.



– “We crash a couple hundred times a day”

# Agenda

- An Introduction to Symbolic Debugging
- Optimizations and Debug Information
- A Comparison



# Symbolic Debugging

The wc program tells you the number of 'w' characters in a file, terribly.

```
static unsigned int
count(char *buffer, size_t n)
{
    unsigned int sum = 0;
    size_t i;

    for (i = 0; i < n; i++)
        sum += buffer[i] == 'w';

    return sum;
}
```

# Symbolic Debugging

The program is compiled and linked into an executable file in a format the underlying operating system understands.

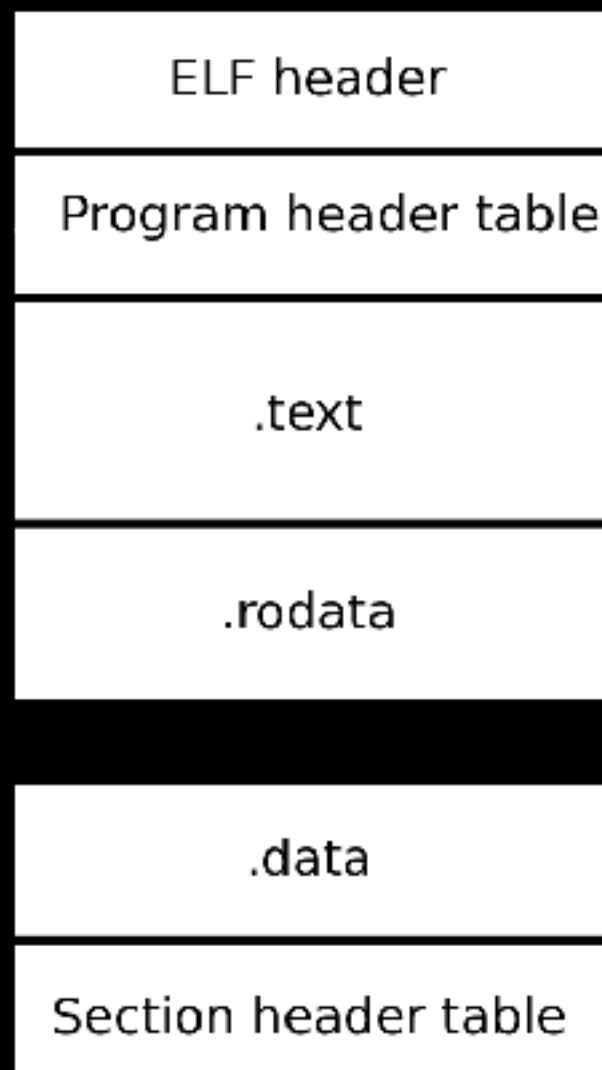
```
$ gcc -o wc wc.c -O2 -ggdb  
$ ./wc /etc/passwd  
16
```

The executable format consists of various sections that tell the kernel **how** the executable should be loaded as well as **executable code**.

# Symbolic Debug

Mention Mach-O and PE.

## ELF for Executable and Linking Format



```
sbahra@broadwell$ readelf -l ./wc
```

Elf file type is EXEC (Executable file)  
Entry point 0x400800  
There are 9 program headers, starting at offset 64

Program Headers:

Type	Offset FileSiz	VirtAddr MemSiz	PhysAddr Flags Align
PHDR	0x0000000000000040	0x0000000000400040	0x0000000000400040
	0x00000000000001f8	0x00000000000001f8	R E 8
INTERP	0x0000000000000238	0x0000000000400238	0x0000000000400238
	0x000000000000001c	0x000000000000001c	R 1
[Requesting program interpreter: /lib64/ld-linux-x86-64.so.2]			
LOAD	0x0000000000000000	0x0000000000400000	0x0000000000400000
	0x00000000000000b24	0x00000000000000b24	R E 200000
LOAD	0x00000000000000e08	0x0000000000600e08	0x0000000000600e08
	0x00000000000000270	0x00000000000000288	RW 200000
DYNAMIC	0x00000000000000e28	0x0000000000600e28	0x0000000000600e28
	0x000000000000001d0	0x000000000000001d0	RW 8
NOTE	0x00000000000000254	0x0000000000400254	0x0000000000400254
	0x00000000000000044	0x00000000000000044	R 4
GNU_EH_FRAME	0x000000000000009b8	0x00000000004009b8	0x00000000004009b8
	0x0000000000000003c	0x0000000000000003c	R 4
GNU_STACK	0x00000000000000000	0x0000000000000000	0x0000000000000000
	0x00000000000000000	0x0000000000000000	RW 10
GNU_RELRO	0x00000000000000e08	0x0000000000600e08	0x0000000000600e08
	0x000000000000001f8	0x000000000000001f8	R 1

# Symbolic Debugging

```
static unsigned int
count(char *buffer, size_t n)
{
    unsigned int sum = 0;
    size_t i;

    for (i = 0; i < n; i++)
        sum += buffer[i] == 'w';

    return sum;
}
```

Source is compiled into executable code that interacts with **memory** and a limited set of **registers**.

Dump of assembler code for function count:

```
0x00000000004008b0 <+0>:      push    %rbx
0x00000000004008b1 <+1>:      lea     0x1000(%rdi),%rdx
0x00000000004008b8 <+8>:      xor     %ebx,%ebx
0x00000000004008ba <+10>:     nopw    0x0(%rax,%rax,1)
0x00000000004008c0 <+16>:     xor     %eax,%eax
0x00000000004008c2 <+18>:     cmpb   $0x77,(%rdi)
0x00000000004008c5 <+21>:     sete   %al
0x00000000004008c8 <+24>:     add     $0x1,%rdi
0x00000000004008cc <+28>:     add     %eax,%ebx
0x00000000004008ce <+30>:     cmp     %rdx,%rdi
0x00000000004008d1 <+33>:     jne     0x4008c0 <count+16>
...
0x00000000004008d8 <+40>:     mov     %ebx,%eax
0x00000000004008da <+42>:     pop     %rbx
0x00000000004008db <+43>:     retq
```

End of assembler dump.

# Symbolic Debugging

A symbolic debugger is able to map the state of registers and memory to a backtrace with variables and type information.

```
sbahra$ bt `pgrep wc`  
PID: 8870  
-----  
Thread 8870  
[ 0] libc-2.21.so          pause (../sysdeps/unix/syscall-template.S:81)  
...  
[ 1] wc                    count (wc.c:19)  
    buffer = (parameter) reference(0, 0x7ffce7b64e10)  
    {char} -><> = string(0x7ffce7b64e10, 32, [root:x:0:0:root:/root:/bin/bash?])  
    sum = 0  
    i = --  
    n = (parameter) 4096  
[ 2] wc                    main (wc.c:42)  
...  
[ 3] libc-2.21.so          __libc_start_main  
...
```

# Symbolic Debugging

Mention that on Windows, other facilities are available.

The first step is to extract register state to determine the current state of the process.

## NAME

ptrace -- process tracing and debugging

## SYNOPSIS

```
#include <sys/types.h>
#include <sys/ptrace.h>
```

int

```
ptrace(int request, pid_t pid, caddr_t addr, int data);
```

Available in various broken ways on Linux and the BSD family of operating systems.

# Symbolic Debugging

The **%rip** register tells us the currently executing instruction in a program.

Thread 6870

[ 0] libc-2.21.so

pause (../sysdeps/unix/syscall-template.S:81)

rax=0xfffffffffffffdfe  
rbx=0  
rbp=0x10b6010  
r9=0  
r12=0x4007b0  
r15=0  
ss=0x2b  
fs=0

rdx=0x7ffce7b65e10  
rsi=0x7ff9b1ad5020  
rsp=0x7ffce7b64df8  
r10=0x3c1  
r13=0x7ffce7b65f10  
**rip=0x7ff9b15b5b70**  
ds=0  
gs=0

rcx=0x7ff9b15b5b70  
rdi=0x7ffce7b65e10  
r8=0x7ff9b1ad5020  
r11=0x246  
r14=0  
cs=0x33  
es=0  
eflags=0x246

# Symbolic Debugging

The kernel provides a mechanism to determine what the process address space looks like.

```
sbahra@broadwell$ cat /proc/`pgrep wc`/maps
00400000-00401000 r-xp 00000000 08:02 6034972      /home/sbahra/projects/backtrace/bangbangcon/wc
00600000-00601000 r--p 00000000 08:02 6034972      /home/sbahra/projects/backtrace/bangbangcon/wc
00601000-00602000 rw-p 00001000 08:02 6034972      /home/sbahra/projects/backtrace/bangbangcon/wc
010b6000-010d7000 rw-p 00000000 00:00 0          [heap]
7ff9b14eb000-7ff9b16ab000 r-xp 00000000 08:02 5246298  /lib/x86_64-linux-gnu/libc-2.21.so
7ff9b16ab000-7ff9b18ab000 ---p 001c0000 08:02 5246298  /lib/x86_64-linux-gnu/libc-2.21.so
7ff9b18ab000-7ff9b18af000 r--p 001c0000 08:02 5246298  /lib/x86_64-linux-gnu/libc-2.21.so
7ff9b18af000-7ff9b18b1000 rw-p 001c4000 08:02 5246298  /lib/x86_64-linux-gnu/libc-2.21.so
7ff9b18b1000-7ff9b18b5000 rw-p 00000000 00:00 0
7ff9b18b5000-7ff9b18d9000 r-xp 00000000 08:02 5246292  /lib/x86_64-linux-gnu/ld-2.21.so
7ff9b1ab1000-7ff9b1ab4000 rw-p 00000000 00:00 0
7ff9b1ad5000-7ff9b1ad8000 rw-p 00000000 00:00 0
7ff9b1ad8000-7ff9b1ad9000 r--p 00023000 08:02 5246292  /lib/x86_64-linux-gnu/ld-2.21.so
7ff9b1ad9000-7ff9b1ada000 rw-p 00024000 08:02 5246292  /lib/x86_64-linux-gnu/ld-2.21.so
7ff9b1ada000-7ff9b1adb000 rw-p 00000000 00:00 0
7ffce7b48000-7ffce7b69000 rw-p 00000000 00:00 0          [stack]
7ffce7bfb000-7ffce7bfb000 r--p 00000000 00:00 0          [vvar]
7ffce7bfb000-7ffce7bfd000 r-xp 00000000 00:00 0          [vdso]
ffffffffffffff600000-ffffffffffffff601000 r-xp 00000000 00:00 0          [vsyscall]
```

rip=0x7ff9b15b5b70



# Symbolic Debugging

**ELF** objects contain various debug-related sections.

Section	Description
.debug_line	Maps memory addresses to line numbers.
.debug_info	Type, variable and function information.
.debug_frame	Unwinding information by memory address.
.eh_frame	Unwinding information by memory address.

These sections and their contents are specified by the **DWARF** format on most UNIX-like systems.

# Symbolic Debugging

**.debug\_line** contains a sequence of operations that are executed by a state machine to form a giant matrix.

```
[0x000000dc] Extended opcode 2: set Address to 0x4008b0
[0x000000e7] Special opcode 12: advance Address by 0 to 0x4008b0 and Line by 7 to 8
[0x000000e8] Special opcode 119: advance Address by 8 to 0x4008b8 and Line by 2 to 10
[0x000000e9] Special opcode 121: advance Address by 8 to 0x4008c0 and Line by 4 to 14
[0x000000ea] Special opcode 200: advance Address by 14 to 0x4008ce and Line by -1 to 13
[0x000000eb] Special opcode 78: advance Address by 5 to 0x4008d3 and Line by 3 to 16
[0x000000ec] Special opcode 77: advance Address by 5 to 0x4008d8 and Line by 2 to 18
[0x000000ed] Advance PC by 4 to 0x4008dc
[0x000000ef] Extended opcode 1: End of Sequence
```

```
[0x000000f2] Extended opcode 2: set Address to 0x4006c0
[0x000000fd] Advance Line by 21 to 22
[0x000000ff] Copy
[0x00000100] Special opcode 131: advance Address by 9 to 0x4006c9 and Line by 0 to 22
[0x00000101] Advance PC by constant 17 to 0x4006da
[0x00000102] Special opcode 38: advance Address by 2 to 0x4006dc and Line by 5 to 27
[0x00000103] Extended opcode 4: set Discriminator to 1
[0x00000107] Set is_stmt to 0
[0x00000108] Special opcode 75: advance Address by 5 to 0x4006e1 and Line by 0 to 27
```

# Symbolic Debugging

**.debug\_line** contains a sequence of operations that are executed by a state machine to form a giant matrix.

Address	Source
0x4000	wc.c:80
0x4008	wc.c:20
0x400d	wc.c:19
0x401a	wc.c:10

# Symbolic Debugging

**.debug\_frame** or **.eh\_frame** contain a sequence of operations that are executed by a state machine to form a giant matrix of register states.

```
00000090 0000000000000002c 00000064 FDE cie=00000030 pc=00000000004006c0..000000000040078d
  DW_CFA_advance_loc: 1 to 00000000004006c1
  DW_CFA_def_cfa_offset: 16
  DW_CFA_offset: r6 (rbp) at cfa-16
  DW_CFA_advance_loc: 1 to 00000000004006c2
  DW_CFA_def_cfa_offset: 24
  DW_CFA_offset: r3 (rbx) at cfa-24
  DW_CFA_advance_loc: 7 to 00000000004006c9
  DW_CFA_def_cfa_offset: 4144
  DW_CFA_advance_loc1: 131 to 000000000040074c
  DW_CFA_remember_state
  DW_CFA_def_cfa_offset: 24
  DW_CFA_advance_loc: 1 to 000000000040074d
  DW_CFA_def_cfa_offset: 16
  DW_CFA_advance_loc: 1 to 000000000040074e
  DW_CFA_def_cfa_offset: 8
  DW_CFA_advance_loc: 1 to 000000000040074f
  DW_CFA_restore_state
  DW_CFA_
```

# Symbolic Debug

Mention David Watson's talk about exception handling.


.eh\_frame is included in programs that need to do exception unwinding.

**.debug\_frame** or **.eh\_frame** contain a sequence of operations that are executed by a state machine to form a giant matrix of register states.

```
000000d0 000000000000000044 000000a4 FDE cie=00000030 pc=4005e0..400645
    LOC          CFA      rbx      rbp      r12      r13      r14      r15      ra
000000000004005e0 rsp+8      u      u      u      u      u      u      c-8
000000000004005e2 rsp+16     u      u      u      u      u      c-16    c-8
000000000004005e4 rsp+24     u      u      u      u      c-24    c-16    c-8
000000000004005e9 rsp+32     u      u      u      c-32    c-24    c-16    c-8
000000000004005eb rsp+40     u      u      c-40    c-32    c-24    c-16    c-8
000000000004005f3 rsp+48     u      c-48    c-40    c-32    c-24    c-16    c-8
000000000004005fb rsp+56     c-56    c-48    c-40    c-32    c-24    c-16    c-8
0000000000040060a rsp+64     c-56    c-48    c-40    c-32    c-24    c-16    c-8
0000000000040063a rsp+56     c-56    c-48    c-40    c-32    c-24    c-16    c-8
0000000000040063b rsp+48     c-56    c-48    c-40    c-32    c-24    c-16    c-8
0000000000040063c rsp+40     c-56    c-48    c-40    c-32    c-24    c-16    c-8
0000000000040063e rsp+32     c-56    c-48    c-40    c-32    c-24    c-16    c-8
00000000000400640 rsp+24     c-56    c-48    c-40    c-32    c-24    c-16    c-8
00000000000400642 rsp+16     c-56    c-48    c-40    c-32    c-24    c-16    c-8
00000000000400644 rsp+8      c-56    c-48    c-40    c-32    c-24    c-16    c-8
```

# Symbolic Debugging

**.debug\_info** specifies how the program is structured. It includes **type**, **variable** and other **structural** information.



# Symbolic Debugging

## DW\_TAG\_compile\_unit

DW\_AT\_name = wc.c

DW\_AT\_ranges = A .. B

## DW\_TAG\_subprogram

DW\_AT\_name = count

DW\_AT\_type =

## DW\_TAG\_formal\_parameter

DW\_AT\_name = buffer

DW\_AT\_type =

## DW\_TAG\_variable

DW\_AT\_name = **sum**

DW\_AT\_type =

## DW\_TAG\_variable

DW\_AT\_name = i

DW\_AT\_type =

## DW\_TAG\_base\_type

DW\_AT\_byte\_size = 8

DW\_AT\_encoding = unsigned

DW\_AT\_name = "long unsigned"

## DW\_TAG\_typedef

DW\_AT\_name = size\_t

DW\_AT\_type =

# Symbolic Debugging

**DWARF** must support all kinds of aggressive compiler optimizations and has a complex expression state machine for expressing variable values.

```
static unsigned int
count(char *buffer, size_t n)
{
    unsigned int sum = 0;
    size_t i;

    for (i = 0; i < n; i++)
        sum += buffer[i] == 'w';

    return sum;
}
```

Dump of assembler code for function count:

0x00000000004008b0	<+0>:	push	%rbx
0x00000000004008b1	<+1>:	lea	0x1000(%rdi),%rdx
0x00000000004008b8	<+8>:	xor	%ebx,%ebx
0x00000000004008ba	<+10>:	nopw	0x0(%rax,%rax,1)
0x00000000004008c0	<+16>:	xor	%eax,%eax
0x00000000004008c2	<+18>:	cmpb	\$0x77,(%rdi)
0x00000000004008c5	<+21>:	sete	%al
0x00000000004008c8	<+24>:	add	\$0x1,%rdi
0x00000000004008cc	<+28>:	add	%eax,%ebx
0x00000000004008ce	<+30>:	cmp	%rdx,%rdi
0x00000000004008d1	<+33>:	jne	0x4008c0 <count+16>
...			
0x00000000004008d8	<+40>:	mov	%ebx,%eax
0x00000000004008da	<+42>:	pop	%rbx
0x00000000004008db	<+43>:	retq	

End of assembler dump.



# Symbolic Debugging

```
< 2><0x00000361>      DW_TAG_variable
                        DW_AT_name          "i"
                        DW_AT_decl_file     0x00000001 /home/sbakra/projects/backtrace/bangbangcon/wc.c
                        DW_AT_decl_line     0x0000000b
                        DW_AT_type          <0x00000029>
...
< 1><0x0000036b>      DW_TAG_subprogram
                        DW_AT_abstract_origin <0x00000332>
                        DW_AT_low_pc        0x004008b0
                        DW_AT_high_pc       <offset-from-lowpc>44
...
< 2><0x0000038f>      DW_TAG_variable
                        DW_AT_abstract_origin <0x00000356>
                        DW_AT_location      <loclist with 3 entries follows>
→ 4008b0 4008ba (DW_OP_lit0; DW_OP_stack_value)
→ 4008ba 4008cc (DW_OP_breg5 (rdi): 0; DW_OP_breg1 (rdx): 0; DW_OP_minus; DW_OP_plus_uconst: 4096; DW_OP_stack_value)
→ 4008cc 4008ce (DW_OP_breg5 (rdi): 0; DW_OP_breg1 (rdx): 0; DW_OP_minus; DW_OP_plus_uconst: 4095; DW_OP_stack_value)
```

```
Dump of assembler code for function count:
0x00000000004008b0 <+0>:      push    %rbx
0x00000000004008b1 <+1>:      lea     0x1000(%rdi),%rdx
0x00000000004008b8 <+8>:      xor     %ebx,%ebx
0x00000000004008ba <+10>:     nopw    0x0(%rax,%rax,1)
0x00000000004008c0 <+16>:     xor     %eax,%eax
0x00000000004008c2 <+18>:     cmpb    $0x77,(%rdi)
0x00000000004008c5 <+21>:     sete    %al
0x00000000004008c8 <+24>:     add     $0x1,%rdi
0x00000000004008cc <+28>:     add     %eax,%ebx
0x00000000004008ce <+30>:     cmp     %rax,%rdi
0x00000000004008d1 <+33>:     jne     0x4008c0 <count+16>
...
0x00000000004008d8 <+40>:     mov     %ebx,%eax
0x00000000004008da <+42>:     pop     %rbx
0x00000000004008db <+43>:     retq
End of assembler dump.
```

# Symbolic Debugging

```
< 2><0x00000361>      DW_TAG_variable
                        DW_AT_name          "i"
                        DW_AT_decl_file     0x00000001 /home/sbakra/projects/backtrace/bangbangcon/wc.c
                        DW_AT_decl_line     0x0000000b
                        DW_AT_type          <0x00000029>
...
< 1><0x0000036b>      DW_TAG_subprogram
                        DW_AT_abstract_origin <0x00000332>
                        DW_AT_low_pc         0x004008b0
                        DW_AT_high_pc        <offset-from-lowpc>44
...
< 2><0x0000038f>      DW_TAG_variable
                        DW_AT_abstract_origin <0x00000356>
                        DW_AT_location       <loclist with 3 entries follows>
4008b0 4008ba (DW_OP_lit0; DW_OP_stack_value)
4008ba 4008cc (DW_OP_breg5 (rdi): 0; DW_OP_breg1 (rdx): 0; DW_OP_minus; DW_OP_plus_uconst: 4096; DW_OP_stack_value)
4008cc 4008ce (DW_OP_breg5 (rdi): 0; DW_OP_breg1 (rdx): 0; DW_OP_minus; DW_OP_plus_uconst: 4095; DW_OP_stack_value)
```

```
(gdb) info locals
sum = 0
i = <optimized out>
(gdb) break *0x4008cc
Breakpoint 2 at 0x4008cc: file wc.c, line 14.
(gdb) r
14                sum += buffer[i] == 'w';
(gdb) info locals
sum = 0
i = 0
```

```
Dump of assembler code for function count:
0x00000000004008b0 <+0>:      push    %rbx
0x00000000004008b1 <+1>:      lea     0x1000(%rdi),%rdx
0x00000000004008b8 <+8>:      xor     %ebx,%ebx
0x00000000004008ba <+10>:     nopw    0x0(%rax,%rax,1)
0x00000000004008c0 <+16>:     xor     %eax,%eax
0x00000000004008c2 <+18>:     cmpb   $0x77,(%rdi)
0x00000000004008c5 <+21>:     sete   %al
0x00000000004008c8 <+24>:     add     $0x1,%rdi
0x00000000004008cc <+28>:     add     %eax,%ebx
0x00000000004008ce <+30>:     cmp     %rdx,%rdi
0x00000000004008d1 <+33>:     jne     0x4008c0 <count+16>
...
0x00000000004008d8 <+40>:     mov     %ebx,%eax
0x00000000004008da <+42>:     pop     %rbx
0x00000000004008db <+43>:     retq
End of assembler dump.
```

# Symbolic Debugging

DWARF is designed to support aggressive compiler optimizations.

So why is my debugger not unwinding correctly or missing crucial information such as variable values?

# Symbolic Debugging



Optimizations

Poor Debug  
Information

# Debug Information Quality

Though a debugger may show a value as optimized out, it may just be a side-effect of bad DWARF being emitted.

```
(gdb) info locals
sum = 0
i = <optimized out>
(gdb) break *0x4008cc
Breakpoint 2 at 0x4008cc: file wc.c, line 14.
(gdb) r
14                                sum += buffer[i] == 'w';
(gdb) info locals
sum = 0
i = 0
```

```
Dump of assembler code for function count:
0x00000000004008b0 <+0>:      push    %rbx
0x00000000004008b1 <+1>:      lea     0x1000(%rdi),%rdx
0x00000000004008b8 <+8>:      xor     %ebx,%ebx
0x00000000004008ba <+10>:     nopw    0x0(%rax,%rax,1)
0x00000000004008c0 <+16>:     xor     %eax,%eax
0x00000000004008c2 <+18>:     cmpb   $0x77,(%rdi)
0x00000000004008c5 <+21>:     sete   %al
0x00000000004008c8 <+24>:     add     $0x1,%rdi
0x00000000004008cc <+28>:     add     %eax,%ebx
0x00000000004008ce <+30>:     cmp     %rdx,%rdi
0x00000000004008d1 <+33>:     jne     0x4008c0 <count+16>
...
0x00000000004008d8 <+40>:     mov     %ebx,%eax
0x00000000004008da <+42>:     pop     %rbx
0x00000000004008db <+43>:     retq
End of assembler dump.
```

# Debug Information Quality

This program stores a loop counter into the same variable UINT64\_MAX times.

```
int
main(void)
{
    uint64_t limit;
    uint64_t vr_index;
    uint64_t vr_ac = 0;

    VOLATILE_STORE(&limit, UINT64_MAX);
    for (vr_index = 0; vr_index < limit; vr_index++)
        VOLATILE_STORE(&vr_ac, vr_index);

    pause();
    return vr_ac;
}
```

# Debug Information Quality

Different compilers emit DWARF at varying levels of **quality** and **accuracy**.

## GCC

```
Thread 11237
[ 1] induction_00      main
    limit = 18446744073709551615
    vr_index = 6676766858
    vr_ac = 6676766857
```

**CORRECT**

## clang

```
[ 1] induction_00      main
    vr_ac = 0
    vr_index = 0
    limit = -- [optimized out]
```

**INVALID DATA**

# Optimizations

Let's review optimizations that will impact availability of debug information.



# Variables and Optimization

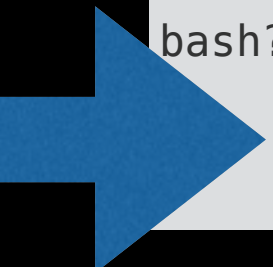
Unfortunately, things do **actually** get optimized out.

```
static unsigned int
count(char *buffer, size_t n)
{
    unsigned int sum = 0;
    size_t i;

    for (i = 0; i < n; i++)
        sum += buffer[i] == 'w';

    return sum;
}
```

```
[ 1] wc.gcc                                count.constprop.0
    buffer = (parameter) reference(0, 0x7ffe8ef83ac0)
    {char} -><> = string(0x7ffe8ef83ac0, 32, [root:x:0:0:root:/root:/bin/
bash?])
    sum = 0
    i = --
    n = (parameter) 4096
```



# Variables and Optimization

Unfortunately, things do **actually** get optimized out.

```
static unsigned int
count(char *buffer, size_t n)
{
    unsigned int sum = 0;
    char *end = buffer + n;

    do {
        sum += *buffer == 'w';
    } while (++buffer != end);

    return sum;
}
```

**i** is optimized out by **induction variable elimination**.

# Variables and Optimization

A good compiler is able to express optimized out constants in debug information.

```
int
main(void)
{
    const int x = 3931;

    return x;
}
```

Dump of assembler code for function main:

```
0x0000000000400400 <+0>:    mov     $0xf5b,%eax
0x0000000000400405 <+5>:    retq
```

**x** is optimized out by **constant value propagation**.

# Variables and Optimization

A good compiler is able to express optimized out constants in debug information.

```
< 2><0x000002a3>      DW_TAG_variable
                        DW_AT_name          "x"
                        DW_AT_decl_file     0x00000001 /home/sbakra/
lol.c
                        DW_AT_decl_line     0x00000006
                        DW_AT_type          <0x000002af>
                        DW_AT_const_value   3931
```

**x** is optimized out by **constant value propagation**.

# Variables and Optimization

Things involving more than one pass will confuse some compilers.

```
int
main(void)
{
    static unsigned int vr_value = 42;
    unsigned int vr_return = vr_value;

    pause();
    return vr_return;
}
```

```
(lldb) frame variable
(unsigned int) vr_value = <no location, value may have been optimized out>
```

Apple LLVM version 9.0.0 (clang-900.0.37)

# Variables and Optimization

Things involving more than one pass will confuse some compilers.

```
0x00000043:      TAG_variable [3]
                  AT_name( "vr_value" )
                  AT_type( {0x0000005b} ( unsigned int ) )
                  AT_decl_file( "/Users/sbakra/Projects/cdqs/src/constant_00.c" )
                  AT_decl_line( 6 )

0x0000004e:      TAG_variable [4]
                  AT_const_value( 0x0000002a )
                  AT_name( "vr_return" )
                  AT_decl_file( "/Users/sbakra/Projects/cdqs/src/constant_00.c" )
                  AT_decl_line( 7 )
                  AT_type( {0x0000005b} ( unsigned int ) )
```

Apple LLVM version 9.0.0 (clang-900.0.37)

# Variables and Optimization

GCC is able to handle transforming debug information along with passes.

```
int
main(void)
{
    static unsigned int vr_value = 42;
    unsigned int vr_return = vr_value;

    pause();
    return vr_return;
}
```

```
(gdb) p vr_return
$1 = 42
(gdb) p vr_value
$2 = 42
```

gcc (Ubuntu 4.8.4-2ubuntu1~14.04.3) 4.8.4

# Variables and Optimization

But, it still gets some things wrong.

```
< 2><0x0000004a>      DW_TAG_variable
                        DW_AT_name          "vr_value"
                        DW_AT_decl_file     0x00000001 /home/sbakra/cs.c
                        DW_AT_decl_line     0x00000004
                        DW_AT_type          <0x00000088>
                        DW_AT_const_value   42
< 2><0x00000056>      DW_TAG_variable
                        DW_AT_name          "vr_return"
                        DW_AT_decl_file     0x00000001 /home/sbakra/cs.c
                        DW_AT_decl_line     0x00000005
                        DW_AT_type          <0x0000008d>
                        DW_AT_const_value   42
```



# Variables and Optimization

But, it still gets some things wrong.

```
(gdb) ptype vr_value  
type = const unsigned int
```

Apple LLVM version 9.0.0 (clang-900.0.37)

```
(gdb) ptype vr_value  
type = unsigned int
```

gcc (Ubuntu 4.8.4-2ubuntu1~14.04.3) 4.8.4

Both GCC and clang get the type information wrong for vr\_value.

# Variables and Optimization

Certain constructs are simply unresolvable, even with optimizations turned off.

```
int
main(void)
{
    int value;

    struct {
        int apple : 2;
        int orange : 4;
        int tomato : 5;
    } x;

    ck_pr_store_int(&value, 1);
    x.apple = value;
    x.orange = value;
    x.tomato = value;
    pause();
    return x.apple;
}
```

```
(lldb) frame variable x
((anonymous struct)) x = <no location, value
```

Apple LLVM version 9.0.0 (clang-900.0.37)

```
(gdb) p x
$1 = {apple = 1, orange = 1, tomato = 1}
```

gcc (Ubuntu 4.8.4-2ubuntu1~14.04.3) 4.8.4

# Variables and Optimization

Other constructs are inaccurate or incomplete, with optimizations turned on.

```
static int vla_length = 3;

int
main(void)
{
    int vla[vla_length];

    vla[0] = 1;
    vla[1] = 1;
    vla[2] = 1;
    pause();
    return vla[0];
}
```

```
(lldb) frame variable vla
(int []) vla = <no location, value may have b
```

Apple LLVM version 9.0.0 (clang-900.0.37)

```
(gdb) p vla
$3 = <optimized out>
```

gcc (Ubuntu 4.8.4-2ubuntu1~14.04.3) 4.8.4

# Variables and Optimization

Live variable analysis

Unfortunately, things are **actually** non-recoverable.

```
int
main(int argc, const char **argv)
{

    VOLATILE_LOAD(argv);
    pause();
    return 0;
}
```

```
[ 1] spill_00      main (spill_00.c:9)
    argc = -- [optimized out]
    argv = -- [optimized out]
```

Registers are **extremely fast**. The compiler's **register allocator** tries to optimally allocate registers for variables.

# Variables and Optimization

Since **arg** is never used once `pause` is invoked, the compiler doesn't **spill** it from the register it lived in.

```
int
main(int argc, const char **argv)
{

    pause();
    VOLATILE_LOAD(argv);
    return 0;
}
```

```
argc = -- [optimized out]
argv = (parameter) reference(0, 0x7ffe022232c8)
    {pointer(const char)} -><> = reference(0x7ffe022232c8, 0x7ffe02224433)
        {const char} -><> = string(0x7ffe02224433, 10, [./spill_00])
```

# Variables and Optimization

+ \* According to the DWARF specification: The default rule for all  
+ \* columns before interpretation of the initial instructions is the  
+ \* undefined rule. However, an ABI authoring body or a compilation  
+ \* system authoring body may specify an alternate default value for any  
+ \* or all columns.

The platform **ABI** determines which registers must be saved across function calls, and whether spilling said register if used.

000000d0 00000000000000044 000000a4 FDE cie=00000030 pc=4005e0..400645

LOC	CFA	rbx	rbp	r12	r13	r14	r15	ra
000000000004005e0	rsp+8	u	u	u	u	u	u	c-8
000000000004005e2	rsp+16	u	u	u	u	u	c-16	c-8
000000000004005e4	rsp+24	u	u	u	u	c-24	c-16	c-8
000000000004005e9	rsp+32	u	u	u	c-32	c-24	c-16	c-8
000000000004005eb	rsp+40	u	u	c-40	c-32	c-24	c-16	c-8
000000000004005f3	rsp+48	u	c-48	c-40	c-32	c-24	c-16	c-8
000000000004005fb	rsp+56	c-56	c-48	c-40	c-32	c-24	c-16	c-8
0000000000040060a	rsp+64	c-56	c-48	c-40	c-32	c-24	c-16	c-8
0000000000040063a	rsp+56	c-56	c-48	c-40	c-32	c-24	c-16	c-8
0000000000040063b	rsp+48	c-56	c-48	c-40	c-32	c-24	c-16	c-8
0000000000040063c	rsp+40	c-56	c-48	c-40	c-32	c-24	c-16	c-8
0000000000040063e	rsp+32	c-56	c-48	c-40	c-32	c-24	c-16	c-8
00000000000400640	rsp+24	c-56	c-48	c-40	c-32	c-24	c-16	c-8
00000000000400642	rsp+16	c-56	c-48	c-40	c-32	c-24	c-16	c-8
00000000000400644	rsp+8	c-56	c-48	c-40	c-32	c-24	c-16	c-8

# Variables and Optimization

Different compilers and debuggers have different semantics for register value propagation in unwinding state machine.

```
(lldb) frame select 1
frame #1: 0x00007fffc5f5eb72 libsystem_c.dylib`nanosleep + 199
libsystem_c.dylib`nanosleep:
    0x7fffc5f5eb72 <+199>: testl    %eax, %eax
    0x7fffc5f5eb74 <+201>: jns     0x7fffc5f5eb3d                ; <+146>
    0x7fffc5f5eb76 <+203>: callq  0x7fffc5f644e4                ; symbol stub for: __error
    0x7fffc5f5eb7b <+208>: cmpl    $0x3c, (%rax)
(lldb) register read
General Purpose Registers:
    rbx = 0x0000000000000000
    rbp = 0x00007fff5fbff690
    rsp = 0x00007fff5fbff660
    r12 = 0x0000000000000000
    r13 = 0x00007fffcda1a20  __stderrp
    r14 = 0x00007fff5fbff6a8
    r15 = 0x00007fff5fbff6b8
    rip = 0x00007fffc5f5eb72  libsystem_c.dylib`nanosleep + 199
13 registers were unavailable.
```

# Variables and Optimization

Different compilers and debuggers have different semantics for register value propagation in unwinding state machine.

```
(gdb) frame 1
#1 [...]
(gdb) info reg
rax      0xfffffffffffffffffc      -4
rbx      0x0      0
rcx      0x7f0c75c5d693      139691492234899
rdx      0x2000      8192
rsi      0xbffbe0      12581856
rdi      0x9      9
rbp      0xbff960      0xbff960
rsp      0x7ffd11a7d290      0x7ffd11a7d290
r8       0x1b58      7000
r9       0x67ff410      109048848
r10      0xffffffff      4294967295
r11      0x293      659
r12      0xbffbec      12581868
r13      0xbffbe0      12581856
r14      0xbffbec      12581868
r15      0x0      0
rip      0x423930      0x423930 <cr_io_enter+48>
eflags   0x293      [ CF AF SF IF ]
```



# Variables and Optimization

Inlining and tail call optimization in particular, will lead to invalid debug information.

```
int
factorial(int ac, int n)
{
    if (--n == 1) { // 3 * 2 * 1 has occurred if we entered here
        pause();
        return ac;
    }

    return factorial(ac * n, n);
}

int
main(void)
{
    return factorial(3, 3);
}
```

# Variables and Optimization

Inlining and tail call optimization in particular, will lead to invalid debug information.

GCC	clang
<pre>  1  constant_03          factorial   n = (parameter) 1   ac = (parameter) 3</pre> <p data-bbox="570 1418 927 1463">INVALID DATA</p>	<pre>  1  constant_03          factorial   ac = (parameter) 3   n = (parameter) 2</pre> <p data-bbox="1671 1418 2027 1463">INVALID DATA</p>

# Variables and Optimization

Inlining and tail call optimization in particular, will lead to invalid debug information.

```
0x0000003d:      TAG_formal_parameter [3]
                AT_location( 0x00000000
                0x00000000100000d00 - 0x00000000100000d08: rdi
                0x00000000100000d08 - 0x00000000100000d79: rbx
                0x00000000100000daf - 0x00000000100000db1: rbx
                0x00000000100000eb3 - 0x00000000100000eba: rbx )
                AT_abstract_origin( {0x0000005c}"ac" )

0x00000046:      TAG_formal_parameter [3]
                AT_location( 0x0000005c
                0x00000000100000d00 - 0x00000000100000d06: rsi
                0x00000000100000d0b - 0x00000000100000eb5: rax )
                AT_abstract_origin( {0x00000067}"n" )
```

# Call stack and Optimizations

Frames might disappear in the presence of **tail call optimization**.

```
Thread 8163
[ 0] libc-2.21.so  pause
[ 1] chain        three
[ 2] chain        two
[ 3] chain        one
[ 4] chain        main
```



```
Thread 8224
[ 0] libc-2.21.so  pause
[ 1] chain        three
[ 2] chain        main
```

Tail calls are disambiguated by **call sites** and some times a debugger will only be able to extract the inner-most tail call in a chain.

# Call stack and Optimizations

Frames might disappear in the presence of **tail call optimization**.

```
main:
  callq one
  ...
one:
  callq two
  ...
two:
  callq three
  ...
three:
  retq
```



```
main:
  callq one
  ...
one:
  jmp two
  ...
two:
  jmp three
  ...
three:
  retq
```

# Call stack and Optimizations

Frames might disappear in the presence of **tail call optimization**.

```
Thread 8163
[ 0] libc-2.21.so  pause
[ 1] chain        three
[ 2] chain        two
[ 3] chain        one
[ 4] chain        main
```



```
Thread 8224
[ 0] libc-2.21.so  pause
[ 1] chain        three
[ 2] chain        main
```

Tail calls are disambiguated by **call sites** and some times a debugger will only be able to extract the inner-most tail call in a chain.

# Call stack and Optimizations

Your callstack can be **erroneous** with basic block commoning.

```
static void
function(const char *string)
{
    ck_pr_load_ptr(&string);
    pause();
    return;
}

static int
f(int x)
{
    if (x == 1) {
        function("a");
    } else if (x == 2) {
        function("b");
    }

    return 0;
}
```

f("a")	2  cbe_00	f (cbe_00.c:21)
f("b")	2  cbe_00	f (cbe_00.c:21)

Due to this optimization, the calls into this function are combined to a single instruction.

# Call stack and Optimizations

Adjacent loads and stores can be hoisted, leading to inaccurate line number information.

```
int
main(void)
{
    struct {
        char one;
        char two;
        char three;
        char four;
    } hoist;

    hoist.one = 1;
    hoist.two = 2;
    hoist.three = 3;
    hoist.four = 4;

    pause();
    ck_pr_rfo(&hoist);
    return 0;
}
```

```
movl    $0x4030201, -0x8(%rbp)    ; imm = 0x4030201
callq   0x100000f94              ; symbol stub for:
```

```
struct {
    char pad[4094];
    char one;
    char two;
    char three;
    char four;
} hoist;
```



# Call stack and Optimizations

NULL pointer checks can be deleted.

```
Thread A {  
    if (X != NULL)  
        do(X);  
  
    signal(B);  
    if (X != NULL)  
        do(X);  
}
```

Make sure to use `stdatomic` or `volatile` for use-cases that demand them.

# Next Time...

**<http://backtrace.io/blog>**

**@0xF390**

How do we measure debuggability?

How do we defeat the optimizer?

How do we marry DWARF with tracing?

...or horror stories so you can avoid creating your own.

# Symbolic Debugging

Some debuggers rely on system-facilities such as **thread\_db** to crawl internal thread data structures. These data structures are **easily corruptible**.

PID: 11756

---

Thread 11756

[ 0]	libpthread-2.21.so	pthread_join (pthread_join.c:90)
[ 1]	crash	crash_environment_create (crash.c:374)
[ 2]	crash	main (crash.c:570)
[ 3]	libc-2.21.so	__libc_start_main

# Symbolic Debugging

Some debuggers rely on system-facilities such as **thread\_db** to crawl internal thread data structures. These data structures are **easily corruptible**.

PID: 11756

Thread 11756

[ 0]	libpthread-2.21.so	pthread_join (pthread_join.c:90)
[ 1]	crash	crash_environment_create (crash.c:374)
[ 2]	crash	main (crash.c:570)
[ 3]	libc-2.21.so	__libc_start_main

Thread 11757

[ 0]	libpthread-2.21.so	__libc_pause (../sysdeps/unix/syscall-template.S:81)
[ 1]	crash	inlined (crash.c:104)
[ 2]	crash	crash_signal (crash.c:433)
[ 3]	libpthread-2.21.so	0x7f527f8f7d10 (../sysdeps/pthread/funlockfile.c:29)
* [ 4]	crash	scenario_smash_fault (smash.c:16)

Signal: [11] SIGSEGV (segmentation violation)  
Reason: SEGV\_ACCERR (invalid permissions for object)  
Faulting address: 0x7f527f215000

[ 5] 0x4141414141414141