

Read, Copy, Update... Then what?


Fedor G Pikus
Chief Scientist
Design2Silicon Division

September 27, 2017

Sometimes high performance is easy...

- To get maximum performance you need concurrency
- Simplest parallel problem:
read-only data+massive computations
 - Results have to be written somewhere...
 - Just lock it if it's infrequent
 - As long as input data does not change
- Many reader (consumer) threads run entirely independently with no contention or synchronization

Sometimes high performance is easy: searching read-only list




Life is rarely that easy

- Read-only data does happen
 - Often in scientific computations
- More often, data is “almost read-only”:
 - Changes happen but infrequently
- Many reader threads, one writer thread
 - Or write operation is locked, so effectively “one writer thread at a time”
- Reader performance should not be affected (ideally)

Almost read-only data

- Even if writes happen rarely, data can change while a reader is reading it
- Without any synchronization, this is a data race
- We definitely don't want the readers to wait on each other
- Ideally we don't want the readers to wait for the writer
- We can't have the readers block the writer indefinitely
- Ideally we don't want the writer to wait at all


Searching read-only list, lock-free (list with atomic shared pointers)




Almost read-only data

- Lock-free solution is generic, allows for multiple writers, scales well, but has overhead even in read-only case
 - Good solution when you need it
 - Wrong choice when updates are rare
- This is what read-write locks are for...

Searching read-only data with r/w locks (pthread_rwlock_t)



Searching changing data with r/w locks (pthread_rwlock_t)



Almost read-only data


- Even if writes happen rarely, data can change while a reader is reading it
- Without any synchronization, this is a data race
- There is reader overhead even if no updates actually happen during the test
 - Simply because they could have happened
- This is not fair!

There has to be a better way


Almost read-only data

- Lock-free solution is generic, allows for multiple writers, scales well, but has overhead even in read-only case
 - Good solution when you need it
 - Wrong choice when updates are rare
- This is what read-write locks are for...

Searching read-only data with r/w locks (pthread_rwlock_t)



Searching changing data with r/w locks (pthread_rwlock_t)




Almost read-only data

- Even if writes happen rarely, data can change while a reader is reading it
- Without any synchronization, this is a data race
- There is reader overhead even if no updates actually happen during the test
 - Simply because they could have happened
- This is not fair!

What is the cheapest atomic operation?

What is the cheapest atomic operation?




What can I do with just reads and writes?

- `std::atomic<node*> head;`
- Reader: `node* p = head; do_search(p);`
- Writer:
`node* new_node = new node(...); // Not visible to readers`
`new_node->next = head->next;`
`head = new_node;`

node or new_node?

// Now it's visible



What can I do with just reads and writes, carefully?

- `std::atomic<node*> head;`
- Reader:
`node* p = head.load(std::memory_order_acquire);`
`do_search(p);`
- Writer:
`node* new_node = new node(...);`
`new_node->next =`
`head.load(std::memory_order_relaxed)->next;`
`head.store(new_node, std::memory_order_release);`
- No writer synchronization – only one writer thread!
 - Or use mutex (or something else)

How does the reader work?

How does the writer work?

- Reader – use atomic read, otherwise nothing special:
`node* p = head.load(std::memory_order_acquire);`
`do_search(p);`

- Writer:

```
node* new_node = new node(...);
```

```
node* next =
```

```
    head.load(std::memory_order_relaxed)->next;
```

```
new_node->next = next;
```


```
head.store(new_node, std::memory_order_release);
```

- **Read, Copy, Update - RCU**

Read current data



Copy it to new data



Update current data



WARNING: confusing terminology ahead!

- “Read, Copy, Update” is actually a very small part of RCU
- RCU does stand for “Read, Copy, Update”
 - Somewhat of a misnomer
- No standard name for the “RCU” part of RCU
 - “Copy-on-write” in Java, but in C++ we use term COW for something entirely different
- Sometimes called “publishing protocol”

RCU, more generally

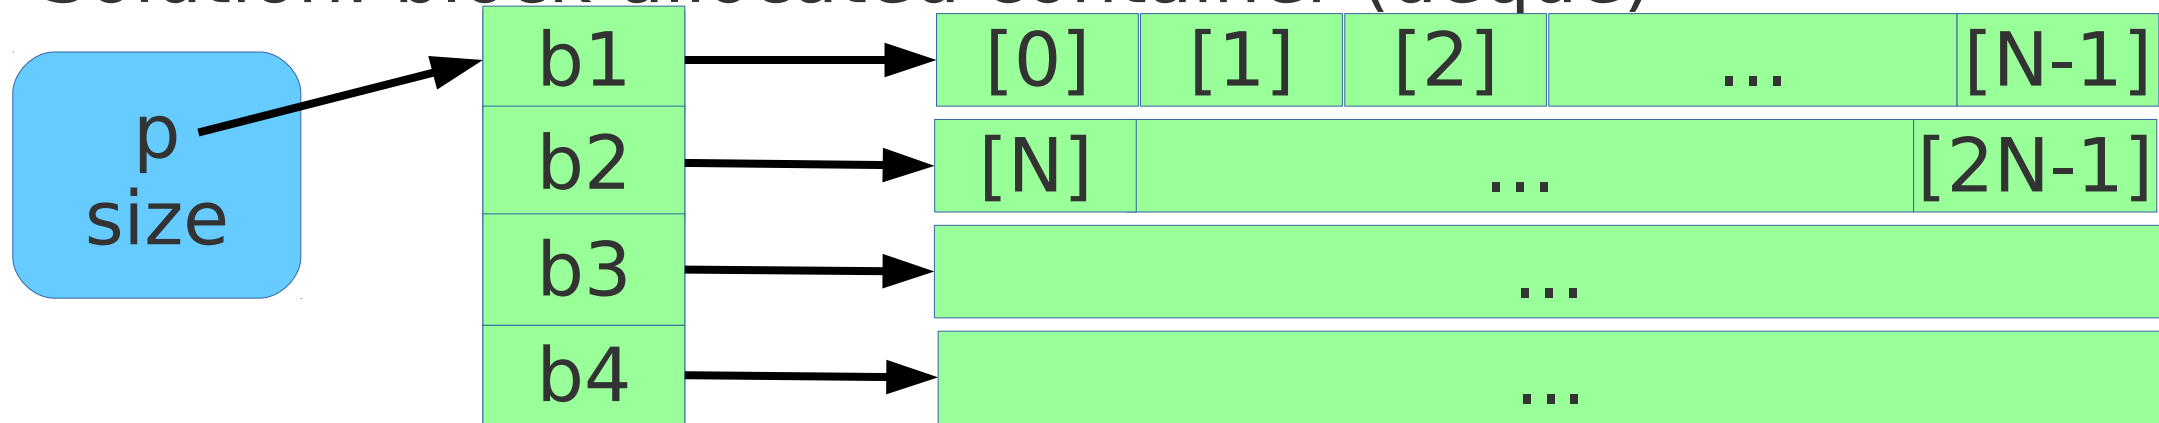
- Data is accessible through “root pointer”
 - Could be index into an array
 - Must be atomic
- Reader: acquire “root pointer” atomically, access data
- Writer: read current data, copy to new data, update new data, and publish it
- Some readers see old data, other readers see new data
 - Normal in concurrent systems

Example: thread-safe “growable array”

- Design an “array-like” container that can grow
 - Resize (grow) can be locked or limited to one thread
 - Array elements that existed before resize remain where they are and can be accessed by any thread, lock-free
- Resize does not invalidate pointers or iterators

First step: how to grow without moving old data?

- Solution: block-allocated container (deque)

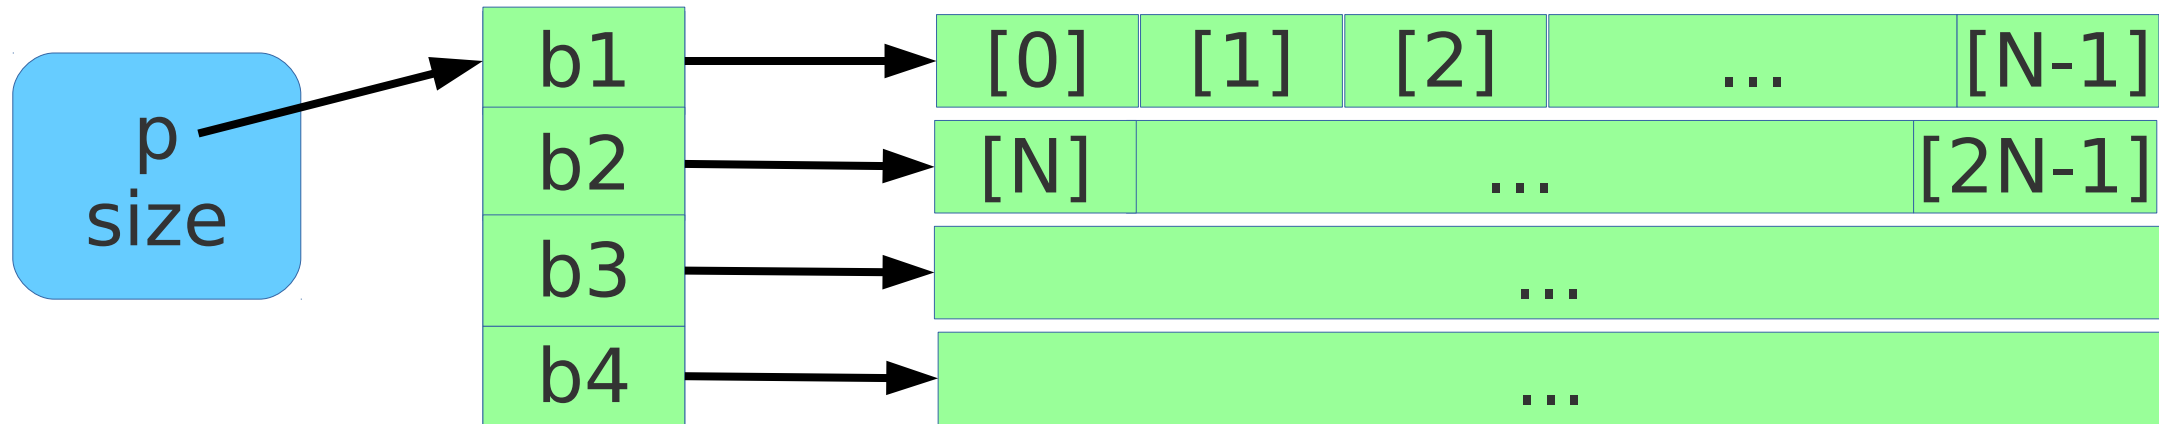


- Resize adds one or more blocks as needed
- Old blocks never move
- Operator[] needs one indirection, C[i]:

$N=2^m$
No % or /
Use & and >>

- read data block pointer from reference block $p[i\%N]$
- access array element in the data block $bx[i/N]$

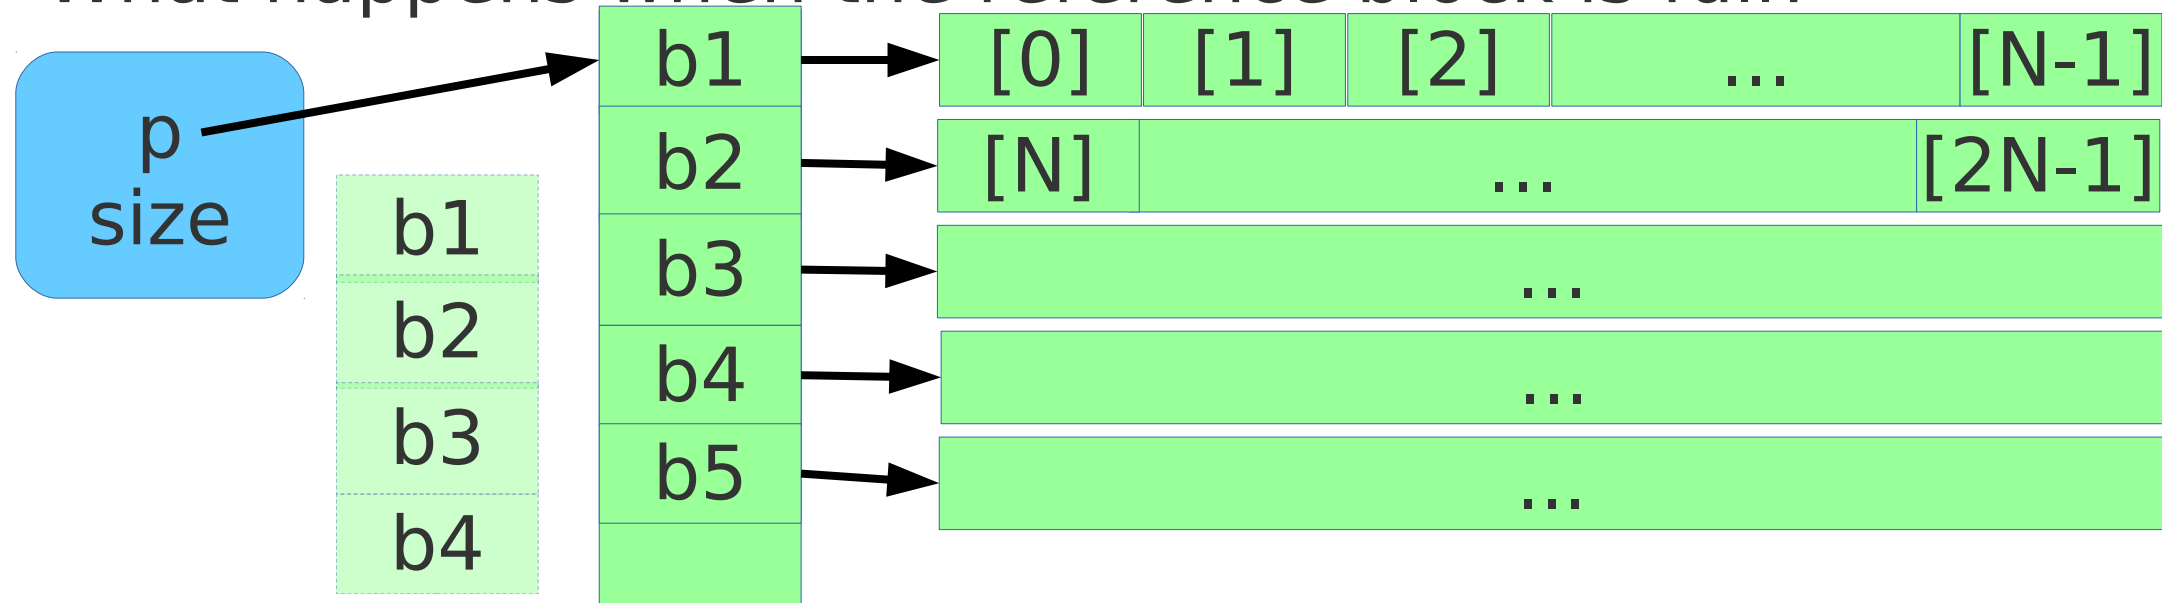
Now, with threads



- Writer **publishes** new blocks then updates size (release)
- Reader acquires size, guaranteed access to $C[0] \dots C[\text{size}-1]$
- Size is the “root pointer” (not p !)
- Resize is locked (or limited to one thread)
- `Operator[]` is not locked – no overhead for readers

Keep growing...

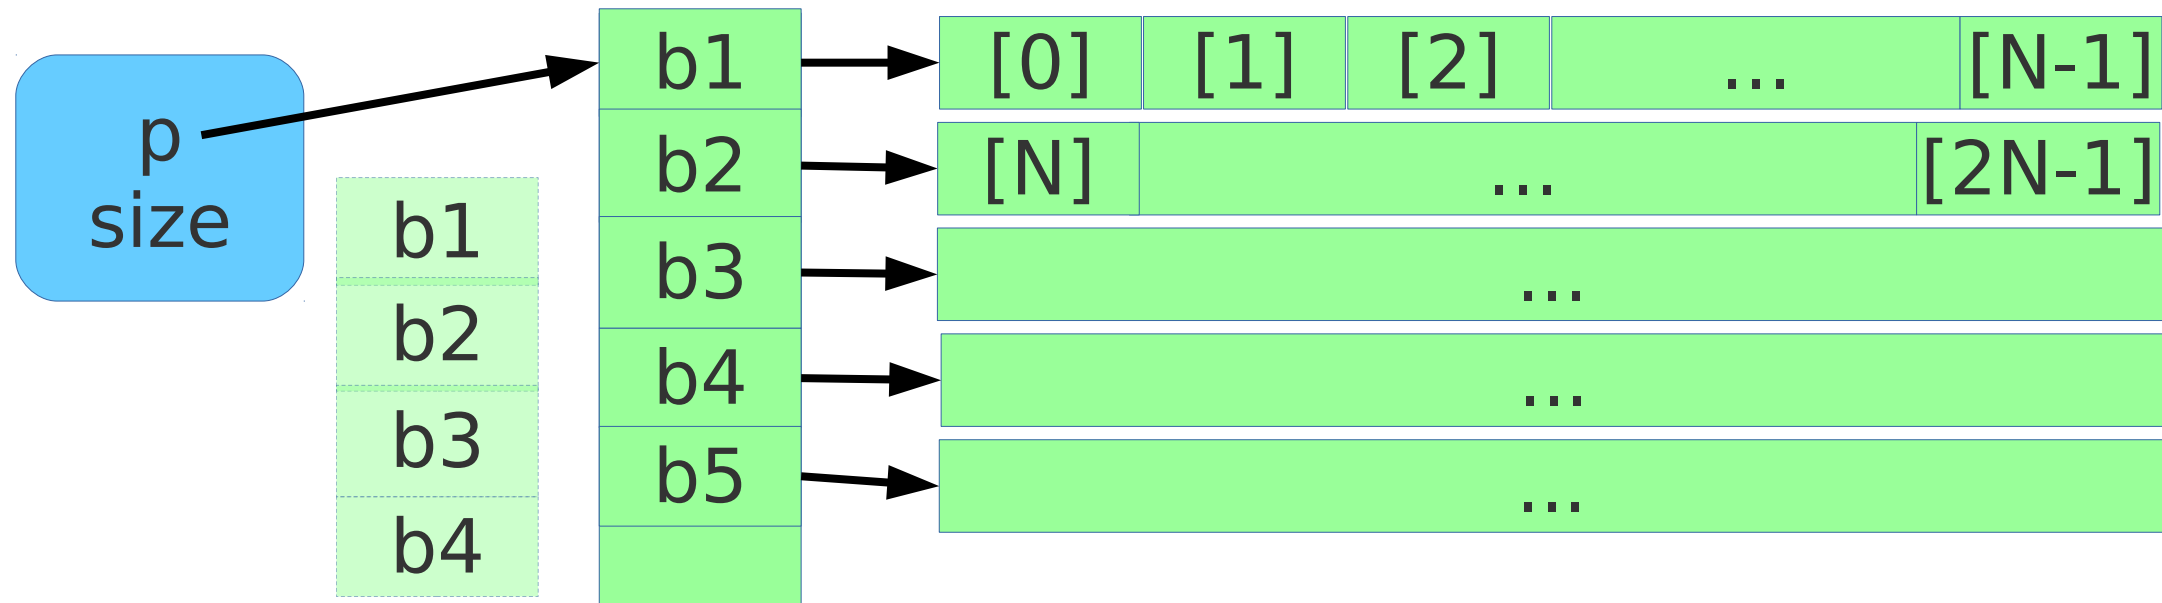
- What happens when the reference block is full?



- Just allocate a larger block, copy the old pointers, add new data block pointer
- RCU – read, copy, update!

RCU in action

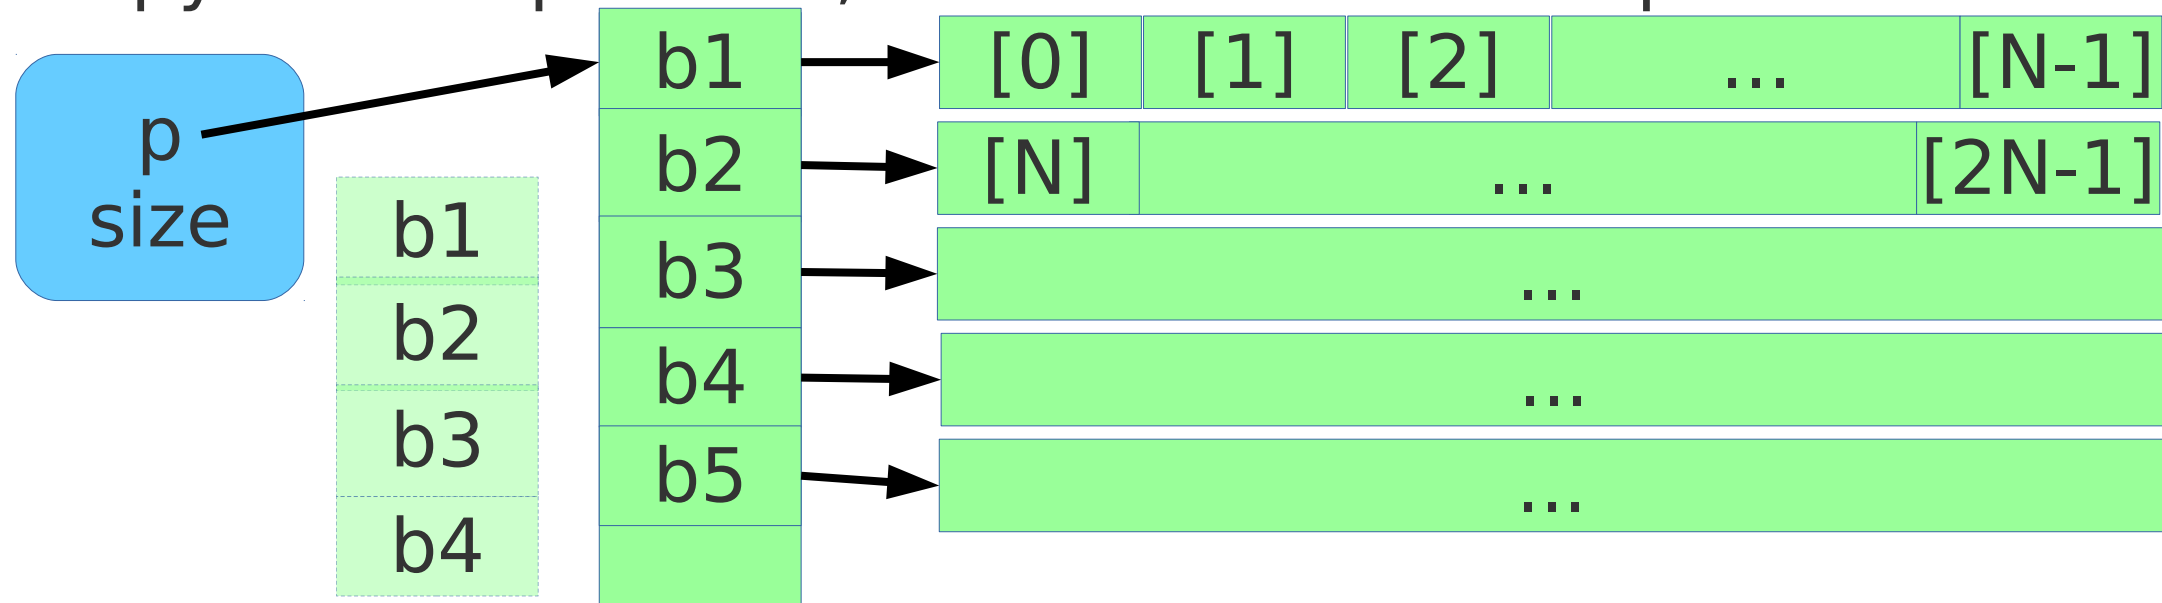
- When the reference block is full, allocate a larger block, copy the old pointers, add new data block pointer
- Still no overhead for readers – RCU rules



- What happens to the old reference block?

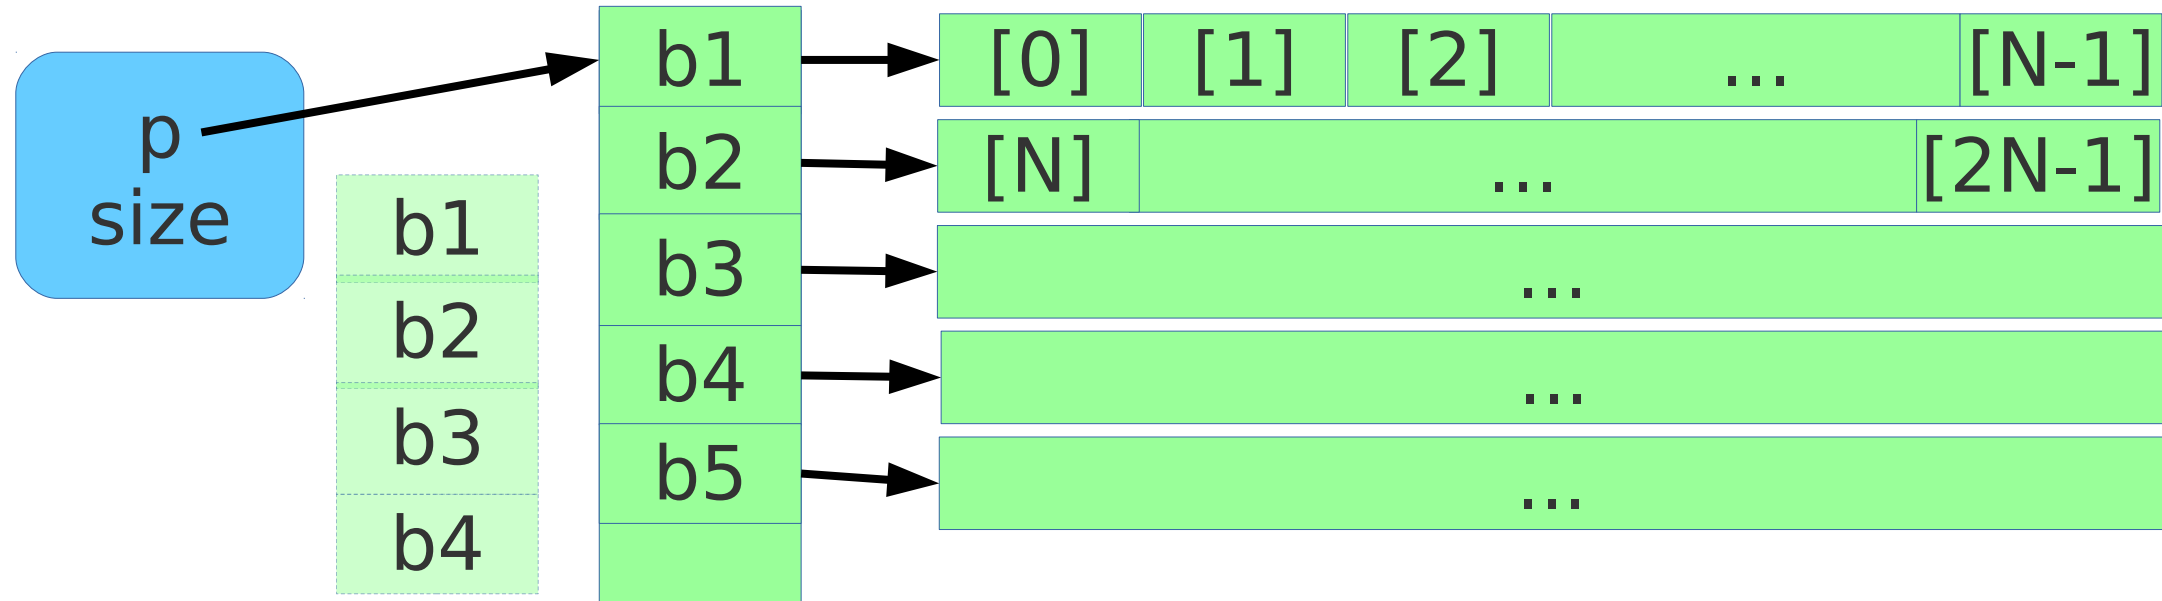
RCU in action

- When the reference block is full, allocate a larger block, copy the old pointers, add new data block pointer



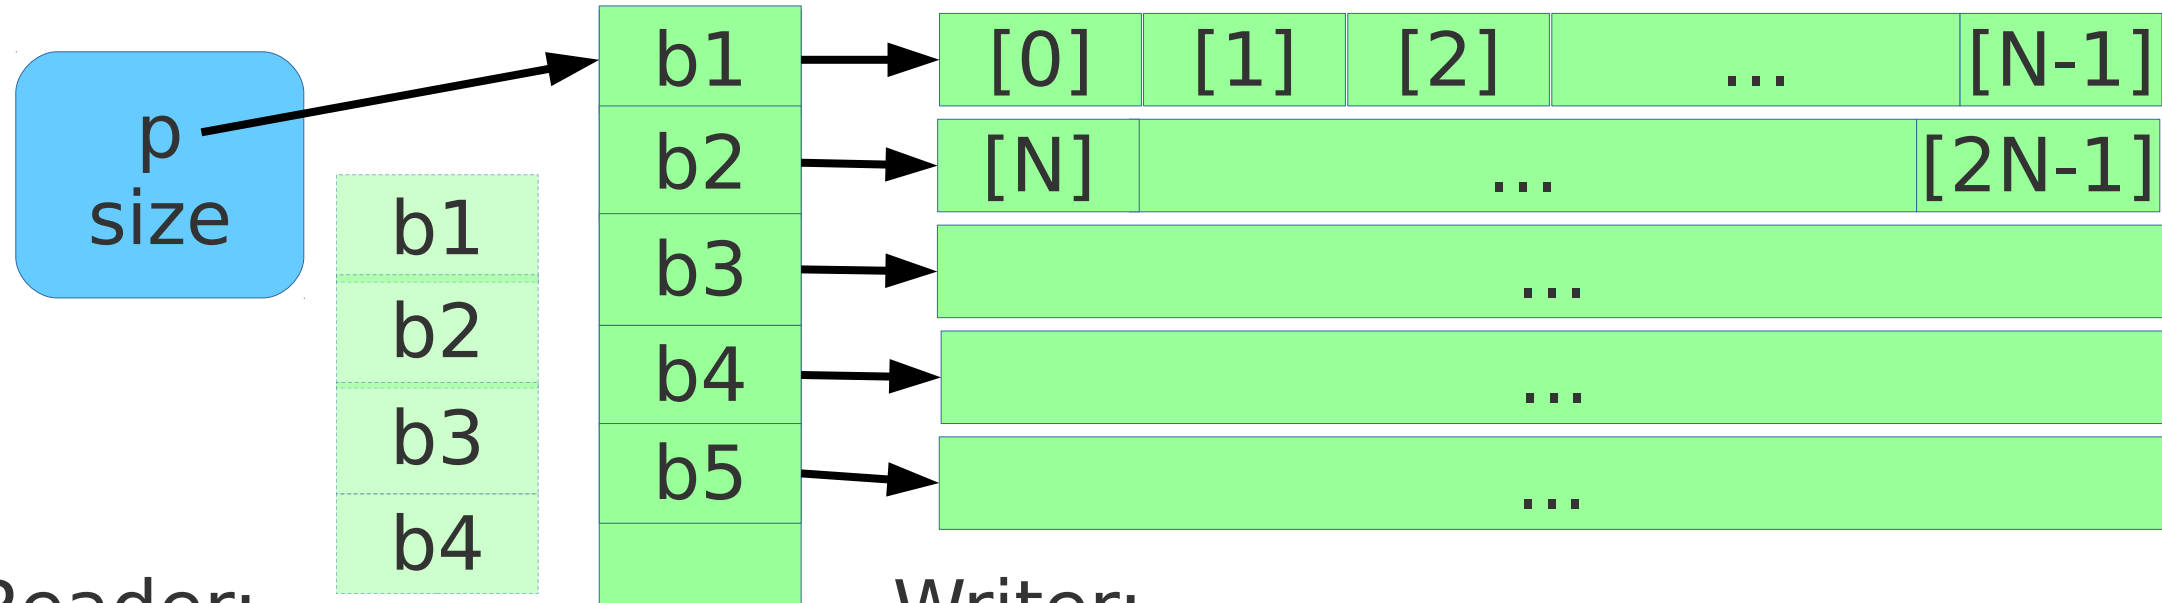
- Old reference block is no longer needed and cannot be reached – delete it

Follow the reader:



- $C[i]$:
split i into block index $ib=i\%N$ and data index $id=i/N$
get $bx=p[ib]$
 $C[i]$ is $bx[id]$

Bellevue, we have a problem

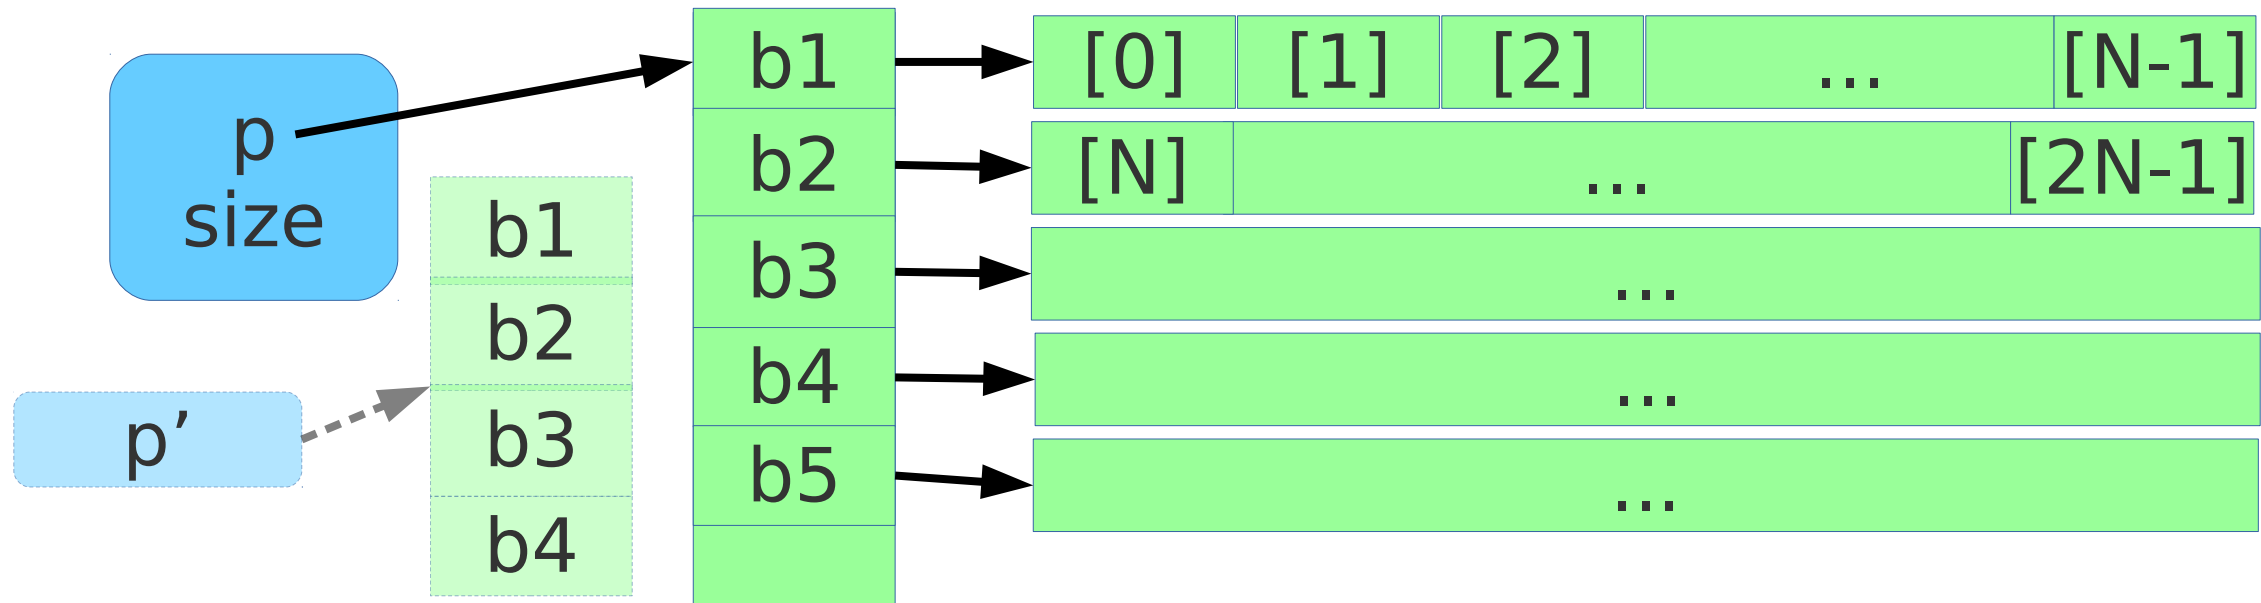


Reader:
compute ib and id
read p (as p')

get $bx = p'[ib]$
 $C[i]$ is $bx[id]$

Writer:
allocate new reference block
copy old bx pointers to new block
store new p, delete old ref block
seat back and watch the fireworks


Follow the reader:



- The writer should not delete the reference block as long as at least one reader can access it
- Soon after the resize, all readers will process the indirection $p'[ib]$, then old reference block can be deleted


It's always the delete

Follow the reader:



- $C[i]$:
split i into block index $ib=i\%N$ and data index $id=i/N$
get $bx=p[ib]$
 $C[i]$ is $bx[id]$

Bellevue, we have a problem




Reader:
compute ib and id
read p (as p')

get $bx=p'[ib]$
 $C[i]$ is $bx[id]$

Writer:
allocate new reference block
copy old bx pointers to new block
store new p , delete old ref block
seat back and watch the fireworks


Follow the reader:



- The writer should not delete the reference block as long as at least one reader can access it
- Soon after the resize, all readers will process the indirection $p'[ib]$, then old reference block can be deleted


It's always the delete

Follow the reader:



- $C[i]$:
split i into block index $ib=i\%N$ and data index $id=i/N$
get $bx=p[ib]$
 $C[i]$ is $bx[id]$

Bellevue, we have a problem




Reader:
compute ib and id
read p (as p')

get $bx=p'[ib]$
 $C[i]$ is $bx[id]$

Writer:
allocate new reference block
copy old bx pointers to new block
store new p , delete old ref block
seat back and watch the fireworks


Follow the reader:



- The writer should not delete the reference block as long as at least one reader can access it
- Soon after the resize, all readers will process the indirection $p'[ib]$, then old reference block can be deleted

It's always the delete

Follow the reader:



- $C[i]$:
split i into block index $ib=i\%N$ and data index $id=i/N$
get $bx=p[ib]$
 $C[i]$ is $bx[id]$

Bellevue, we have a problem



Reader:
compute ib and id
read p (as p')

get $bx=p'[ib]$
 $C[i]$ is $bx[id]$

Writer:
allocate new reference block
copy old bx pointers to new block
store new p , delete old ref block
seat back and watch the fireworks

It's always the delete



Reader:
compute ib and id
read p (as p')

get $bx=p'[ib]$
 $C[i]$ is $bx[id]$

Writer:
allocate new reference block
copy old bx pointers to new block
store new p , delete old ref block
seat back and watch the fireworks

It's always the delete

- Deletion of old data (reclaiming memory) is the main challenge in almost all lock-free programs
 - Inserting data is easy
- The problem is that other threads (readers) may hold “stale” reference to the data that the writer already made inaccessible from the “root pointer”
- The problem is not unique to the RCU but the solution is

WARNING: confusing terminology ahead!

- “Read, Copy, Update” - the publishing protocol – does not make your concurrency synchronization scheme an RCU
 - You can use publishing protocol with atomic shared pointers or hazard pointers
- RCU is distinguished by the specific way of memory reclamation
- RCU also implies that instead of changing data in place, the writer publishes a new copy of the data
 - Read old data, copy it, and update it
 - Then delete the old data “the RCU way”

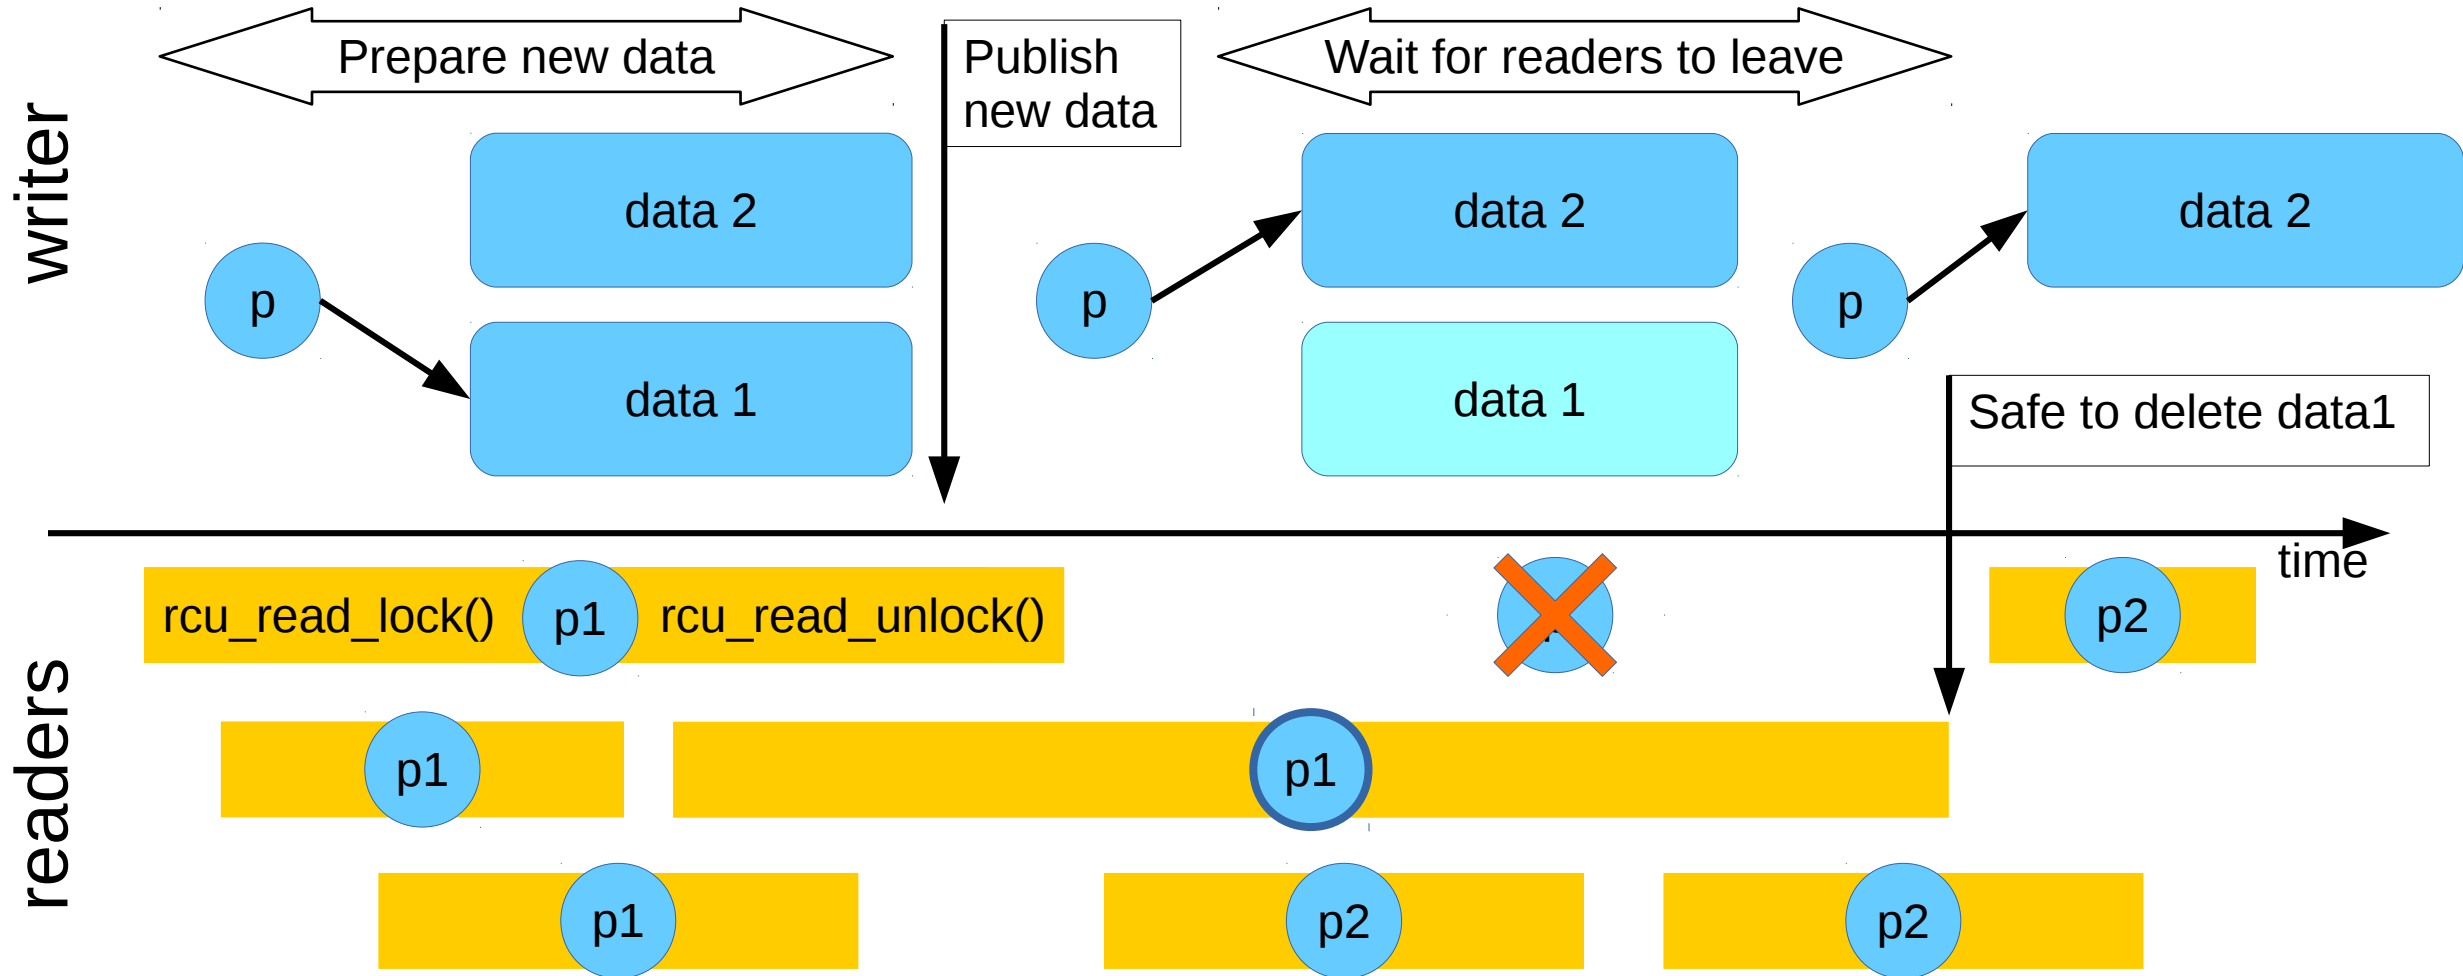
RCU memory reclamation - the real RCU (no Read, Copy, or Update)

- RCU uses cooperative protocol to track when it is safe to reclaim memory (when no reader can access it)
- Readers **MUST** follow these steps to access shared data:
 - 1) Call `rcu_read_lock()` to request access
 - 2) Get the root pointer (not use the old copy)
 - 3) Call `rcu_read_unlock()` to announce end of access
- Readers may access shared data only between the calls to `rcu_read_lock()` and `rcu_read_unlock()`
 - reader-side critical section
 - readers in “quiescent state” don’t read shared data

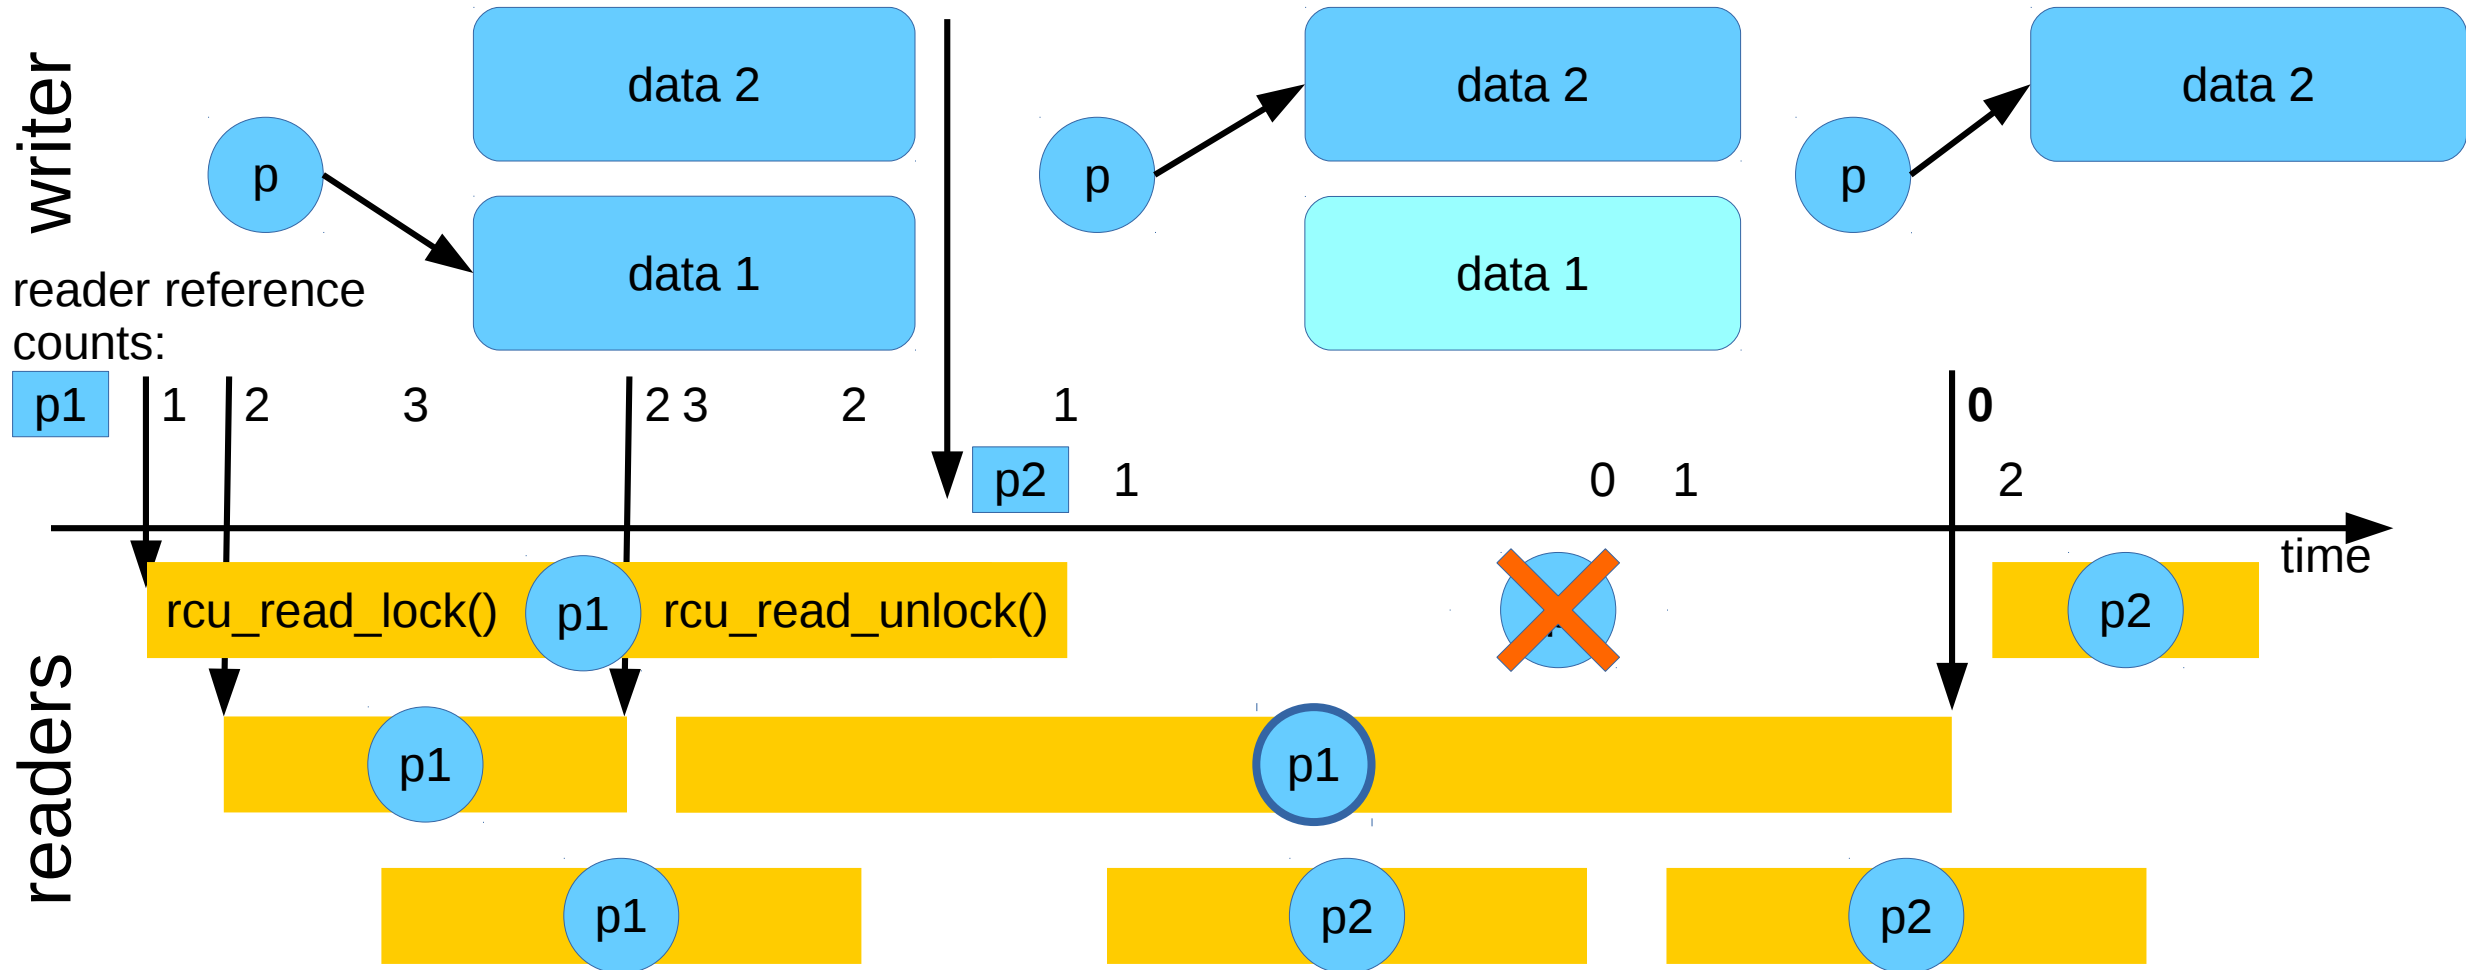
RCU memory reclamation - the real RCU (no Read, Copy, or Update)

- RCU uses cooperative protocol to track when it is safe to reclaim memory (when no reader can access it)
- Writer MUST follow these steps to modify shared data:
 - 1) Make old shared data inaccessible from the root
 - 2) Call `synchronize_rcu()` to wait for all readers who called `rcu_read_lock()` before step 1 to call `rcu_read_unlock()`
 - 3) Delete old data and reclaim the memory
- We don't need to wait for all readers to exit critical section, only those who acquired the old root pointer

RCU protocol



RCU implementation - the main idea



RCU under any other name

- RCU implementations use different function names:
- Readers entering critical section: `rcu_read_lock`, `rcu_enter`
- Readers leaving critical section: `rcu_read_unlock`, `rcu_leave`
- Writer waiting for readers before reclaiming memory: `synchronize_rcu`, `wait_for_readers_to_leave`
- C++ standard proposal WG21/P0461R1 uses names `rcu_read_lock`, `rcu_read_unlock`, and `synchronize_rcu`
- RCU implementations may use additional functions to register threads, defer deletion callbacks, etc


User-space RCU vs kernel RCU

- RCU is used in Linux kernel extensively
- Kernel RCU has several advantages
 - Kernel knows when context switches happen (readers leave critical section)
 - Kernel knows how many threads are there, which ones are running, etc
 - Kernel RCU does not need (extra) memory barriers
- The basic idea is the same but it's important to understand the implicit assumptions


RCU vs GC

- RCU does memory reclamation when that memory is unreachable (just like garbage collection)
- Strictly speaking, GC implies “automatic” (by definition)
 - “User-driven GC” is a contradiction in terms
- “User-driven GC” is a term that’s often used and I don’t know of a good alternative term
 - E.g. allocate a lot of objects on a memory pool then free the entire pool without deallocating each object
 - Memory reclamation in RCU is a kind of user-driven GC

Searching read-only (or not) data with RCU



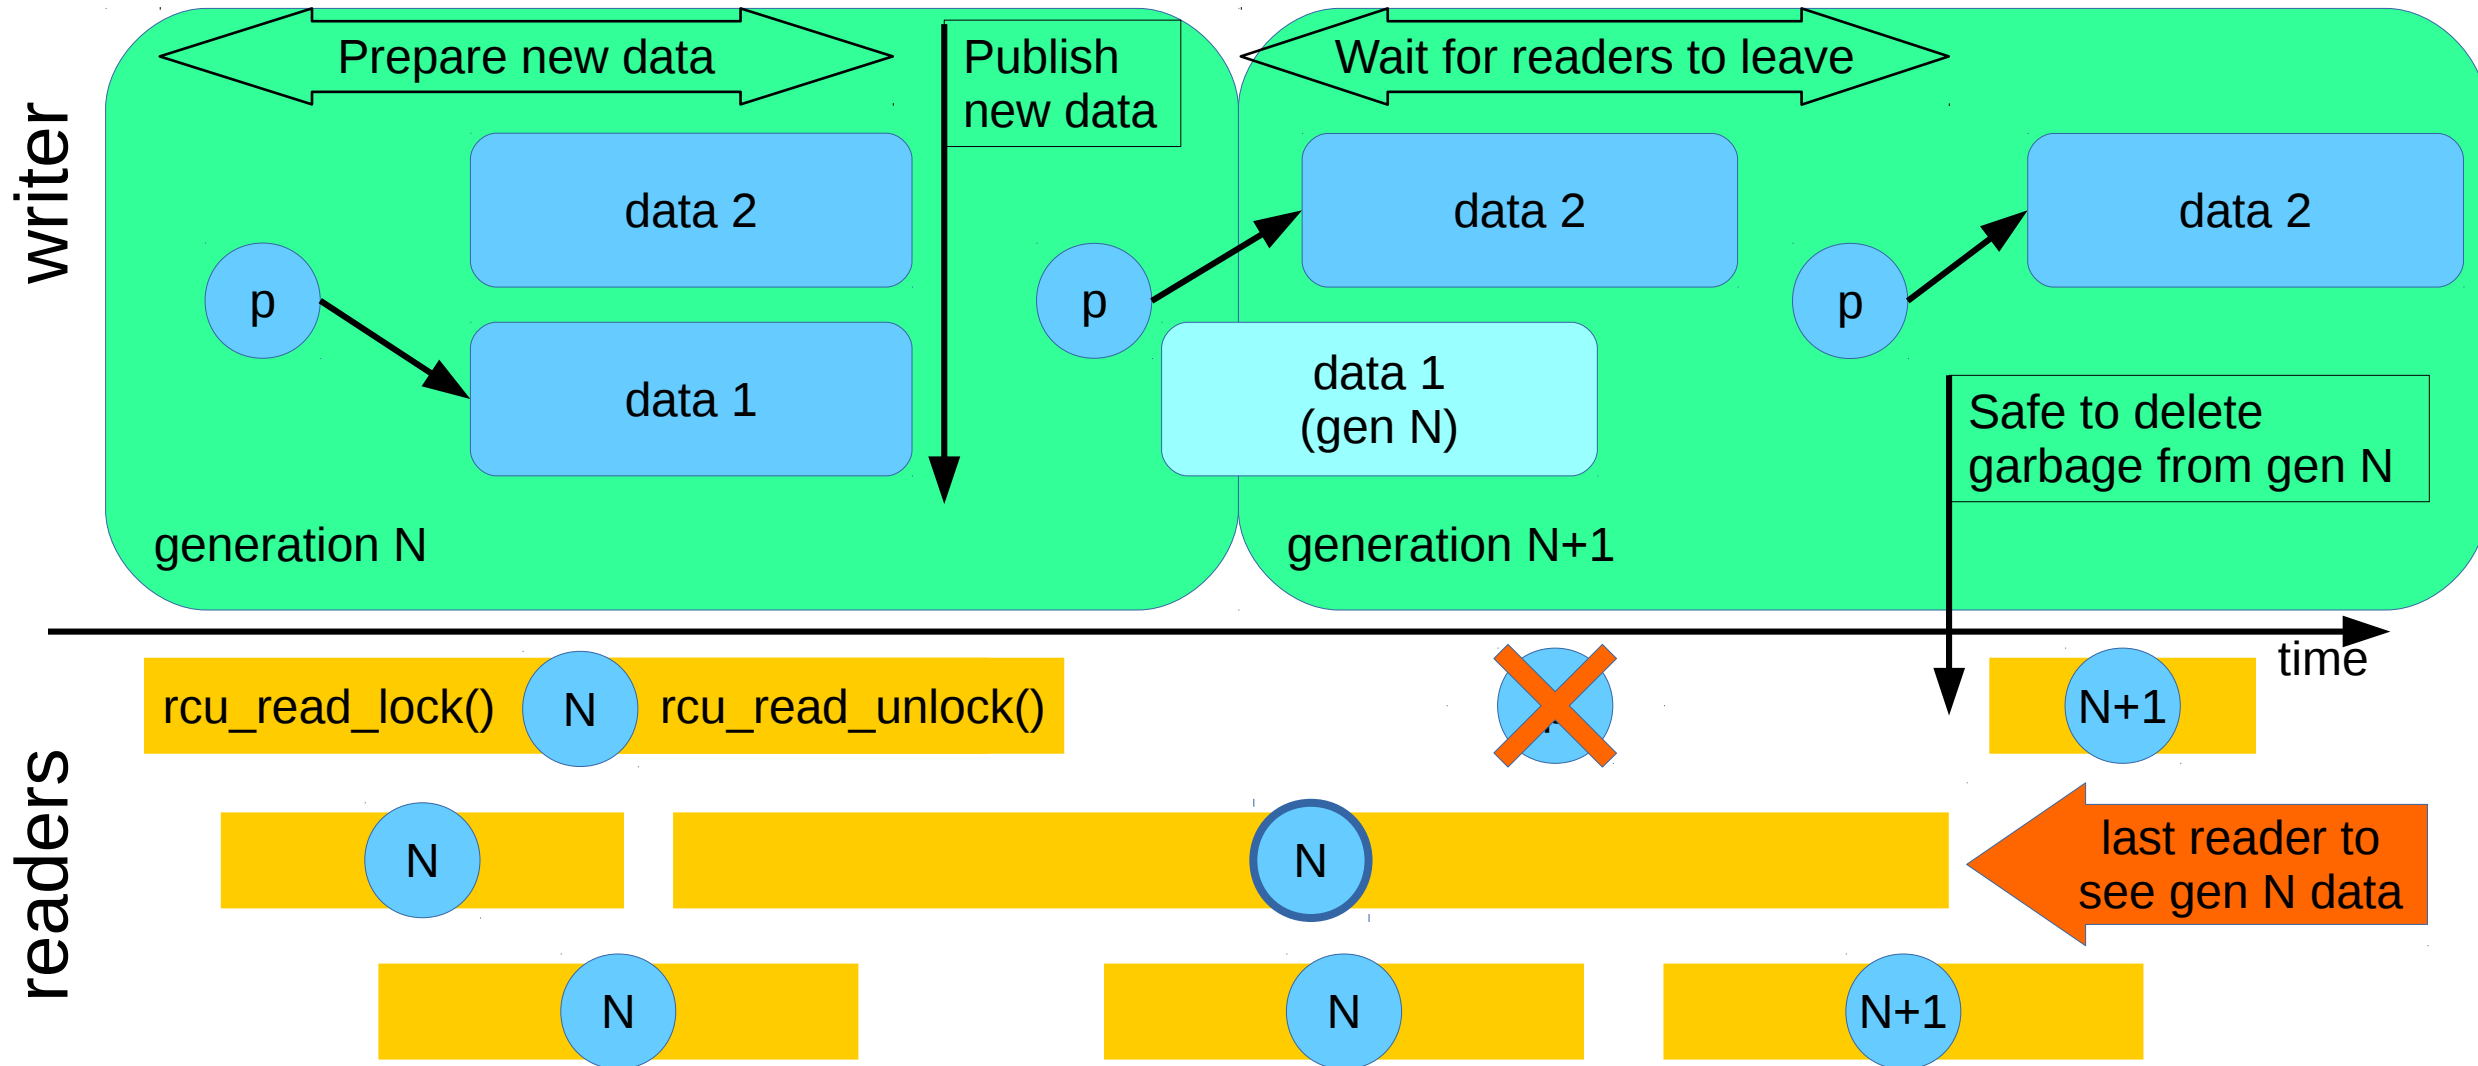
Searching data with (another) RCU



User-space RCU implementation (one of many)

- Writer maintains a global atomic generation (epoch)
`std::atomic<unsigned long> generation;`
- All data currently live belong to the current generation
- When the writer does an update, the old data is placed in a garbage queue
- There is one garbage queue per generation
- At some point the writer increments the generation
- Readers can access only current generation data
 - But can keep accessing it after it becomes garbage
- Writer tracks how many readers access each generation

User-space generational RCU



User-space RCU implementation:

Reader interface

```
std::atomic<unsigned long> generation;  
std::atomic<unsigned long> refcount[max_generations];  
class handle_t {...}; // Contains reader generation  
handle_t rcu_read_lock() {  
    size_t cg=generation;  
    ++refcount[cg];  
    return handle_t(cg);  
}  
void rcu_read_unlock(handle_t handle) {  
    --refcount[handle.get()];  
}
```

User-space RCU careful implementation:

Reader interface

```
handle_t rcu_read_lock() {  
    size_t cg=generation;  
    ++refcount[cg];  
    return handle_t(cg);  
}
```

- But aren't atomics `memory_order_seq_cst` by default? And isn't `memory_order_seq_cst` expensive? - Yes, and can be.
- Real implementation will use acquire/release fences and `memory_order_relaxed` as much as possible
 - Minimum necessary barriers depend on details of the implementation (one `seq_cst` is usually unavoidable)



User-space RCU implementation: Writer interface

```
std::atomic<unsigned long> generation;
std::atomic<unsigned long> refcount[max_generations];
std::queue<data_t*> garbage[max_generations];
unsigned long last_gc_gen = 0;
void synchronize_rcu() {
    unsigned long last_gen = generation++;
    while (last_gc_gen < last_gen) {
        while (refcount[last_gc_gen] > 0) wait();
        delete_garbage(garbage[last_gc_gen]);
        ++last_gc_gen;
    }
}
```

RCU implementations

- All RCU implementations will have `rcu_read_lock()`, `rcu_read_unlock()`, and `synchronize_rcu()`
 - Or `enter()`, `leave()`, and `wait_for_readers_to_leave()`
- Then you can get creative...

RCU implementations: readers

- Reader implementations may try to minimize `read_lock()` and `read_unlock()` overhead
 - Possibly by imposing some restrictions
- Why reader overhead?
 - Read generation number
 - Increment reference count  read-modify-write shared data
 - Read root pointer  acquire memory barrier
- Cache contention on the reference count is the main overhead (in user-space RCU)

RCU implementations: readers

- Reference count is shared between all readers
- Possible solution: give each reader its own count

```
handle_t rcu_read_lock(size_t reader_id) {  
    size_t cg=generation;  
    ++refcount[reader_id][cg]; // Padded to cache line  
    return handle_t(cg);  
}
```
- Reader ID is a thread ID: 0, 1, 2...
- Reader threads must register in advance with the RCU
- Writer must loop over all reader slots to add up the count

RCU implementations: readers

- Can reader critical sections be nested? If not, each reader may access only one version of shared data at a time, no need for a reference count array:

```
void rcu_read_lock(size_t reader_id) {  
    size_t cg=generation;  
    reader_gen[reader_id] = cg; // Padded to cache line  
}  
  
const size_t NO_READER = 0; // ULONG_MAX, etc  
void rcu_read_unlock(size_t reader_id) {  
    reader_gen[reader_id] = NO_READER;  
}
```


RCU implementations: readers

- What if thread registration is impossible?
- Can still reduce reader contention by using an array of reference counts, index is a hash of thread ID
- Hash collisions can happen, so reference counts must be atomically incremented
 - But contention is greatly reduced
 - Remember false sharing, one count per cache line!

RCU implementations: writers


- `synchronize_rcu()` vs `call_rcu()`
- RCU writer may defer deletion callbacks and queue them:

```
data_t* old_data = current_data;  
data_t* new_data = new data_t(*old_data);  
new_data->update();  
current_data = new_data;  
call_rcu(old_data, deleter); // Deferred until readers leave
```
- Callbacks are executed in batches by the writer thread or a special cleanup thread
 - Deletion callbacks are deferred until all readers leave the generation to which callbacks belong


RCU implementations: writers

- Granularity of updates vs granularity of cleanup
- Writer may process many updates in one generation
 - Old data from earlier updates cannot be reclaimed until generation changes
 - readers can enter current generation any time, only old generations are not re-readable once all readers leave
- Guard object – each update of each root pointer effectively advances the generation
 - More overhead, less unclaimed garbage


Short-lived vs long-lived generations



RCU performance - slowly changing data



RCU performance - rapidly changing data



RCU implementations: writers

- Readers do not wait for writers, ever – check
- Writer can do updates without waiting for readers – check
- Eventually writer will run out of memory...
- Ideally writer can reclaim memory between updates
- No-wait `synchronize_rcu()`, or `GC()`, or `do_callbacks()` - reclaim memory that no reader can see, right now
 - If the garbage queue is long, process a batch then return, to avoid delaying updates
- `wait_for_readers_to_leave(timeout)` – blocking wait with a timeout, reclaim what you can while waiting

RCU implementations: writers

- RCU itself does nothing for multiple writers
 - RCU works well when updates are infrequent – one writer thread should be able to handle the load
 - Or one writer thread at a time – lock or CAS
- Preparing new data may be time-consuming, so multiple writer threads are often desirable
- Some implementations support concurrent calls to `synchronize_rcu()` - makes multiple writers easier
 - Garbage queue may be per-thread or global, per-generation or sentinel-delimited

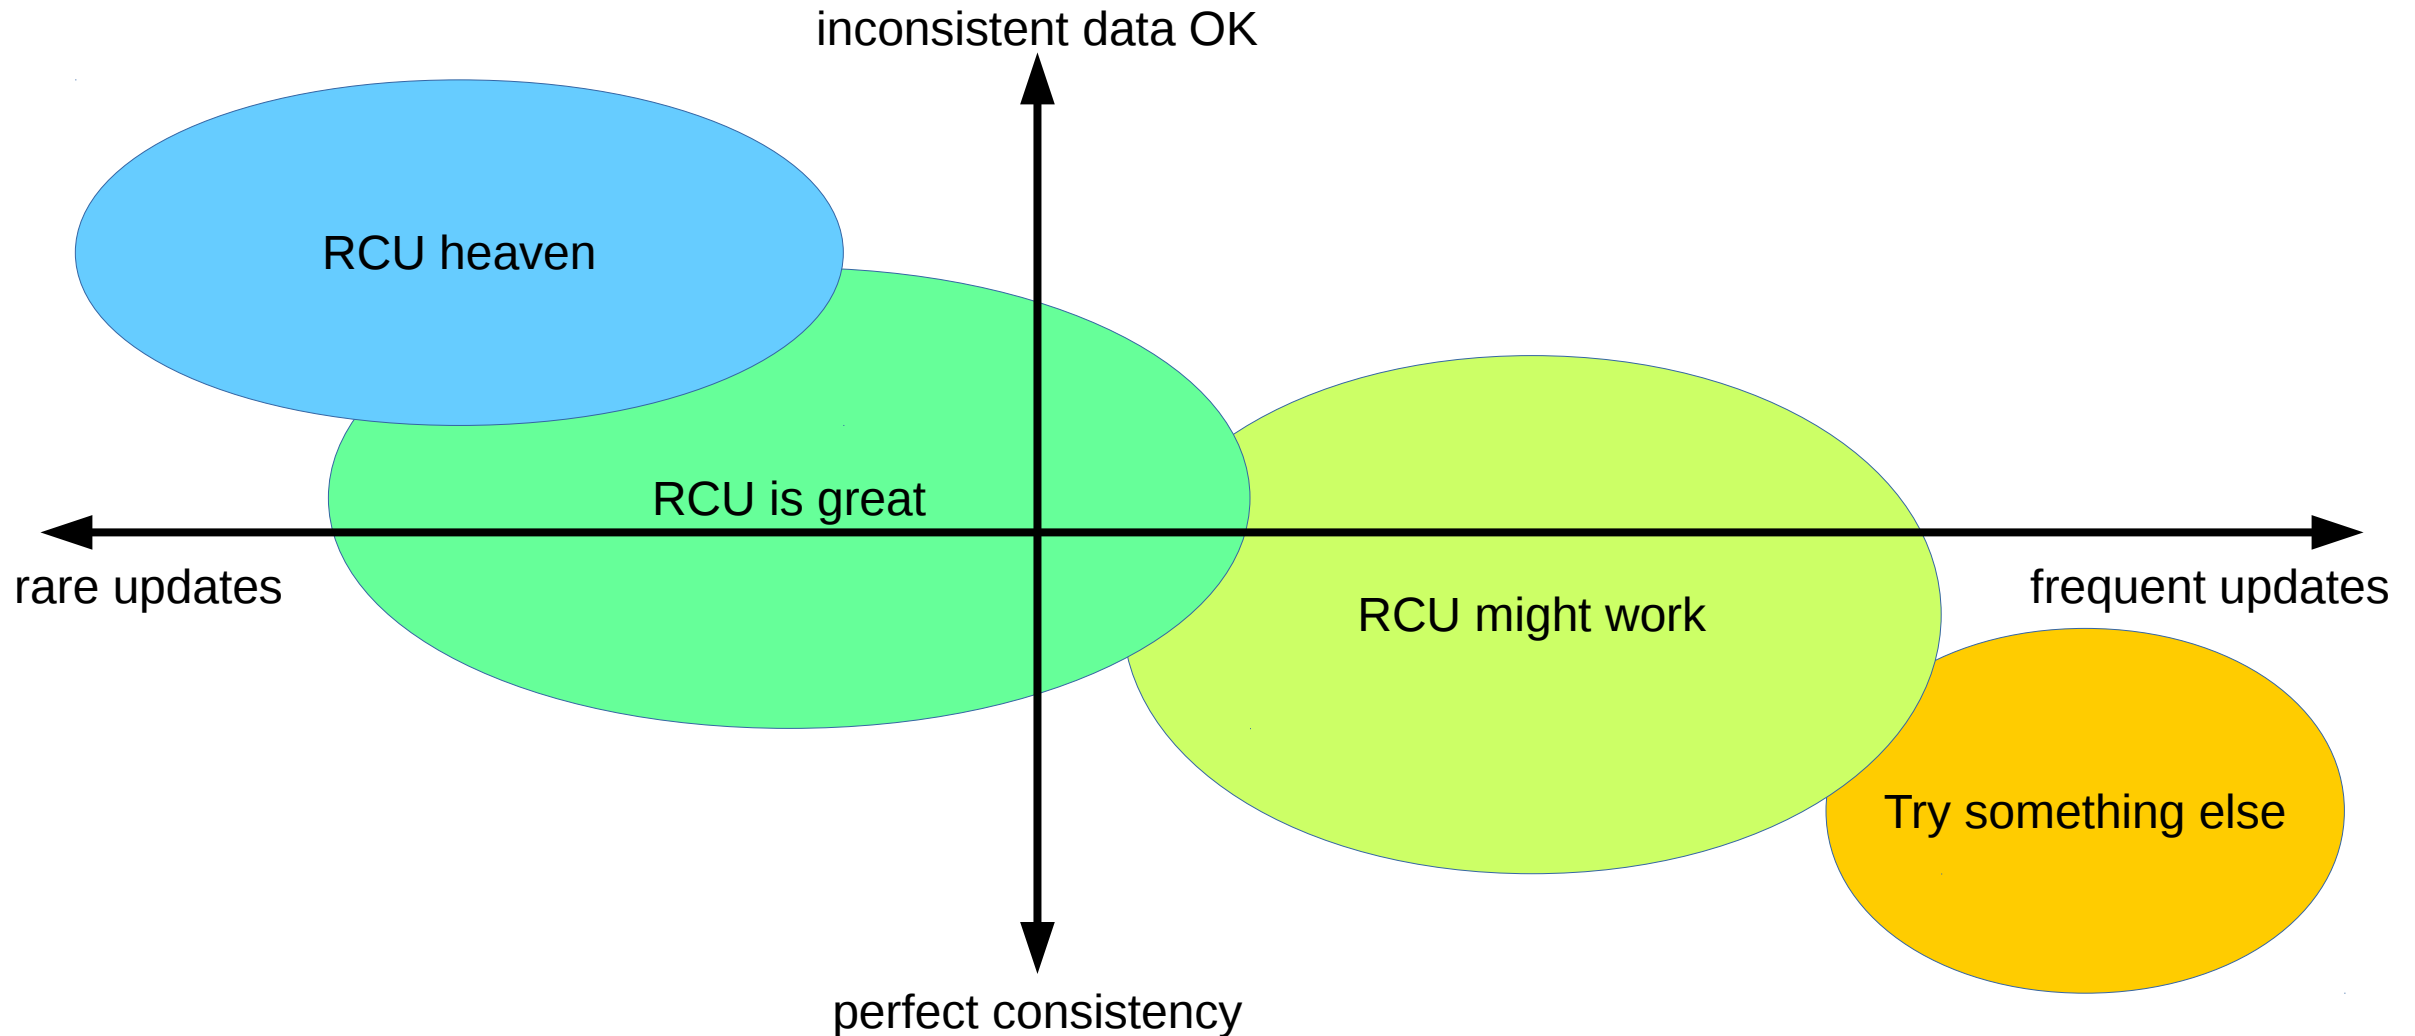
RCU implementations: writers and readers

- Reader can hold a reference to the shared data forever
 - Writer can avoid blocking waits but cannot reclaim data
- In some applications writers can force-reclaim (after a grace period) – readers will be left in an undefined state

RCU implementations: mean writers and meek (or just fault-tolerant) readers

- Memory is reused but not freed, readers can handle inconsistent data – RCU paradise
- Memory is reused, readers need consistent data
`bool rcu_read_unlock(handle_t handle);`
 - Returns false if the handle has expired
 - Everything done in this critical section must be redone
- Readers can crash and it's OK
 - Separate reader processes, stateless, easy to restart

When to use RCU



RCU and alternatives

	RCU	Hazard pointers	Atomic shared_ptr
Readers	wait-free population-oblivious	lock-free (wait-free?)	lock-free (slow)
Writers	single writer (BYOL)	lock-free	lock-free
Reclamation	blocking (or memory grows)	non-blocking	lock-free
Garbage	unbounded (or writers must block)	bounded by $N_{\text{threads}} * N_{\text{HP}}$	None
Ease of use	very easy	hard	easy (watch out for cycles)

```
rcu_questions_unlock();  
do { answer_questions() }  
while ( !wait_for_readers_to_leave(); )
```



www.mentor.com