

RUNTIME POLYMORPHISM: BACK TO THE BASICS

LOUIS DIONNE, CPPCON 2017

WHAT IS RUNTIME POLYMORPHISM AND WHEN DO YOU NEED IT?

CONSIDER THE FOLLOWING

```
struct Car {  
    void accelerate();  
};
```

```
struct Truck {  
    void accelerate();  
};
```

```
struct Plane {  
    void accelerate();  
};
```

RETURNING RELATED TYPES FROM A FUNCTION

```
??? getVehicle(std::istream& user) {  
    std::string choice;  
    user >> choice;  
    if      (choice == "car")    return Car{...};  
    else if (choice == "truck") return Truck{...};  
    else if (choice == "plane") return Plane{...};  
    else                          die();  
}
```

STORING RELATED TYPES IN A CONTAINER

```
int main() {  
    // Should store anything that has an accelerate() method  
    std::vector<???> vehicles;  
  
    vehicles.push_back(Car{...});  
    vehicles.push_back(Truck{...});  
    vehicles.push_back(Plane{...});  
  
    for (auto& vehicle : vehicles) {  
        vehicle.accelerate();  
    }  
}
```

variant SOMETIMES DOES THE TRICK

- But it only works for closed set of types
- Using visitation is sometimes (often?) not convenient

BOTTOM LINE:
MANIPULATING AN OPEN SET OF RELATED TYPES
WITH DIFFERENT REPRESENTATIONS

C++ HAS A SOLUTION FOR THAT!

INHERITANCE

```
struct Vehicle {  
    virtual void accelerate() = 0;  
    virtual ~Vehicle() { }  
};
```

```
struct Car : Vehicle {  
    void accelerate() override;  
};
```

```
struct Truck : Vehicle {  
    void accelerate() override;  
};
```

```
struct Plane : Vehicle {  
    void accelerate() override;  
};
```

UNDER THE HOOD

```
Vehicle* ptr;
```

Car:

```
__vtable* __vptr;  
string make;  
int year;  
...
```

Car virtual table:

```
void (*accelerate)(Vehicle* __this);  
void (*__dtor)(Vehicle* __this);  
...
```

ASIDE

INHERITANCE HAS MANY PROBLEMS

BAKES IN REFERENCE SEMANTICS

```
void foo(Vehicle* vehicle) {  
    Vehicle* copy = vehicle;  
    ...  
    copy->accelerate();  
    ...  
}
```

HEAP ALLOCATIONS

```
std::unique_ptr<Vehicle> getVehicle(std::istream& user) {  
    std::string choice;  
    user >> choice;  
    if      (choice == "car")    return std::make_unique<Car>(...);  
    else if (choice == "truck") return std::make_unique<Truck>(...);  
    else if (choice == "plane") return std::make_unique<Plane>(...);  
    else                        die();  
}
```

BAKES IN NULLABLE SEMANTICS

```
std::unique_ptr<Vehicle> vehicle = getVehicle(std::cin);  
// can vehicle be null?
```

OWNERSHIP HELL

```
Vehicle*           getVehicle(std::istream& user);  
std::unique_ptr<Vehicle> getVehicle(std::istream& user);  
std::shared_ptr<Vehicle> getVehicle(std::istream& user);
```

DOESN'T PLAY WELL WITH ALGORITHMS

```
std::vector<std::unique_ptr<Vehicle>> vehicles;  
vehicles.push_back(std::make_unique<Car>(...));  
vehicles.push_back(std::make_unique<Truck>(...));  
vehicles.push_back(std::make_unique<Plane>(...));  
  
std::sort(vehicles.begin(), vehicles.end()); // NOT what you wanted!
```


INTRUSIVE

```
namespace lib {  
    struct Motorcycle { void accelerate(); };  
}  
  
void foo(Vehicle& vehicle) {  
    ...  
    vehicle.accelerate();  
    ...  
}  
  
Motorcycle bike;  
foo(bike); // can't work!
```

LISTEN TO SEAN PARENT, NOT ME

<https://youtu.be/QGcVXgEVMJg>

I JUST WANTED THIS!

```
interface Vehicle { void accelerate(); };

namespace lib {
    struct Motorcycle { void accelerate(); };
}
struct Car { void accelerate(); };
struct Truck { void accelerate(); };

int main() {
    std::vector<Vehicle> vehicles;
    vehicles.push_back(Car{...});
    vehicles.push_back(Truck{...});
    vehicles.push_back(lib::Motorcycle{...});

    for (auto& vehicle : vehicles) {
        vehicle.accelerate();
    }
}
```

HOW MIGHT THAT WORK?

WITH INHERITANCE

```
Vehicle* ptr;
```

Car:

```
__vtable* __vptr;  
string make;  
int year;  
...
```

Car virtual table:

```
void (*accelerate)(Vehicle* __this);  
void (*__dtor)(Vehicle* __this);  
...
```

GOAL:

INDEPENDENT STORAGE AND METHOD DISPATCH

- Storage *policy*
- VTable *policy*

REMOTE STORAGE

Vehicle:

```
vtable const* vptr_;  
void* ptr_;
```

Car "virtual table":

```
void (*accelerate)(void*);  
void (*delete_)(void*);  
...
```

Car:

```
string make;  
int year;  
...
```

HOW THAT'S IMPLEMENTED

```
class Vehicle {
    vtable const* const vptr_;
    void* ptr_;

public:
    template <typename Any>
        // enabled only when vehicle.accelerate() is valid
    Vehicle(Any vehicle)
        : vptr_{&vtable_for<Any>}
        , ptr_{new Any(vehicle)}
    { }

    Vehicle(Vehicle const& other); // implementation omitted

    void accelerate()
    { vptr_->accelerate(ptr_); }

    ~Vehicle()
    { vptr_->delete_(ptr_); }
};
```


THE VTABLE

```
struct vtable {  
    void (*accelerate)(void* this_);  
    void (*delete_)(void* this_);  
};  
  
template <typename T>  
vtable const vtable_for = {  
    [](void* this_) {  
        static_cast<T*>(this_)->accelerate();  
    },  
  
    [](void* this_) {  
        delete static_cast<T*>(this_);  
    }  
};
```

WITH DYNO

```
struct Vehicle {  
    template <typename Any>  
    Vehicle(Any vehicle) : poly_{vehicle} { }  
  
    void accelerate()  
    { poly_.virtual_("accelerate"_s)(poly_); }  
  
private:  
    dyno::poly<IVehicle, dyno::remote_storage> poly_;  
    // ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^  
};
```

DYNO'S VTABLE

```
struct IVehicle : decltype(dyno::requires(
    dyno::CopyConstructible{},
    dyno::Destructible{},
    "accelerate"_s = dyno::function<void(dyno::T&)>
)) { };

template <typename T>
auto dyno::default_concept_map<IVehicle, T> = dyno::make_concept_map(
    "accelerate"_s = [](T& vehicle) { vehicle.accelerate(); }
);
```

STRENGTHS AND WEAKNESSES

- ✓ Simple model, similar to classic inheritance
- ✗ Always requires an allocation

THE *SMALL BUFFER OPTIMIZATION* (SBO)

Vehicle:

```
vtable const* vptr_;  
bool on_heap_;  
union {  
    void* ptr_;  
    char buffer_[N] {  
        Car:  
        string make;  
        int year;  
        ...  
    }  
};
```

Car "virtual table":

```
void (*accelerate)(void*);  
void (*delete_)(void*);  
void (*dtor)(void*);  
...
```

Car:

```
string make;  
int year;  
...
```

HOW THAT'S IMPLEMENTED

```
struct Vehicle {
    vtable const* const vptr_;
    union { void* ptr_;
           std::aligned_storage_t<16> buffer_; };
    bool on_heap_;

    template <typename Any>
    Vehicle(Any vehicle) : vptr_{&vtable_for<Any>} {
        if (sizeof(Any) > 16) {
            on_heap_ = true;
            ptr_ = new Any(vehicle);
        } else {
            on_heap_ = false;
            new (&buffer_) Any{vehicle};
        }
    }

    void accelerate()
    { vptr_->accelerate(on_heap_ ? ptr_ : &buffer_); }
};
```

ALTERNATIVE IMPLEMENTATION 1

Vehicle:

```
vtable const* vptr_;  
union {  
    void* ptr_;  
    char buffer_[N] {  
        Car:  
        string make;  
        int year;  
        ...  
    }  
};
```

Car "virtual table":

```
bool on_heap;  
void (*accelerate)(void*);  
void (*delete_)(void*);  
void (*dtor)(void*);  
...
```

Car:

```
string make;  
int year;  
...
```


ALTERNATIVE IMPLEMENTATION 2

(seems to be the fastest)

Vehicle:

`vtable const* vptr_;`

`void* storage_;`

`char buffer_[N] {`

Car:

`string make;`

`int year;`

`...`

`}`

Car "virtual table":

`void (*accelerate)(void*);`

`void (*delete_)(void*);`

`void (*dtor)(void*);`

`...`

Car:

`string make;`

`int year;`

`...`

WITH DYNO

```
struct Vehicle {  
    template <typename Any>  
    Vehicle(Any vehicle) : poly_{vehicle} { }  
  
    void accelerate()  
    { poly_.virtual_("accelerate"_s)(poly_); }  
  
private:  
    dyno::poly<IVehicle, dyno::sbo_storage<16>> poly_;  
    //  
};
```

STRENGTHS AND WEAKNESSES

- ✓ Does not always require allocating
- ✗ Takes up more space
- ✗ Copy/move/swap is more complicated
- ✗ Dispatching may be more costly

ALWAYS-LOCAL STORAGE

Vehicle:

```
vtable const* vptr_;  
char buffer_[N] {
```

Car:

```
string make;  
int year;  
...
```

```
}
```

Car "virtual table":

```
void (*accelerate)(void*);  
void (*dtor)(void*);  
...
```

DOESN'T FIT? DOESN'T COMPILE!

HOW THAT'S IMPLEMENTED

```
class Vehicle {
    vtable const* const vptr_;
    std::aligned_storage_t<64> buffer_;

public:
    template <typename Any>
    Vehicle(Any vehicle) : vptr_{&vtable_for<Any>} {
        static_assert(sizeof(Any) <= sizeof(buffer_),
            "can't hold such a large object in a Vehicle");
        new (&buffer_) Any(vehicle);
    }

    void accelerate()
    { vptr_->accelerate(&buffer_); }

    ~Vehicle()
    { vptr_->dtor(&buffer_); }
};
```

WITH DYNO

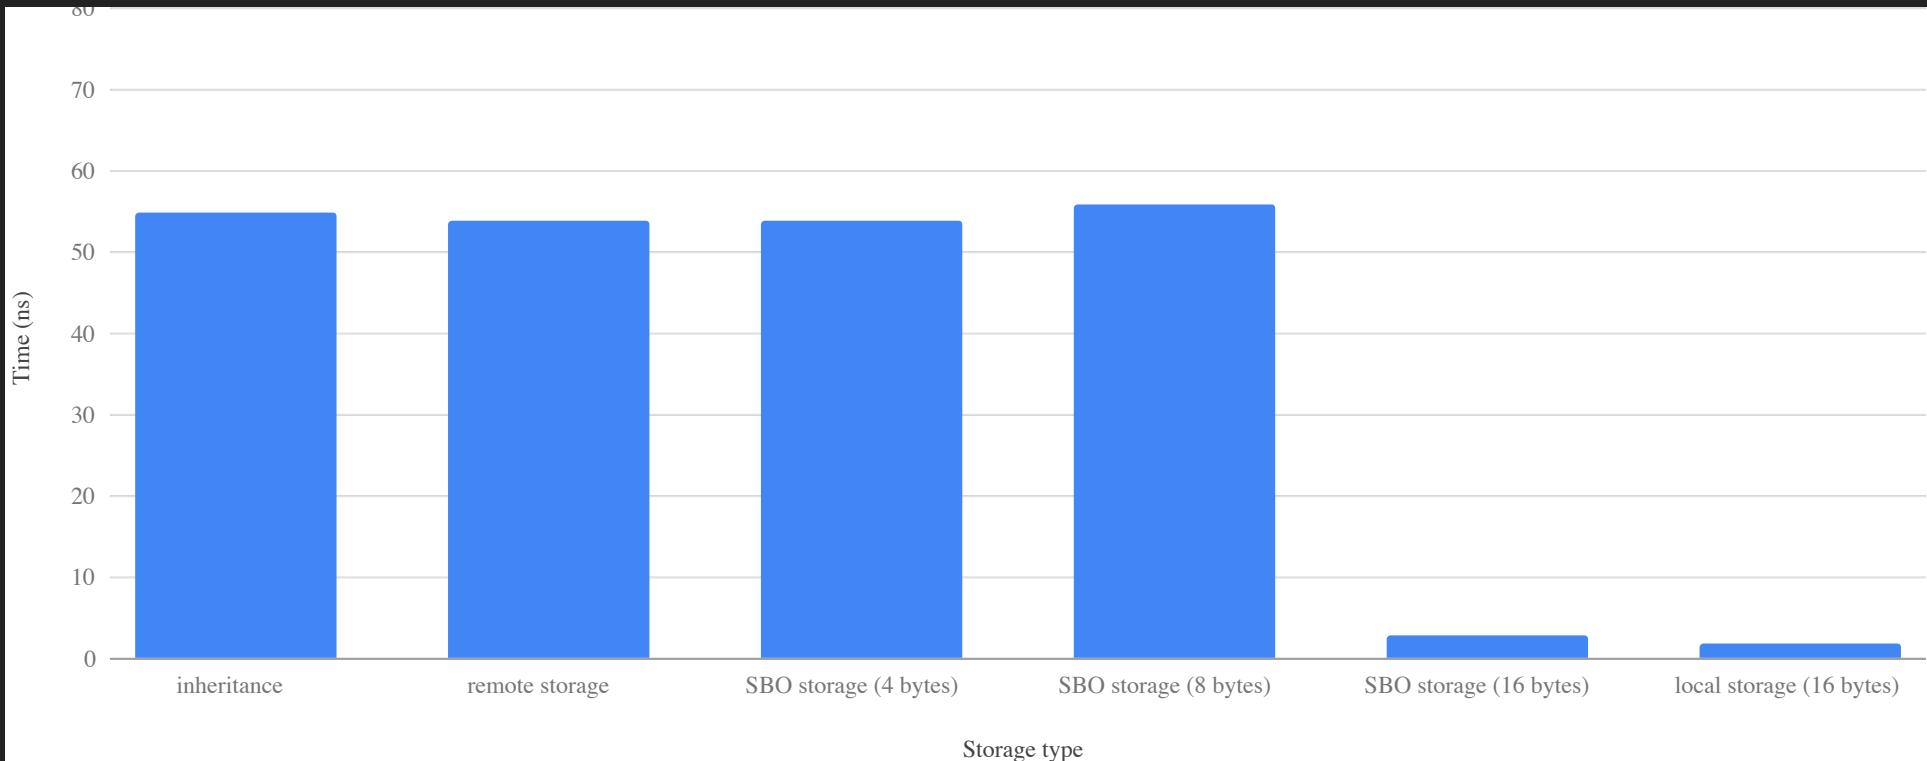
```
struct Vehicle {  
    template <typename Any>  
    Vehicle(Any vehicle) : poly_{vehicle} { }  
  
    void accelerate()  
    { poly_.virtual_("accelerate"_s)(poly_); }  
  
private:  
    dyno::poly<IVehicle, dyno::local_storage<64>> poly_;  
    //  
};
```

STRENGTHS AND WEAKNESSES

- ✓ No allocation – ever
- ✓ Simple dispatching
- ✗ Takes up more space

SOME BENCHMARKS

Creating many 16 bytes objects



Accessing many 4 bytes objects (10 x 3 method calls, SBO with pointer)

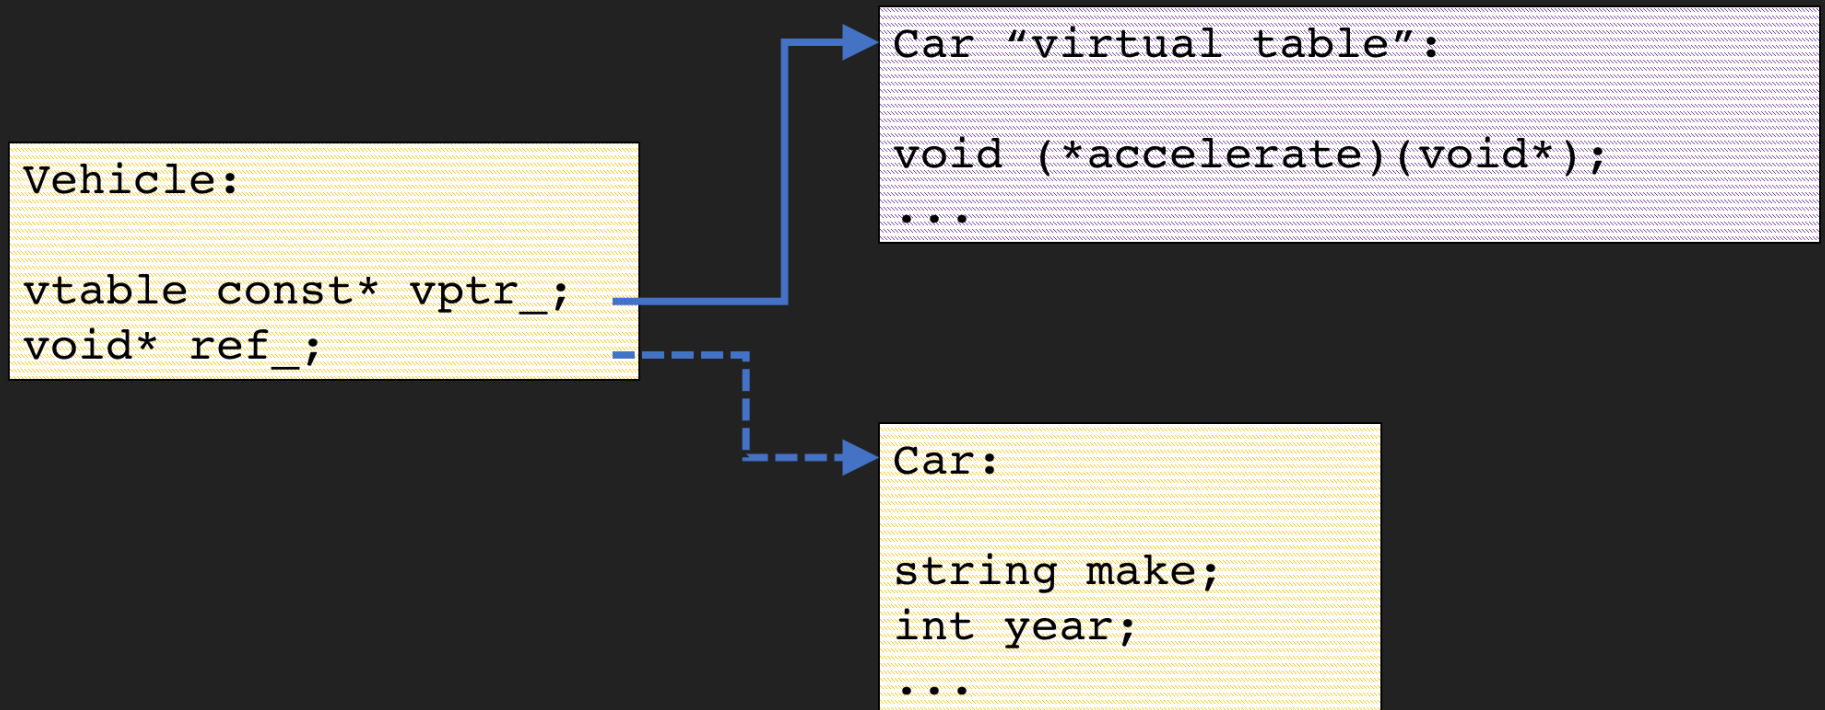


GUIDELINES

- Use local storage whenever you can afford it
- Otherwise, use SBO with the largest reasonable size
- Use purely-remote storage only when
 - Object sizes are so scattered SBO wouldn't help

NON-OWNING STORAGE

(reference semantics, not value semantics)



BASICALLY A POLYMORPHIC VIEW

```
void process(Vehicle vehicle) {  
    ...  
    vehicle.accelerate();  
    ...  
}  
  
int main() {  
    Truck truck{...};  
    process(truck); // No copy!  
}
```

SHARED REMOTE STORAGE

Car:

```
string make;  
int year;  
...
```

Car "virtual table":


```
void (*accelerate)(void*);  
...
```

Vehicle:

```
vtable const* vptr_  
shared_ptr<void> ptr_;
```

Vehicle:

```
vtable const* vptr_  
shared_ptr<void> ptr_;
```



NOW, LET ME SHOW YOU WHY YOU CARE

HAVE YOU HEARD OF THE FOLLOWING?

- `std::function`
- `inplace_function`
- `function_view`

CONSIDER THIS

```
template <typename Signature, typename StoragePolicy>
struct basic_function;

template <typename R, typename ...Args, typename StoragePolicy>
struct basic_function<R(Args...), StoragePolicy> {
    template <typename F>
    basic_function(F&& f) : poly_{std::forward<F>(f)} { }

    R operator()(Args ...args) const
    { return poly_.virtual_("call"_s)(poly_, args...); }

private:
    dyno::poly<Callable<R(Args...)>, StoragePolicy> poly_;
};
```


HERE'S ALL OF THEM:

```
template <typename Signature>
using function = basic_function<Signature,
                               dyno::sbo_storage<16>>>;
```

```
template <typename Signature, std::size_t Size = 32>
using inplace_function = basic_function<Signature,
                                         dyno::local_storage<Size>>>;
```

```
template <typename Signature>
using function_view = basic_function<Signature,
                                     dyno::non_owning_storage>;
```

```
template <typename Signature>
using shared_function = basic_function<Signature,
                                       dyno::shared_remote_storage>;
```

**WE'VE TALKED ABOUT STORAGE
WHAT ABOUT VTABLES?**

NORMALLY, IT IS REMOTE

Vehicle:

```
vtable const* vptr_  
... storage ...
```

Car "virtual table":

```
void (*accelerate)(void*);  
void (*dtor)(void*);  
...
```

TURNS OUT WE HAVE SOME CHOICES

INLINING THE VTABLE IN THE OBJECT

Vehicle:

```
vtbl vtbl_ {
```

```
    Car "virtual table":
```

```
    void (*accelerate)(void*);
```

```
    void (*dtor)(void*);
```

```
    ...
```

```
}
```

```
... storage ...
```

HOW THAT'S IMPLEMENTED

```
struct Vehicle {  
    template <typename Any>  
    Vehicle(Any vehicle)  
        : vtbl_{vtable_for<Any>}  
        , ptr_{new Any(vehicle)}  
    { }  
  
    void accelerate()  
    { vtbl_.accelerate(ptr_); }  
  
    ~Vehicle()  
    { vtbl_.delete_(ptr_); }  
  
private:  
    vtable const vtbl_; // <= not a pointer!  
    void* ptr_;  
};
```

WITH DYNO

```
struct Vehicle {  
    template <typename Any>  
    Vehicle(Any vehicle) : poly_{vehicle} { }  
  
    void accelerate()  
    { poly_.virtual_("accelerate"_s)(poly_); }  
  
private:  
    using VTable = dyno::vtable<dyno::local<dyno::everything>>;  
    // ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^  
    dyno::poly<IVehicle, dyno::remote_storage, VTable> poly_;  
};
```


USUALLY A PESSIMIZATION

PARTIAL VTABLE INLINING

Vehicle:

```
hybrid_vtable vtbl_ {
```

```
    vtable const* remote;  
    void (*accelerate)(void*);
```



```
}
```

```
... storage ...
```

Car "virtual table":

```
void (*delete_)(void*);  
...
```

THE VTABLE — REMOTE PART

```
struct vtable {  
    void (*delete_)(void* this_);  
    // ...  
};  
  
template <typename T>  
vtable const vtable_for = {  
    [](void* this_) {  
        delete static_cast<T*>(this_);  
    }  
    // ...  
};
```

THE VTABLE — LOCAL PART

```
struct joined_vtable {  
    vtable const* const remote;  
    void (*accelerate)(void* this_);  
};  
  
template <typename T>  
joined_vtable const joined_vtable_for = {  
    &vtable_for<T>,  
    [](void* this_) {  
        static_cast<T*>(this_)->accelerate();  
    }  
};
```

THE POLYMORPHIC WRAPPER

```
class Vehicle {
    joined_vtable const vtbl_;
    void* ptr_;

public:
    template <typename Any>
    Vehicle(Any vehicle)
        : vtbl_{joined_vtable_for<Any>}
        , ptr_{new Any(vehicle)}
    { }

    void accelerate()
    { vtbl_.accelerate(ptr_); }

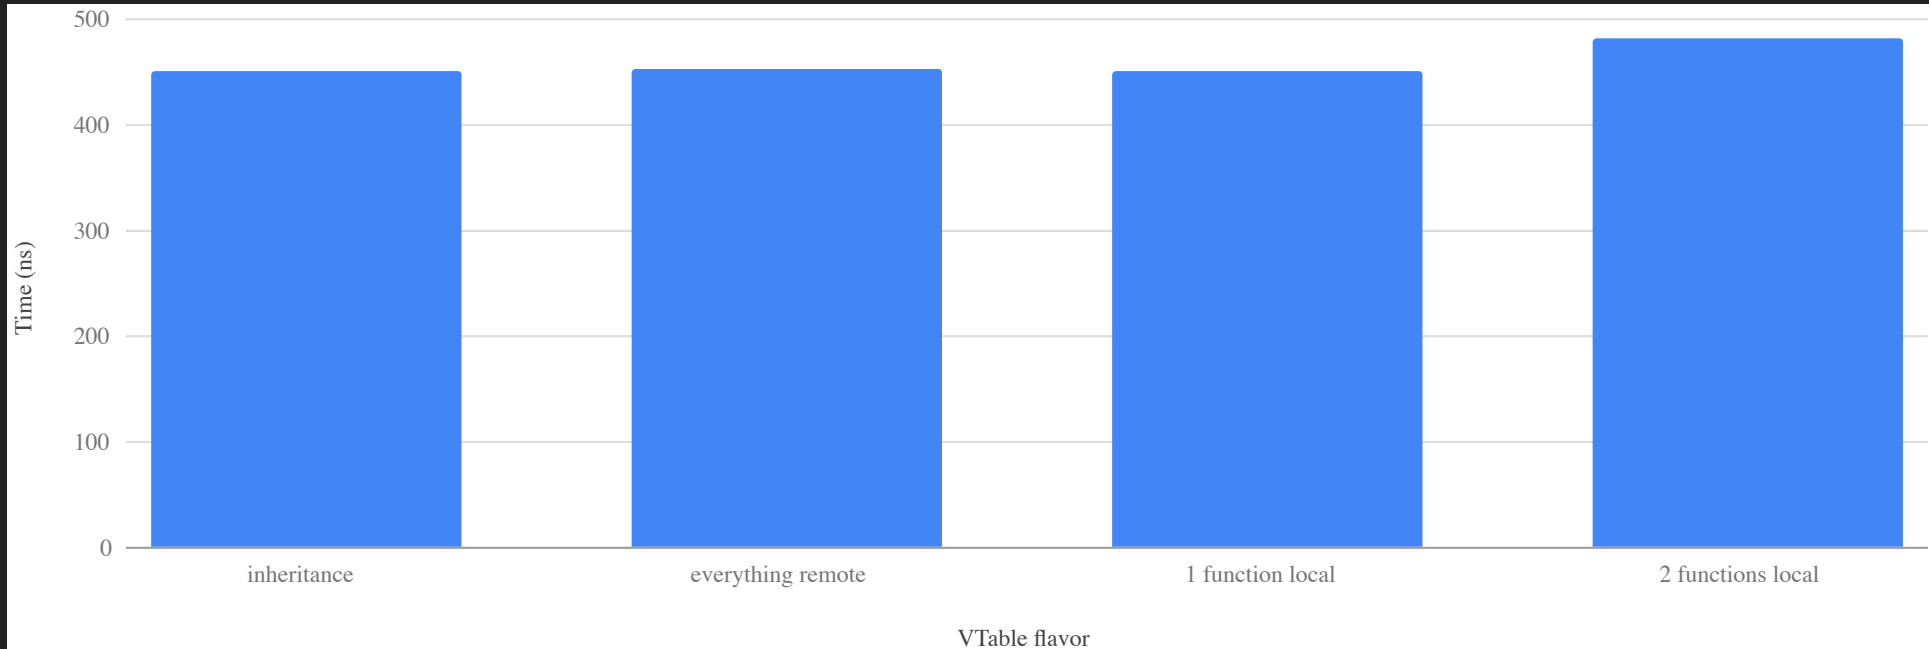
    ~Vehicle()
    { vtbl_.remote->delete_(ptr_); }
};
```

WITH DYNO

```
struct Vehicle {  
    template <typename Any>  
        Vehicle(Any vehicle) : poly_{vehicle} { }  
  
    void accelerate()  
    { poly_.virtual_("accelerate"_s)(poly_); }  
  
private:  
    using VTable = dyno::vtable<  
        dyno::local<dyno::only<decltype("accelerate"_s)>>,  
            dyno::remote<dyno::everything_else>>;  
        // ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^  
    dyno::poly<IVehicle, dyno::remote_storage, VTable> poly_;  
};
```

SOME BENCHMARKS

Calling 2 virtual functions (100 x 2 calls)



NOT VERY CONCLUSIVE
LET'S LOOK AT ASSEMBLY

A▼



```
1 struct VTable {
2     void (*f1)(void*);
3     void (*f2)(void*);
4     void (*f3)(void*);
5     void (*f4)(void*);
6 };
7
8 template <typename T>
9 extern VTable const vtable;
10
11
12 struct remote_any {
13     void f1() { vptr_->f1(self_); }
14     void f2() { vptr_->f2(self_); }
15     void f3() { vptr_->f3(self_); }
16     void f4() { vptr_->f4(self_); }
17     VTable const* const vptr_;
18     void* self_;
19 };
20
21 struct local_any {
22     void f1() { vtbl_.f1(self_); }
23     void f2() { vtbl_.f2(self_); }
24     void f3() { vtbl_.f3(self_); }
25     void f4() { vtbl_.f4(self_); }
26     VTable const vtbl_;
27     void* self_;
28 };
29
```

x86-64 clang 5.0.0 ▼

-O3 -std=c++14

11010

.LX0:

.text

//

\s+

Intel

Demangle

A▼



1

Edit on C++ Compiler Explorer

(1)


A STORY OF INLINING

```
template <typename AnyIterator, typename It>
__attribute__((noinline)) AnyIterator make(It it) {
    return AnyIterator{std::move(it)};
}

template <typename AnyIterator>
void benchmark_any_iterator(benchmark::State& state) {
    std::vector<int> input{...};
    std::vector<int> output{...};


    while (state.KeepRunning()) {
        auto first = make<AnyIterator>(input.begin());
        auto last = make<AnyIterator>(input.end());
        auto result = make<AnyIterator>(output.begin());

        for (; !(first == last); ++first, ++result) {
            *result = *first;
        }
    }
}
```



NOW, JUST A SMALL TWEAK

```
template <typename AnyIterator, typename It>
// __attribute__((noinline))
AnyIterator make(It it) {
    return AnyIterator{std::move(it)};
}
```



WHAT HAPPENED?

Inheritance:

```
Vehicle* ptr;
```

Car:

```
__vtable* __vptr;  
string make;  
int year;  
...
```

Car virtual table:

```
void (*accelerate)(Vehicle* __this);  
void (*__dtor)(Vehicle* __this);  
...
```

Dyno's remote vtable:

Vehicle:

```
vtable const* vptr_  
... storage ...
```

Car "virtual table":

```
void (*accelerate)(void*);  
void (*dtor)(void*);  
...
```

WHAT'S THE LESSON?

- Reducing pointer hops can lead to unexpected inlining
- When that happens, giant optimizations become possible

GUIDELINES

- By default, all methods are in the remote vtable
- Consider inlining some methods if
 - you have slack space
 - you know you're calling them often
 - you've measured it makes a difference

THE FUTURE WITH REFLECTION?

```
struct Vehicle {  
    void accelerate();  
};  
  
int main() {  
    std::vector<poly<Vehicle>> vehicles;  
    vehicles.push_back(Car{...});  
    vehicles.push_back(Truck{...});  
    vehicles.push_back(lib::Motorcycle{...});  
  
    for (auto& vehicle : vehicles) {  
        vehicle.accelerate();  
    }  
}
```


THE FUTURE WITH METACLASSES?

```
interface Vehicle {  
    void accelerate();  
};  
  
int main() {  
    std::vector<Vehicle> vehicles;  
    vehicles.push_back(Car{...});  
    vehicles.push_back(Truck{...});  
    vehicles.push_back(lib::Motorcycle{...});  
  
    for (auto& vehicle : vehicles) {  
        vehicle.accelerate();  
    }  
}
```

SUMMARY

- Inheritance model is just one option amongst others
 - Don't bake that choice in
- Many ways of storing polymorphic objects
 - As always, space/time tradeoff
- Vtables can be inlined (measure!)
- Type erasure is tedious to do manually
 - Reflection will be there to help

THE DYNO LIBRARY IS AVAILABLE

<https://github.com/ldionne/dyno>

USEFUL LINKS AND RELATED MATERIAL

- Sean Parent's NDC 2017 talk:
<https://youtu.be/QGcVXgEVMJg>
- Zach Laine's CppCon 2014 talk:
<https://youtu.be/0l0FD3N5cgM>
- Boost.TypeErasure:
http://www.boost.org/doc/libs/release/doc/html/boost_typeerasure.html
- Adobe Poly:
https://stlab.adobe.com/group__poly__related.html
- Eraserface:
<https://github.com/badair/eraserface>
- liberasure:
<https://github.com/atomgalaxy/liberasure>
- 2004 thread on interfaces:
<https://goo.gl/zaBN6X>

THANK YOU

<https://ldionne.com>