

# UNDEFINED BEHAVIOR IS AWESOME

---

Piotr Padlewski

[piotr.padlewski@gmail.com](mailto:piotr.padlewski@gmail.com), @PiotrPadlewski







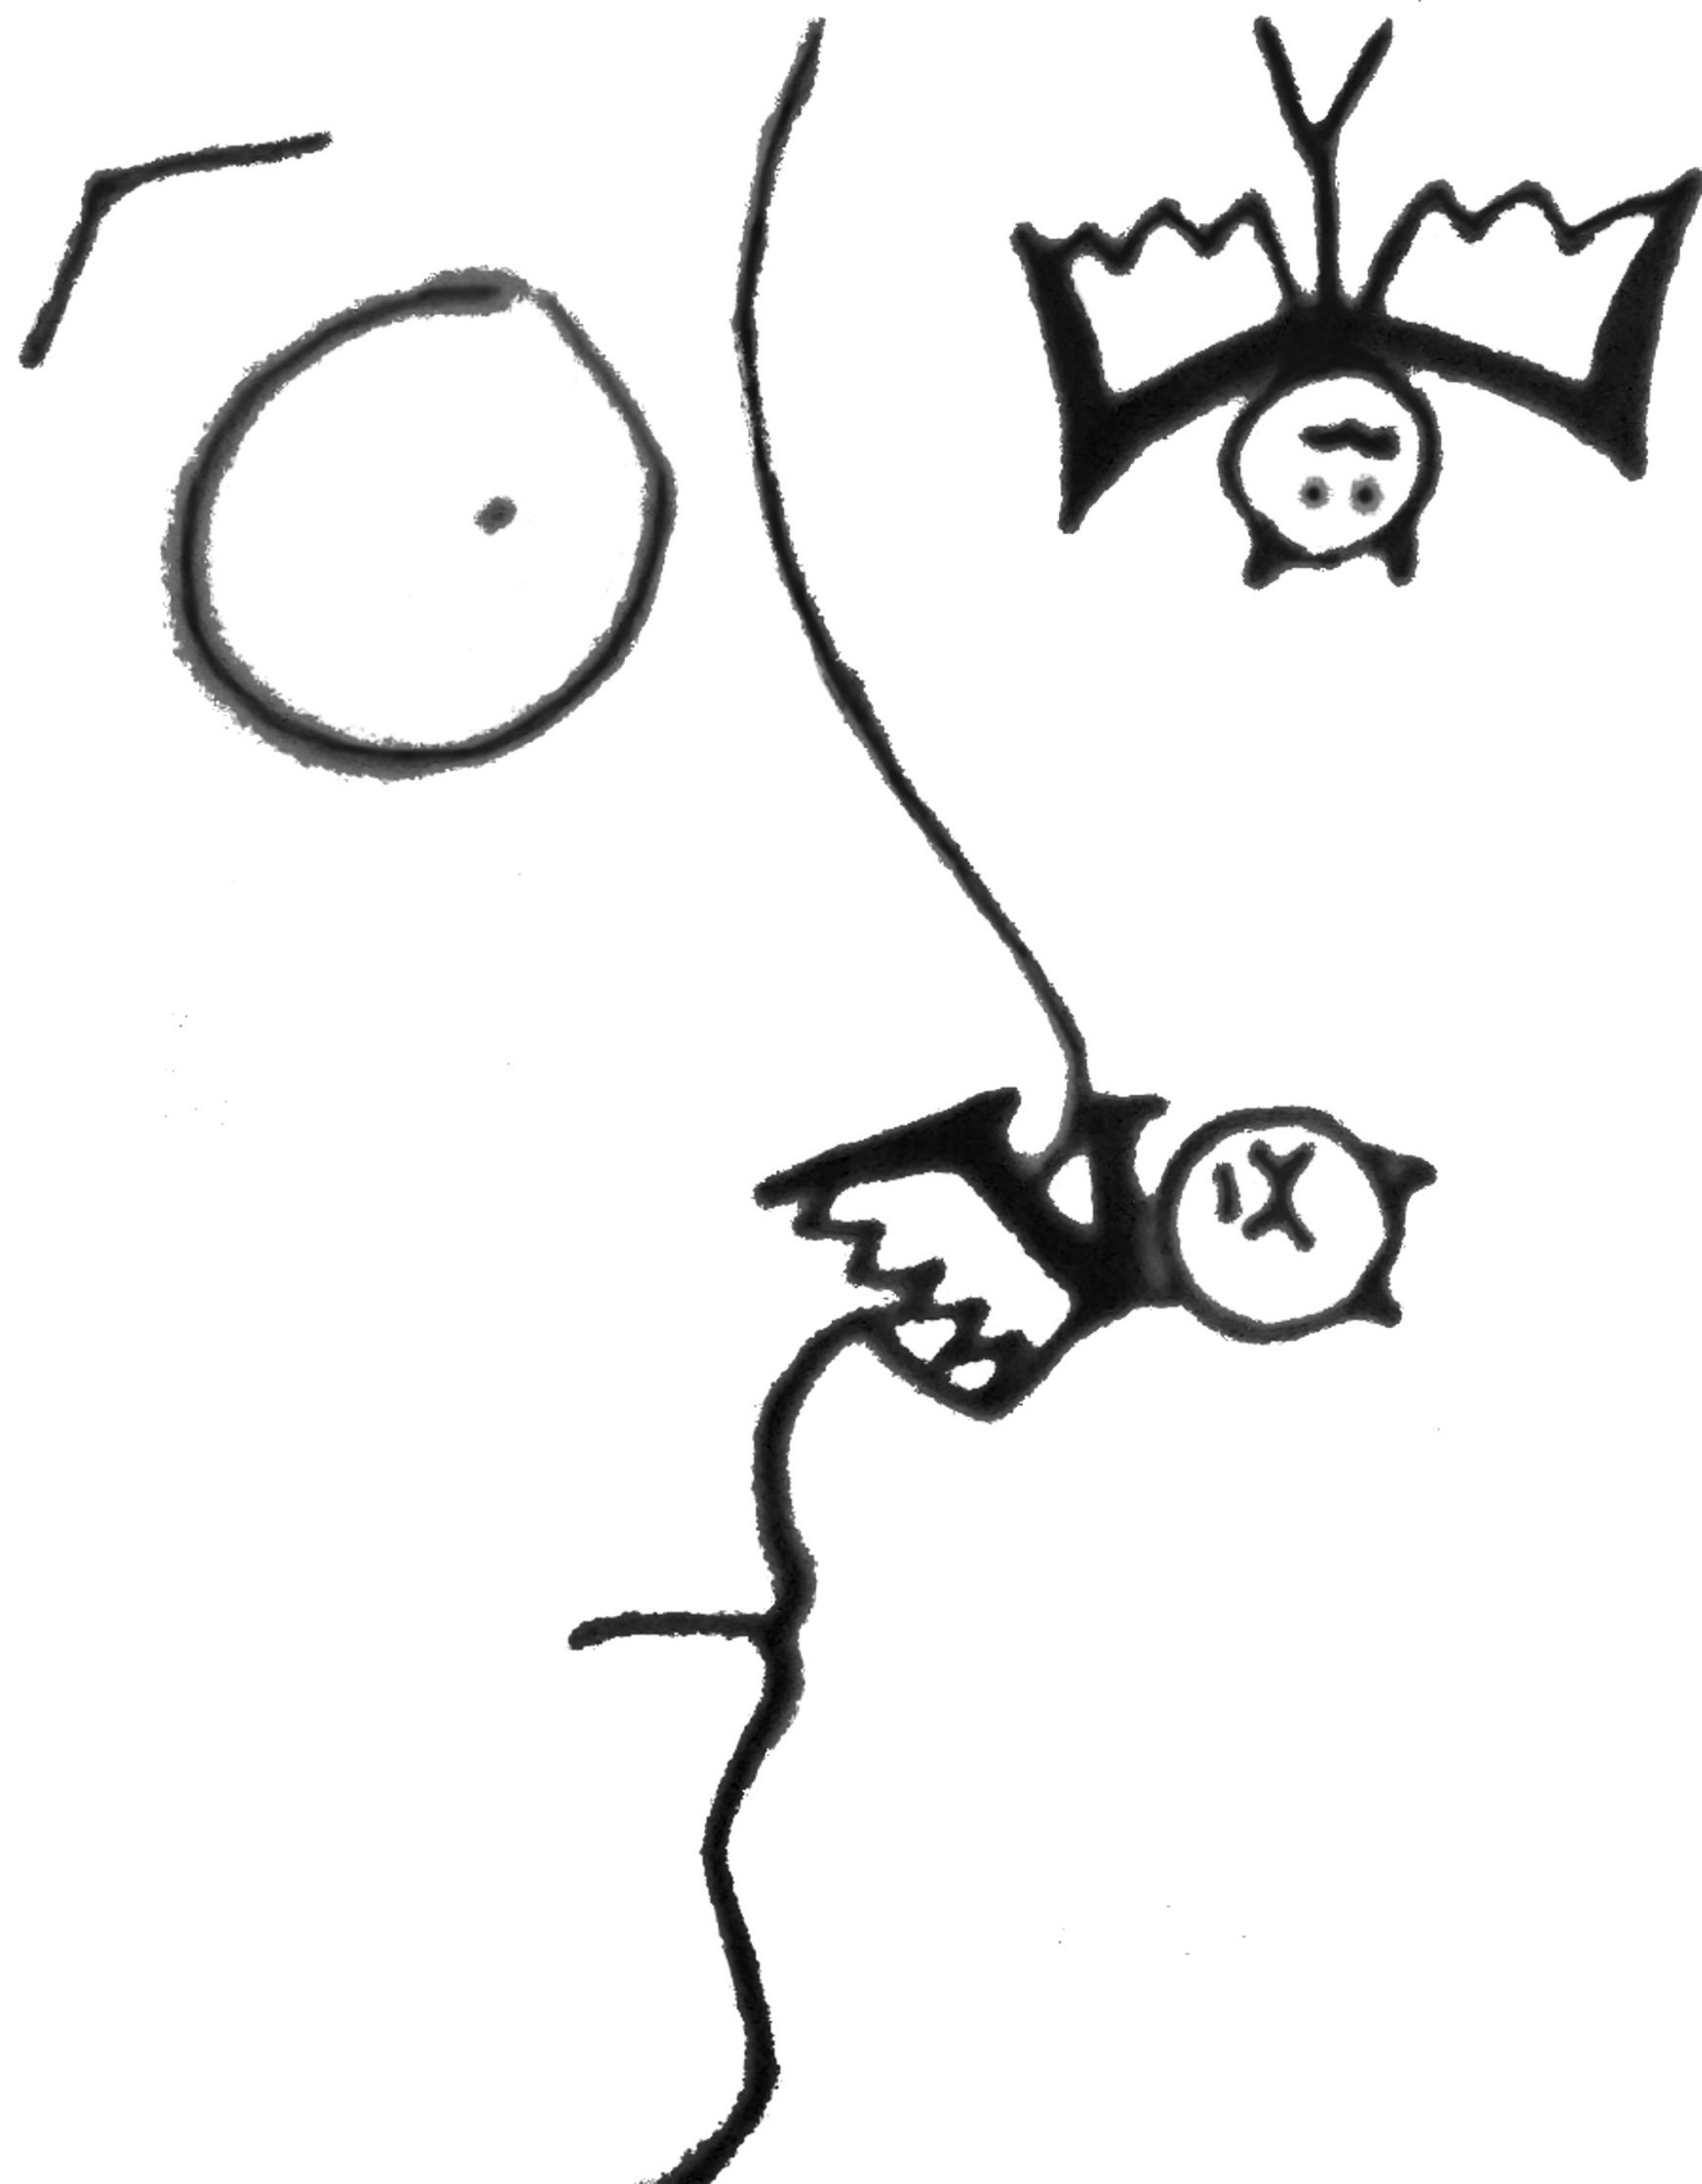
## OUTLINE

- ▶ What is UB
- ▶ Why it sucks
- ▶ How to fight with it
- ▶ Why we need it

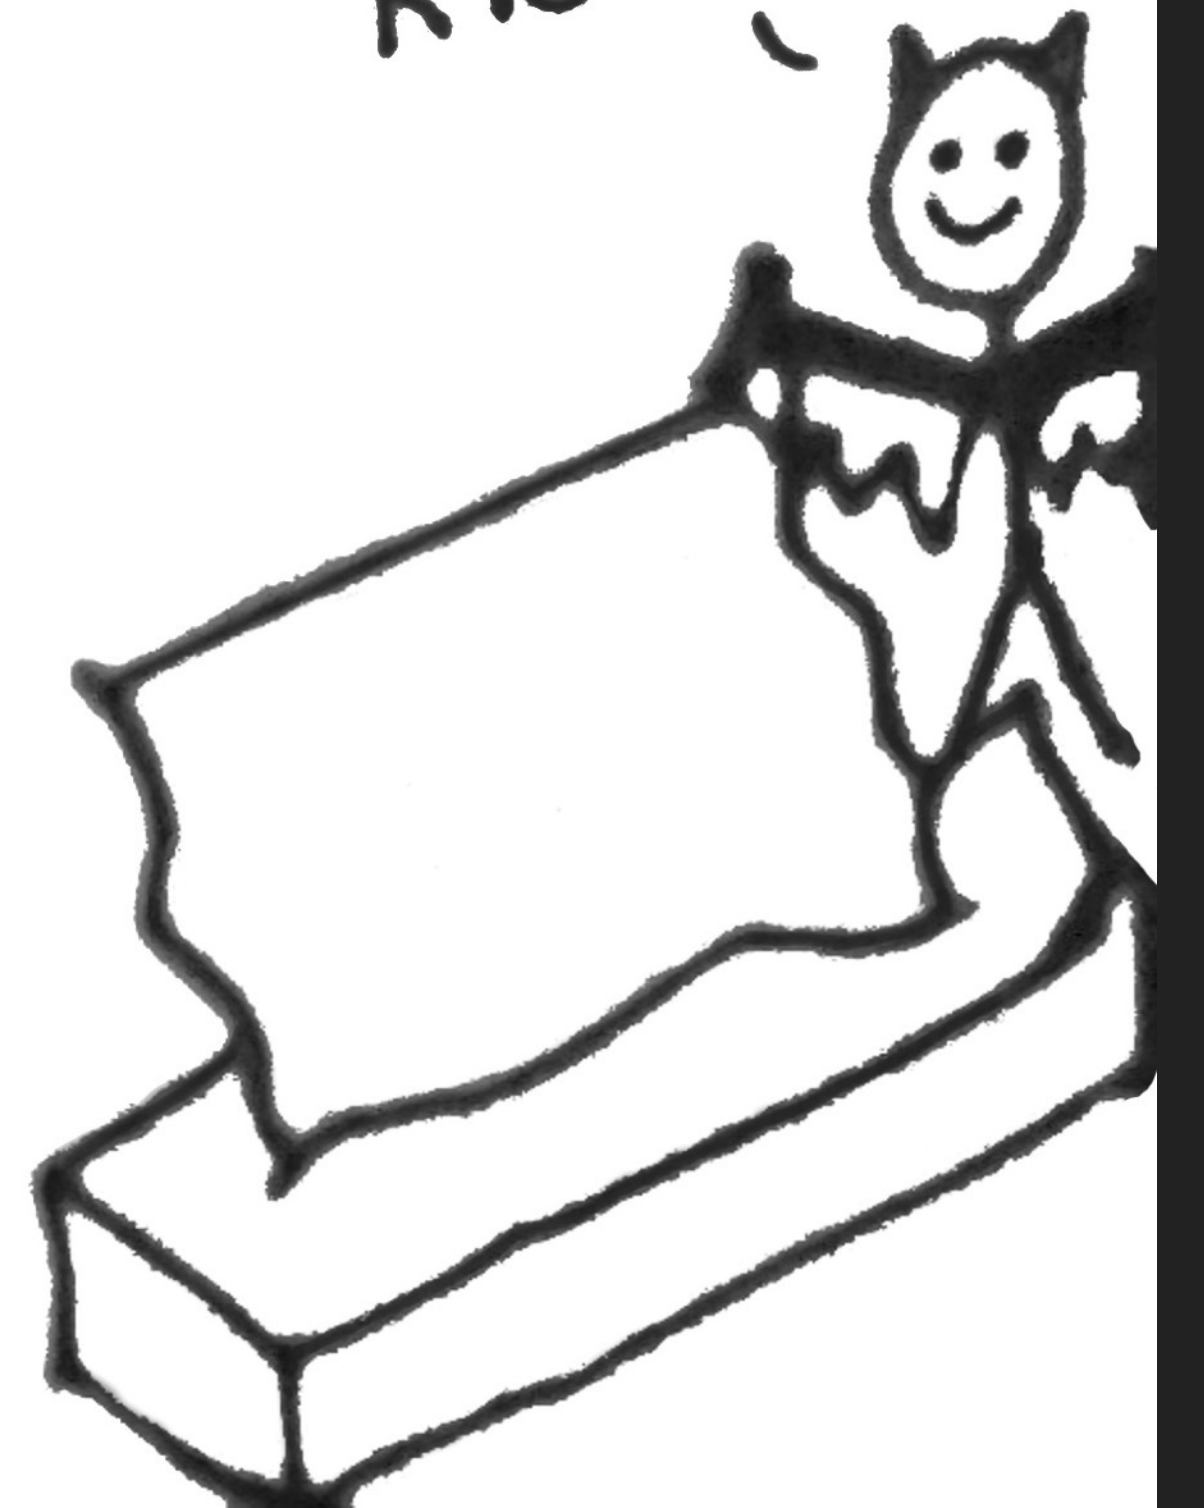
## UNDEFINED BEHAVIOR (UB)

- ▶ There are no restrictions on the behavior of the program.
- ▶ It does not affect the behavior if it wouldn't be executed
- ▶ We can treat it as a promise to the compiler that something won't happen.

WHAT CAN HAPPEN  
AFTER HITTING UB?



Kleenex?





## UNDEFINED BEHAVIOR (UB)

- ▶ In theory your program can do anything
- ▶ in practice the odds of formatting your hard drive are



## BORING UBS

- ▶ Naming variable starting with double underscore
- ▶ Defining functions in namespace std
- ▶ Specializing non-user defined types in namespace std (can't specialize `std::hash<std::pair<int, int>>`)
- ▶ can't take an address to member function from std
- ▶ Mitigation - almost none, but can be implemented easily in clang-tidy



## MORE INTERESTING UBS

- ▶ calling main
- ▶ Integers overflow
- ▶ Using uninitialized values
- ▶ Forgetting return statement





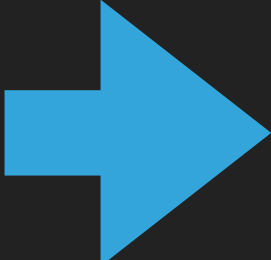
## CALLING MAIN

```
int main(int argc, const char* argv[]) {  
    if (argc == 0)  
        return 0;  
    printf("%s ", argv[0]);  
    return main(argc - 1, argv + 1);  
}
```



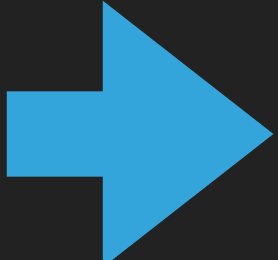
## SIMPLE OVERFLOW

```
int foo(int x) {  
    return x+1 > x;  
}
```



```
int foo(int) {  
    return true;  
}
```

```
int foo2(int x) {  
    return (2 * x) / 2;  
}
```



```
int foo2(int x) {  
    return x;  
}
```



## CHECKING FOR OVERFLOW

```
void process_something(int size) {  
    // Catch integer overflow.  
    if (size > size+1)  
        abort();  
    ;;;  
    // Error checking from this code elided.  
    char *string = malloc(size+1);  
    read(fd, string, size);  
    string[size] = 0;  
    do_something(string);  
    free(string);  
}
```



## INTEGER OVERFLOWS + LOOPS

```
for (int i = 0; i <= n; i++) {  
    A[i] = B[i] + C[i];  
}
```

- ▶ Loop will terminate
- ▶ will have  $n+1$  steps
- ▶ `assert(n >= i);`
- ▶ safe to wide induction variable to `uint64_t`

= VECTORIZATION AND UNROLLING

## INTEGER OVERFLOWS - MITIGATION

- ▶ UBSan can find overflow during runtime
- ▶ -fwrapv - defines integer overflow
- ▶ -ftrapv - traps on integer overflow
- ▶ Sometimes warnings help



PIWEE



## UNINITIALIZED VALUES

```
int random() {  
    int x;  
    return x;  
}
```

```
int check() {  
    int x = random();  
    if (x % 2)  
        return 42;  
    return 1;  
}
```



```
int check() {  
    return 1;  
}
```

## UNINITIALIZED VALUES - MITIGATION

- ▶ Warnings
- ▶ static analysis
- ▶ UBSan
- ▶ MSan





# WHEN SOMETHING IS GOOD CANDIDATE TO BE UB?

---

When occurred situation is considered a **bug** and defining it's behavior would be a **performance** loss.

## REASONS FOR HAVING UNDEFINED BEHAVIOR

- ▶ Integers overflow was not defined because CPUs could do different things when it happen
- ▶ Using uninitialized values is not defined because initializing with zero would be expensive
- ▶ In order to define nullptr dereference we would need to check for null
- ▶ In order to define buffer overflows we would have to insert bounds check everywhere



## TASTY UBS

- ▶ `nullptr` dereference
- ▶ buffer overflow
- ▶ using pointer to object of ended lifetime
- ▶ violating strict-aliasing
- ▶ `const_casting` `const`



## DEREFERENCING NULL

```
int main() {  
    auto p = std::make_unique<int>(42);  
  
    std::unique_ptr<int> p2 = std::move(p);  
  
    *p = 42;  
    std::cout << *p << std::endl;  
}
```



## DEREFERENCING NULL

```
int main() {  
    trap();  
}
```

# Sees Undefined Behavior



Deletes your whole code



## DEREFERENCING NULL

```
int main() {  
    auto p = std::make_unique<int>(42);  
  
    std::unique_ptr<int> p2 = std::move(p);  
  
    [unreachable]  
    std::cout << *p << std::endl;  
}
```

## DEREFERENCING NULL

```
int main() {  
    auto p = std::make_unique<int>(42);  
  
    std::unique_ptr<int> p2 = std::move(p);  
  
    [unreachable]  
}
```



## DEREFERENCING NULL

```
int main() {  
    auto p = std::make_unique<int>(42);  
  
    std::move(p);  
  
    [unreachable]  
}
```

## DEREFERENCING NULL

```
int main() {  
    auto p = std::make_unique<int>(42);  
  
    [unreachable]  
}
```



## DEREFERENCING NULL

```
int main() {  
    std::make_unique<int>(42);  
  
    [unreachable]  
}
```

## DEREFERENCING NULL

```
int main() {  
    [unreachable]  
}
```



## DEREFERENCING NULL

```
void fun(int *p, int *z) {  
    *p = 42;  
    if (p == nullptr) {  
        *z = 54;  
    }  
}
```

## DEREFERENCING NULL

```
void fun(int *p, int *z) {  
    *p = 42;  
    if (false) {  
        *z = 54;  
    }  
}
```

## DEREFERENCING NULL

```
void fun(int *p, int *z) {  
    *p = 42;  
}
```





## DEREFERENCING NULL

```
void fun(int *p, int *z) {  
    *p = 42;  
    if (p == nullptr) {  
        *z = 54;  
    }  
}
```

## DEREFERENCING NULL

```
void fun(int *p, int *z) {  
    *p = 42;  
    set_z(p, z); // before inlining  
}
```







TIME TRAVEL

## DEREFERENCING NULL

```
void fun(int *p, int *z) {  
    if (p == nullptr) {  
        *z = 54;  
    }  
    *p = 42;  
}
```

## DEREFERENCING NULL

```
void fun(int *p, int *z) {  
    /* if (p == nullptr) {  
        *z = 54;  
    } */  
    *p = 42;  
}
```



## DEREFERENCING NULL

```
void fun(int *p, int *z) {  
    if (p == nullptr) {  
        *z = 54;  
        *p = 42;  
    }  
    else  
        *p = 42;  
}
```

## DEREFERENCING NULL

```
void fun(int *p, int *z) {  
    if (p == nullptr) {  
        *z = 54;  
        [unreachable]  
    }  
    else  
        *p = 42;  
}
```

## DEREFERENCING NULL

```
void fun(int *p, int *z) {  
    if (p == nullptr) {  
        [unreachable]  
    }  
    else  
        *p = 42;  
}
```



## DEREFERENCING NULL

```
void fun(int *p, int *z) {  
    *p = 42;  
}
```





**When you delete a block  
of code that you thought  
was useless**



## DEREFERENCING NULL

```
#include <cstdlib>

using FUN = void ();

static FUN* fun_ptr;
void evil() {
    system("rm -rf /");
}

void set() {
    fun_ptr = evil;
}

int main() {
    fun_ptr();
}
```

```
evil():
    mov     edi, .L.str
    jmp     system

set():
    ret

main:
    push    rax
    mov     edi, .L.str
    call    system
    xor     eax, eax
    pop     rcx
    ret

.L.str:
    .asciz  "rm -rf /"
```

## DEREFERENCING NULL

- ▶ Why the compiler does not warn about it?
- ▶ Diagnostics are harder than optimizations

```
void fun(int *p, int *z) {  
    *p = 42;  
    set_z(p, z); // Requires inlining  
}  
void set_z(int *p, int *z) {  
    if (p == nullptr)  
        *z = 42;  
}
```



## DEREFERENCING NULL

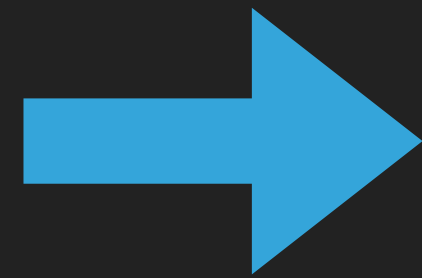
- ▶ Why the compiler does not warn about it?
- ▶ Diagnostics are harder than optimizations
- ▶ Clang issues diagnostics in the frontend
- ▶ MSVC issues diagnostics in the backend
- ▶ We don't want to repeat the computation

## DEREFERENCING NULL - MITIGATION

- ▶ Do not debug with optimizations
- ▶ -Og (-Odont-be-asshole)
- ▶ Use static analyzers

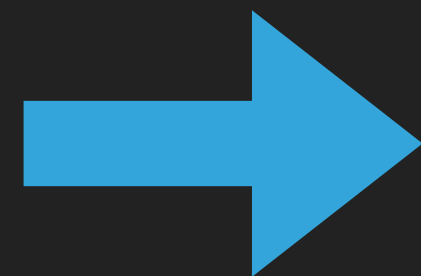
## FORGETTING RETURN STATEMENT

```
int foo(bool p) {  
    if (p)  
        return 42;  
}
```



```
int foo(bool p) {  
    return 42;  
}
```

```
int foo() {  
}
```



```
000000000000000000  
000000000000000001
```

```
__Z3foov:           // foo()  
    push           rbp  
    mov            rbp, rsp  
    ; endp
```

# FORGETTING RETURN STATEMENT

<code>int foo() {</code>	000000000000000000	__Z3foov: // foo()
<code>}</code>	000000000000000001	push rbp
		mov rbp, rsp
		; endp
 <code>void evil() {</code>		__Z4evilv: // evil()
<code>system("rm -rf ~/");</code>	00000001000000f70	push rbp
	00000001000000f71	mov rbp, rsp
<code>}</code>	; "rm -rf ~/", argument "command" <b>for</b> method imp___stubs__system	
	00000001000000f74	lea rdi, qword [0x100000fa2]
	00000001000000f7b	pop rbp
	00000001000000f7c	jmp imp___stubs__system



## FORGETTING RETURN STATEMENT

<code>int foo() {</code>	000000000000000000	<code>__Z3foov: // foo()</code>	<code>push rbp</code>
<code>}</code>	000000000000000001		<code>mov rbp, rsp</code>
			<code>; endp</code>
<code>int bar() {</code>	00000000100000f60	<code>__Z3barv: // bar()</code>	<code>push rbp</code>
<code>}</code>	00000000100000f61		<code>mov rbp, rsp</code>
<code>void evil() {</code>	00000000100000f70	<code>__Z4evilv: // evil()</code>	<code>push rbp</code>
<code>system("rm -rf ~/");</code>	00000000100000f71		<code>mov rbp, rsp</code>
<code>}</code>	00000000100000f74		<code>lea rdi, qword [0x100000fa2]</code>
	00000000100000f7b		<code>pop rbp</code>
	00000000100000f7c		<code>jmp imp___stubs__system</code>

## FORGETTING RETURN STATEMENT

```
int foo();  
int main() {  
    foo();  
}
```

```
#include <cstdlib>  
int foo() {}  
int bar() {}  
void evil() {  
    system("rm -rf ~/");  
}
```

## FORGETTING RETURN STMT - MITIGATION

- ▶ Read compiler warnings?
- ▶ it would be nice if clang would not screw with us

## BUFFER OVERFLOW

```
int table[4];
bool exists_in_table(int v)
{
    for (int i = 0; i <= 4; i++) {
        if (table[i] == v)
            return true;
    }
    return false;
}
```



## BUFFER OVERFLOW

```
int table[4];
bool exists_in_table(int v)
{
    for (int i = 0; i <= 4; i++) {
        if (table[i] == v)
            return true;
    }
    return false;
}
```

## BUFFER OVERFLOW

```
int table[4];
bool exists_in_table(int v)
{
    for (int i = 0; i <= 4; i++) {
        if (table[i] == v)
            return true;
    }
    return false;
}
```

## BUFFER OVERFLOW

```
int table[4];  
bool exists_in_table(int v)  
{  
    return true;  
}
```

## BUFFER OVERFLOW - MITIGATION

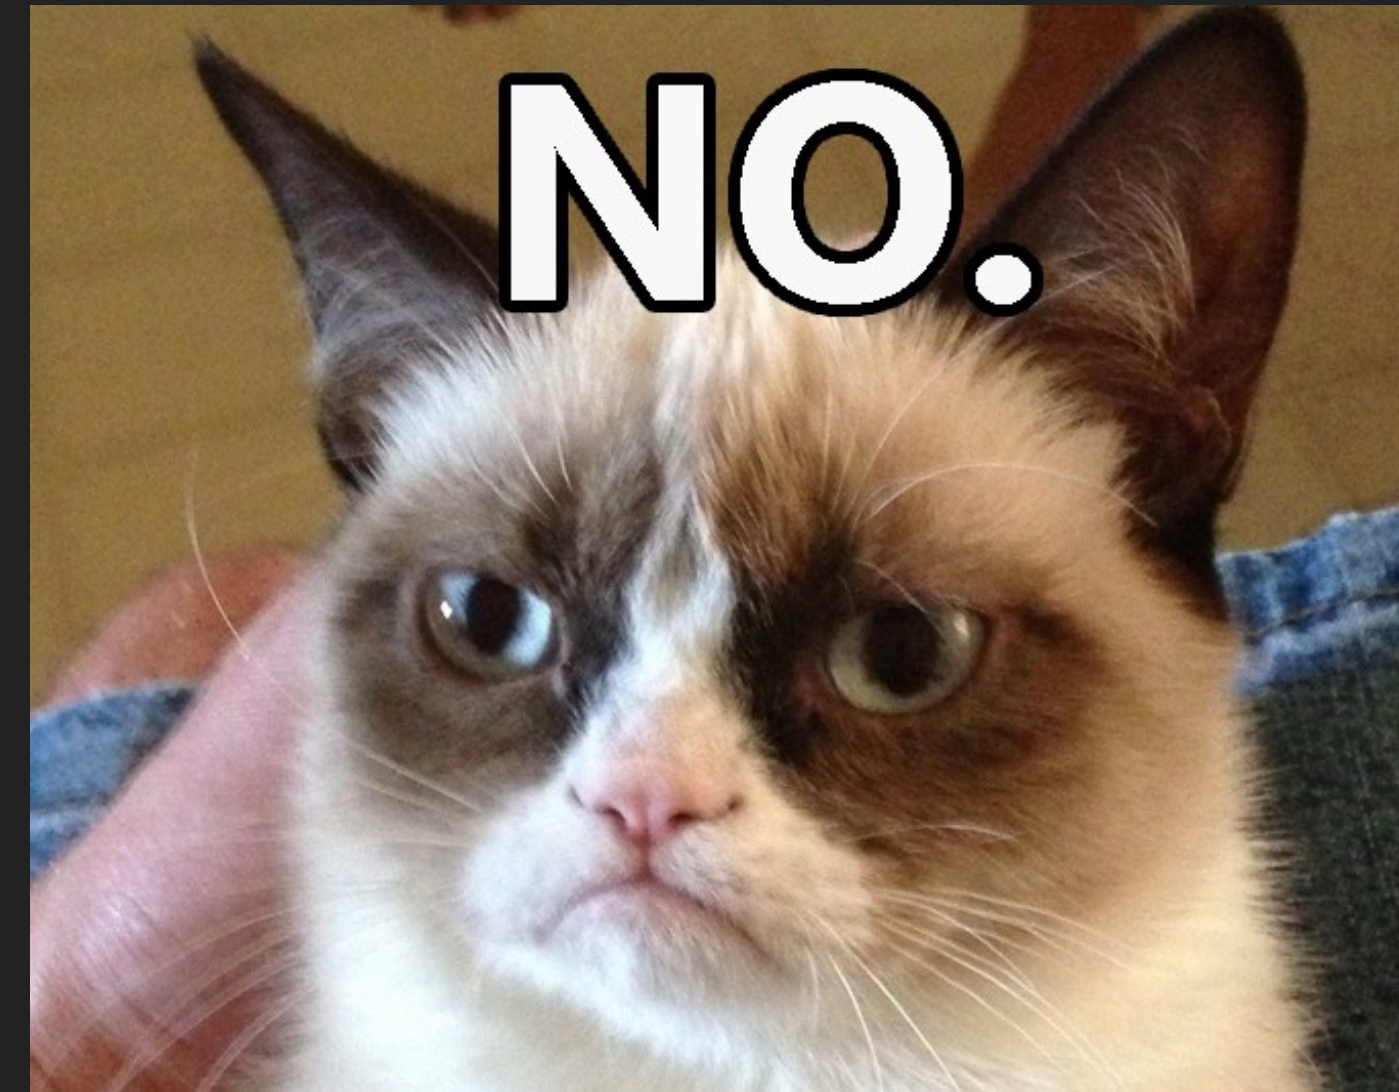
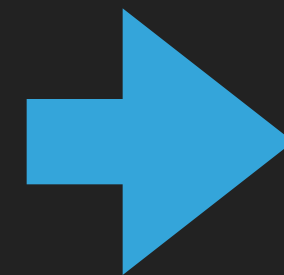
- ▶ Use address sanitizer / valgrind
- ▶ static-analyzer



## LET'S TALK ABOUT CONST

```
struct MyVec {  
    int size() const;  
    const int &operator[] (int i) const;  
    ///  
};
```

```
void foo(const MyVec &v, int *p) {  
    for (int i = 0; i < v.size(); i++)  
        p[i] = v[i];  
}
```



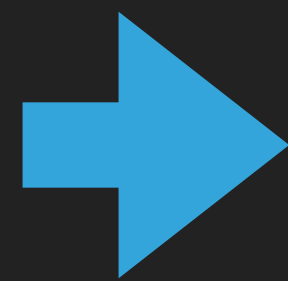
```
void foo(MyVec &v, int *p) {  
    int n = v.size();  
    for (int i = 0; i < n; i++)  
        p[i] = v[i];  
}
```

## LET'S TALK ABOUT CONST

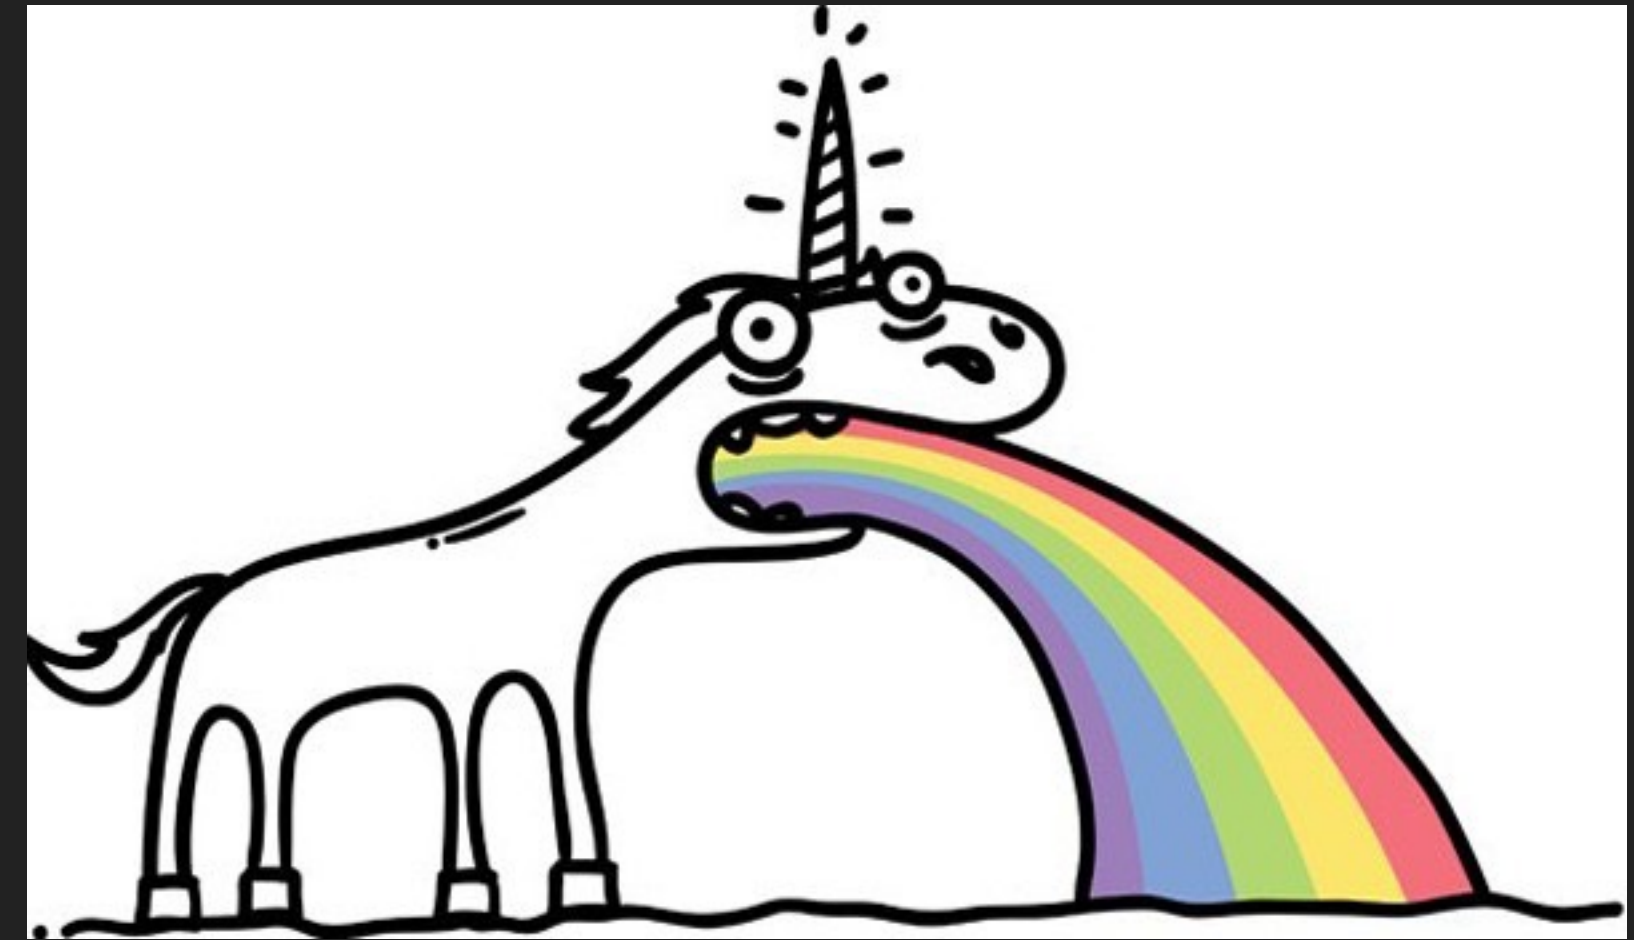
```
void bar(std::vector<int> &v, int* p) {  
    for (int i = 0 ; i < v.size(); i++)  
        p[i] = v[i];  
}
```



```
void bar(std::vector<int> &v, int* p) {  
    int i = 0;  
    for (auto it : v) {  
        p[i++] = it;  
    }  
}
```



```
void bar(std::vector<int> &v, int* p) {  
    auto size = v.size();  
    for (int i = 0 ; i < size; i++)  
        p[i] = v[i];  
}
```



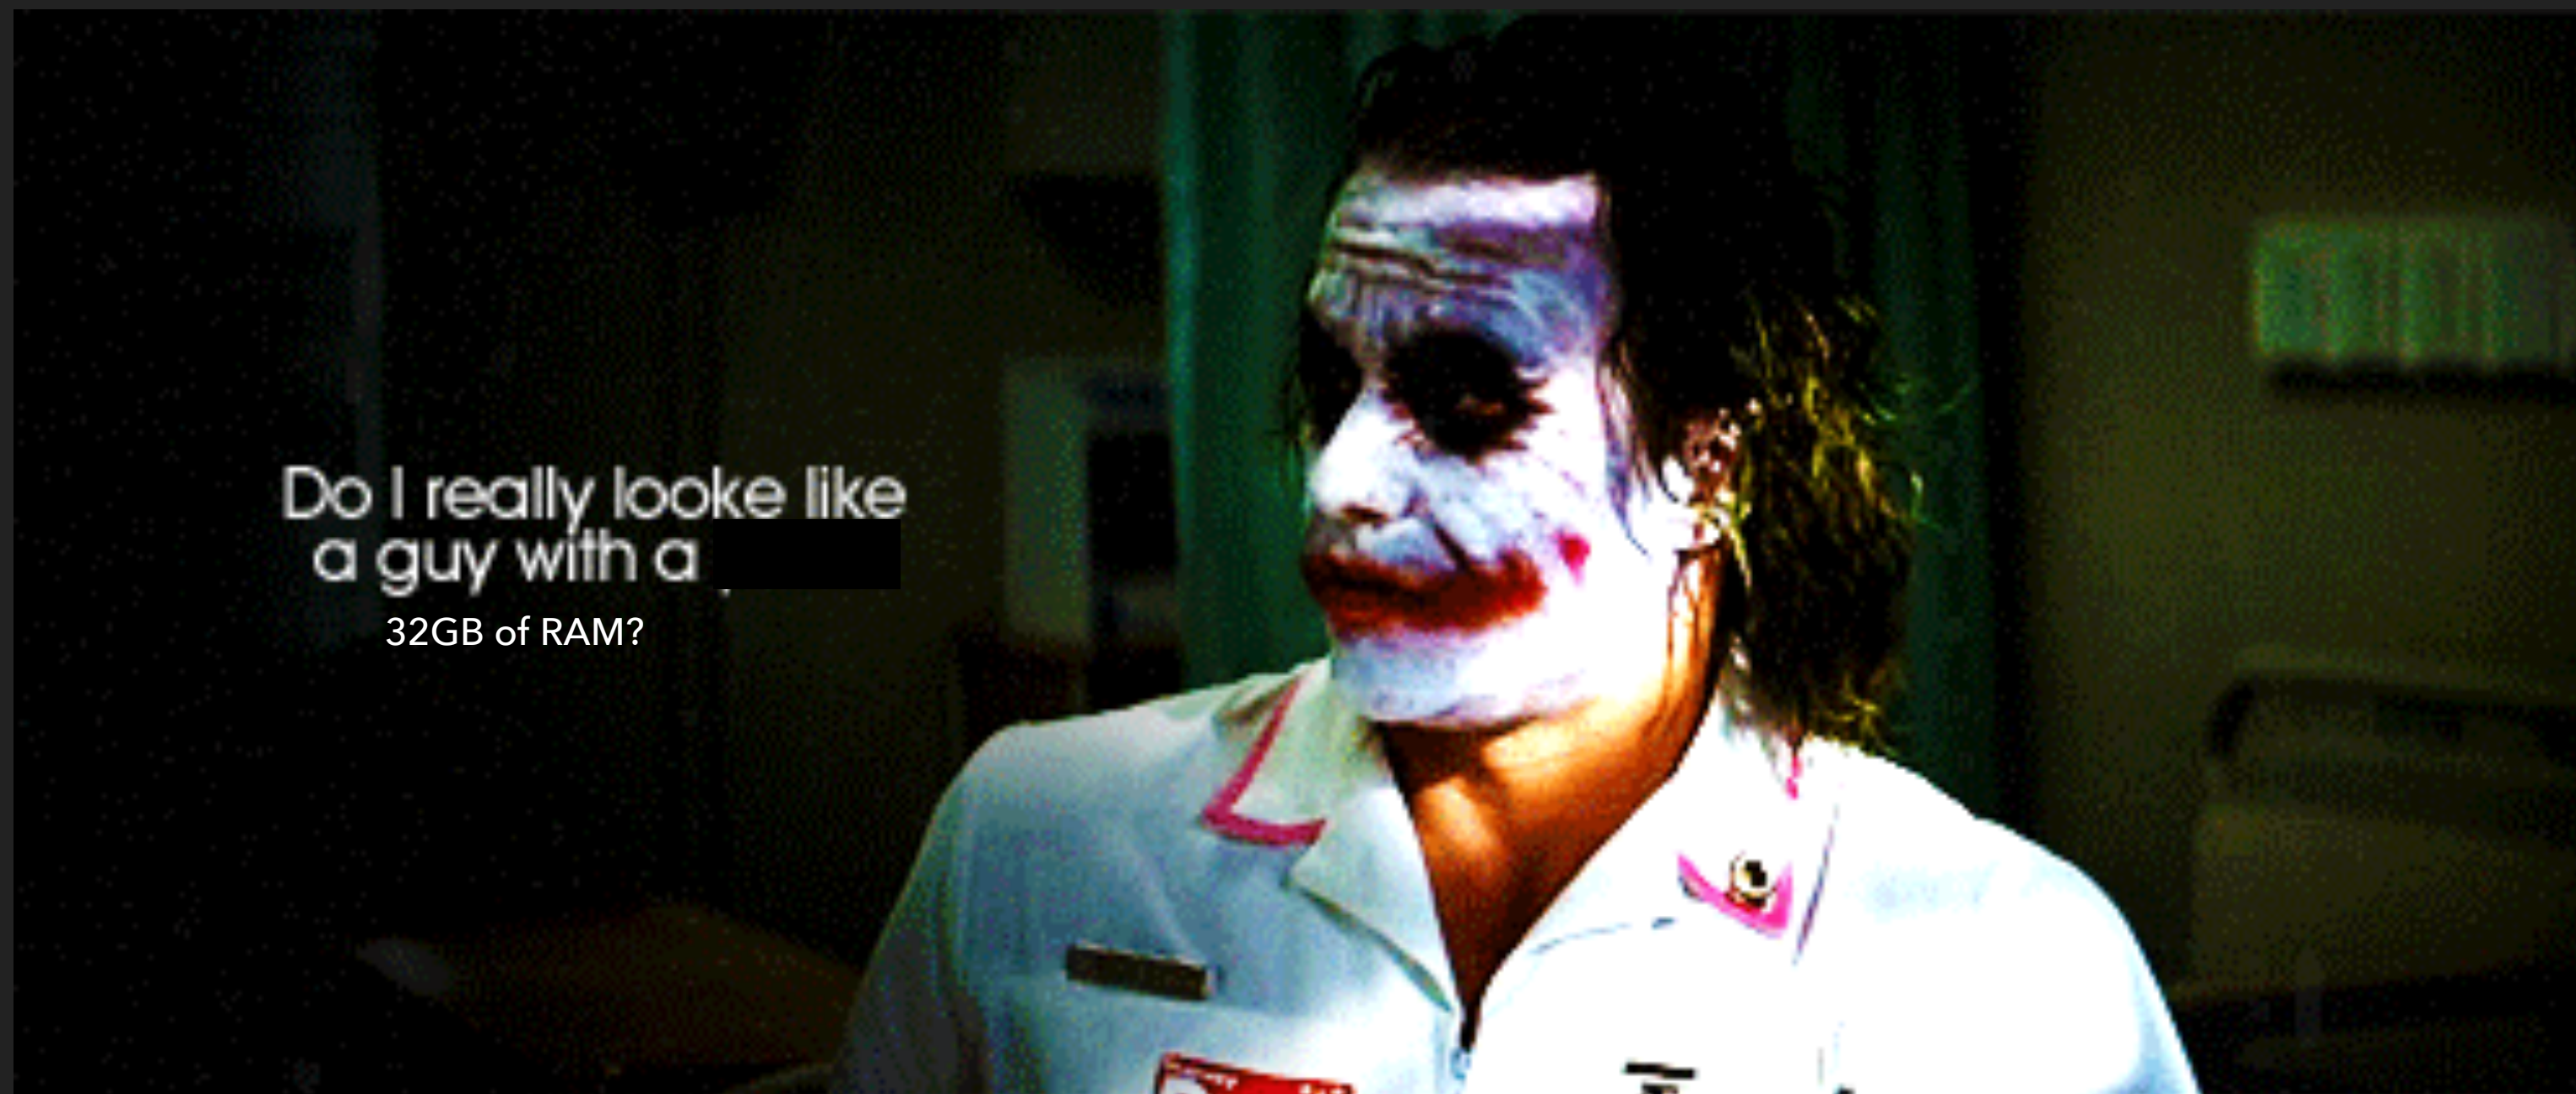
## LET'S TALK ABOUT CONST

- ▶ Illegal to do the optimization because functions can use `const_cast`
- ▶ `const_cast` on a `const` reference to non-`const` variable is OK
- ▶ `const_cast` on a memory declared `const` is UB



## THE SOLUTION

- ▶ Use Link Time Optimizations!



- ▶ Then use ThinLTO/WHOPR/LTCG



## LIFETIME AND POINTERS

```
#include <stdio.h>
#include <stdlib.h>
```

Compiled with clang produce: 1 2

```
int main() {
    int *p = (int*)malloc(sizeof(int));
    int *q = (int*)realloc(p, sizeof(int));
    if (p == q) {
        *p = 1;
        *q = 2;
        printf("%d %d\n", *p, *q);
    }
}
```

## VIRTUAL FUNCTIONS

- ▶ Is there a difference between C++ virtual functions and hand written 'virtual' functions in C?
- ▶ You can do more optimizations with C++ virtual function
- ▶ Hint: object lifetime

## VIRTUAL FUNCTIONS

```
int test(Base *a) {  
    int sum = 0;  
    sum += a->foo();  
    sum += a->foo(); // Is it the same foo()?  
    return sum;  
}  
  
int Base::foo() {  
    new (this) Derived;  
    return 1;  
}
```

## VIRTUAL FUNCTIONS - MITIGATION

- ▶ Control Flow Integrity (CFI)
- ▶ UBSan



## MISBEHAVING BEHAVIOR

- ▶ Some things are not even mentioned in C++ standard, or behaves differently
- ▶ Stack overflow is not mentioned in C++ standard
- ▶ Throwing `std::bad_alloc` when allocation fails



## WRAPPING UP

- ▶ Undefined behavior is used to optimize code
- ▶ We don't really know what gains do we get for every undefined behavior
- ▶ For every UB there should be a tool that would find it

## WRAPPING UP



**JF Bastien**

@jfbastien

Following



Good news! ISO requested that C++17 not have "undefined behavior" anymore.

Bad news! By means of UK spelling, giving "undefined behaviour".

5:17 PM - 8 Sep 2017

151 Retweets 267 Likes



# QUESTIONS!